

Enhancing constraint programming with machine learning

Current challenges and future opportunities

CP 2024 - Girona



POLYTECHNIQUE
MONTRÉAL

CORAIL



Mila



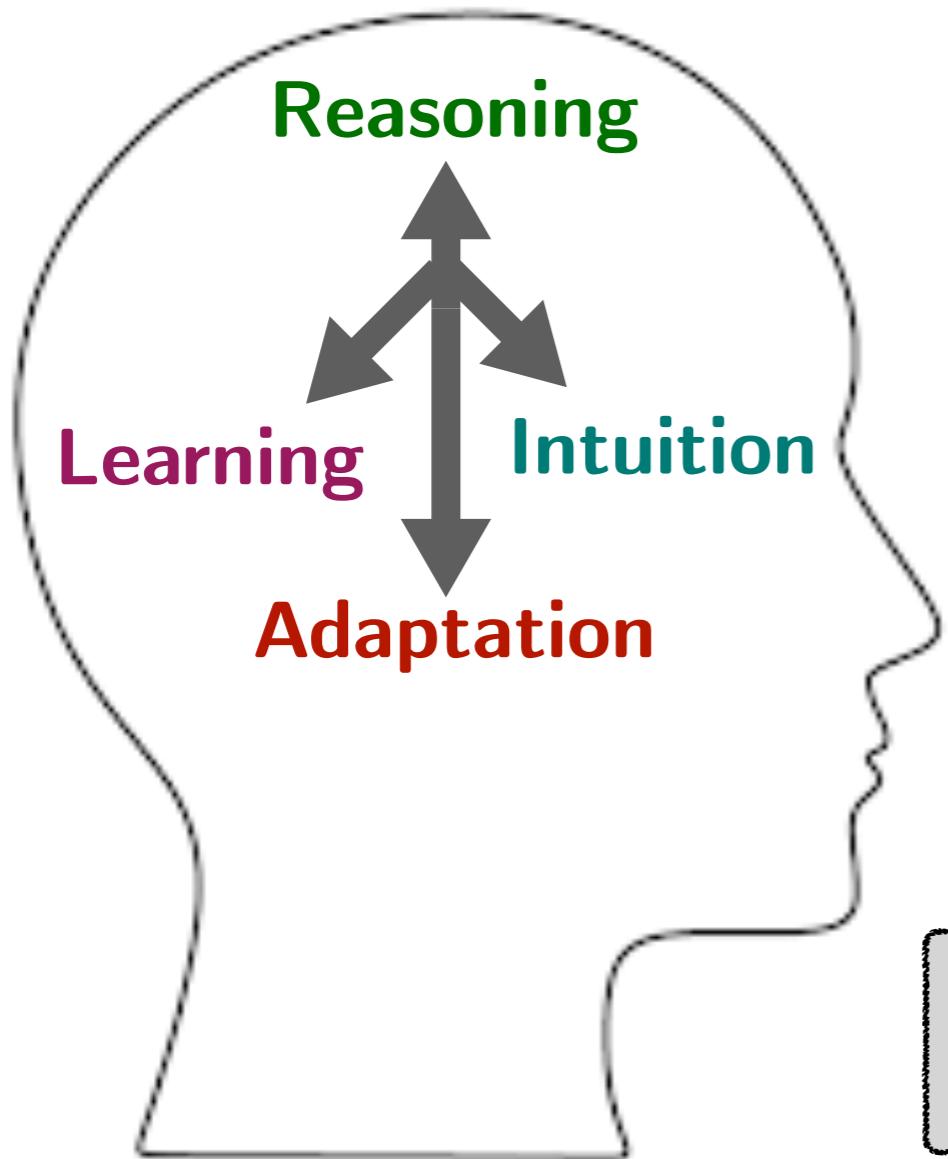
Combinatorial Optimization and
Reasoning in
Artificial Intelligence
Laboratory



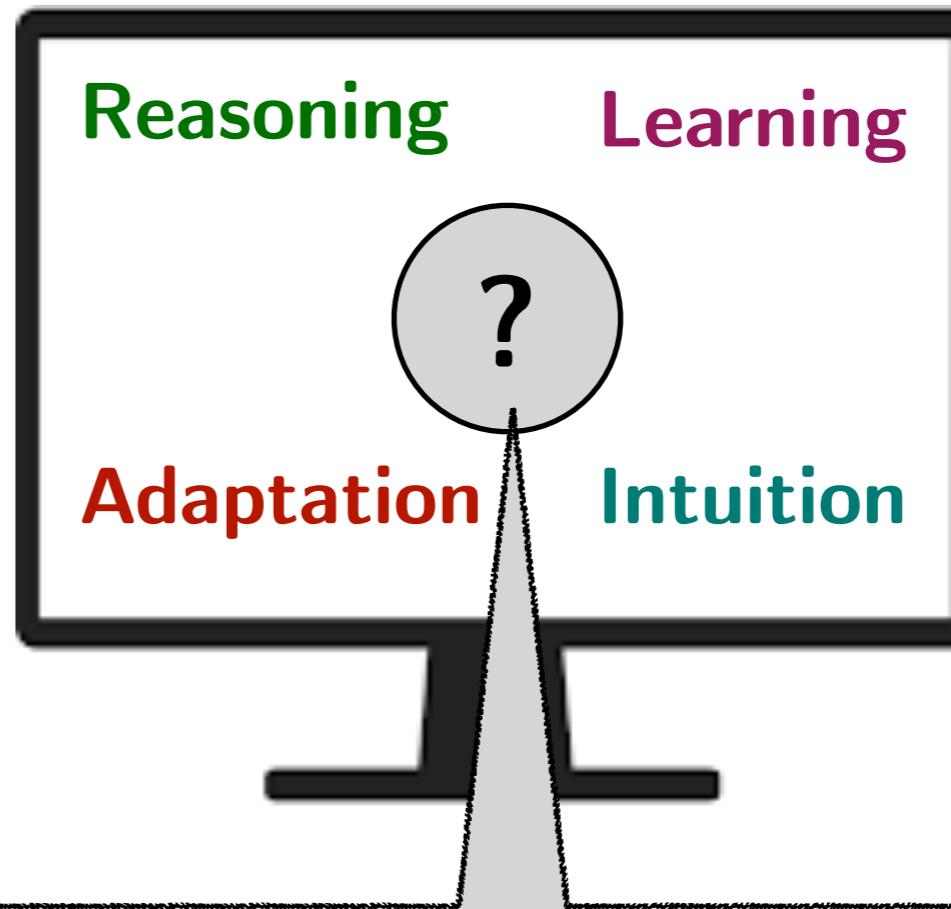
ACP
Quentin Cappart

Human intelligence versus artificial intelligence

Human intelligence



Artificial intelligence

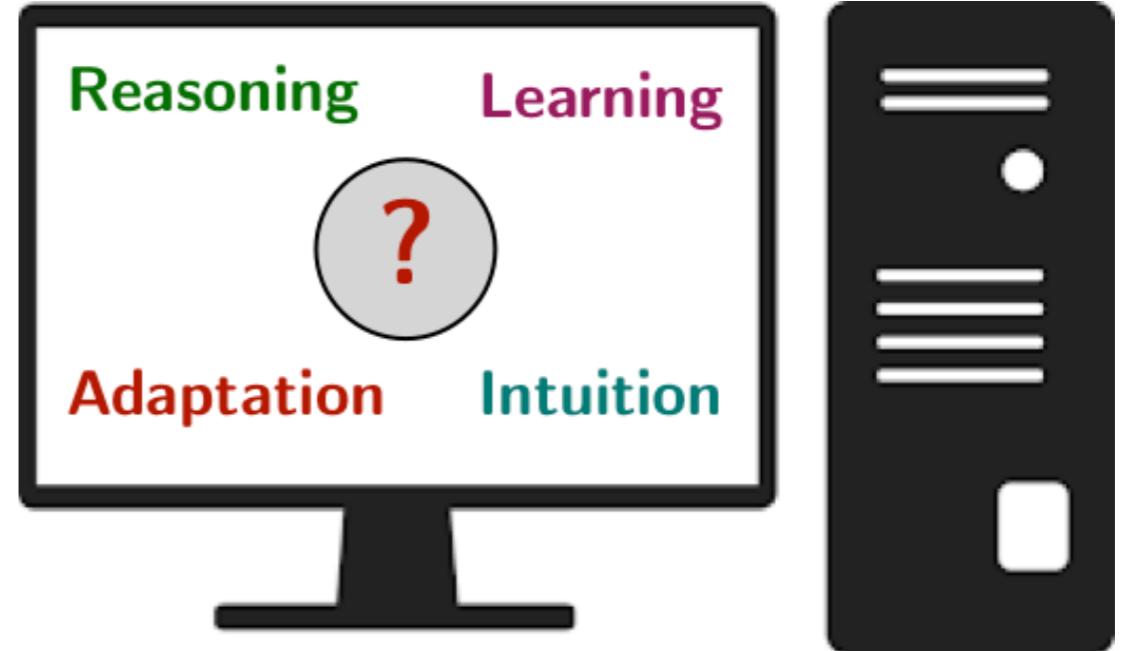


This connection is not yet established

Long-term research plan: building an AI with these connections

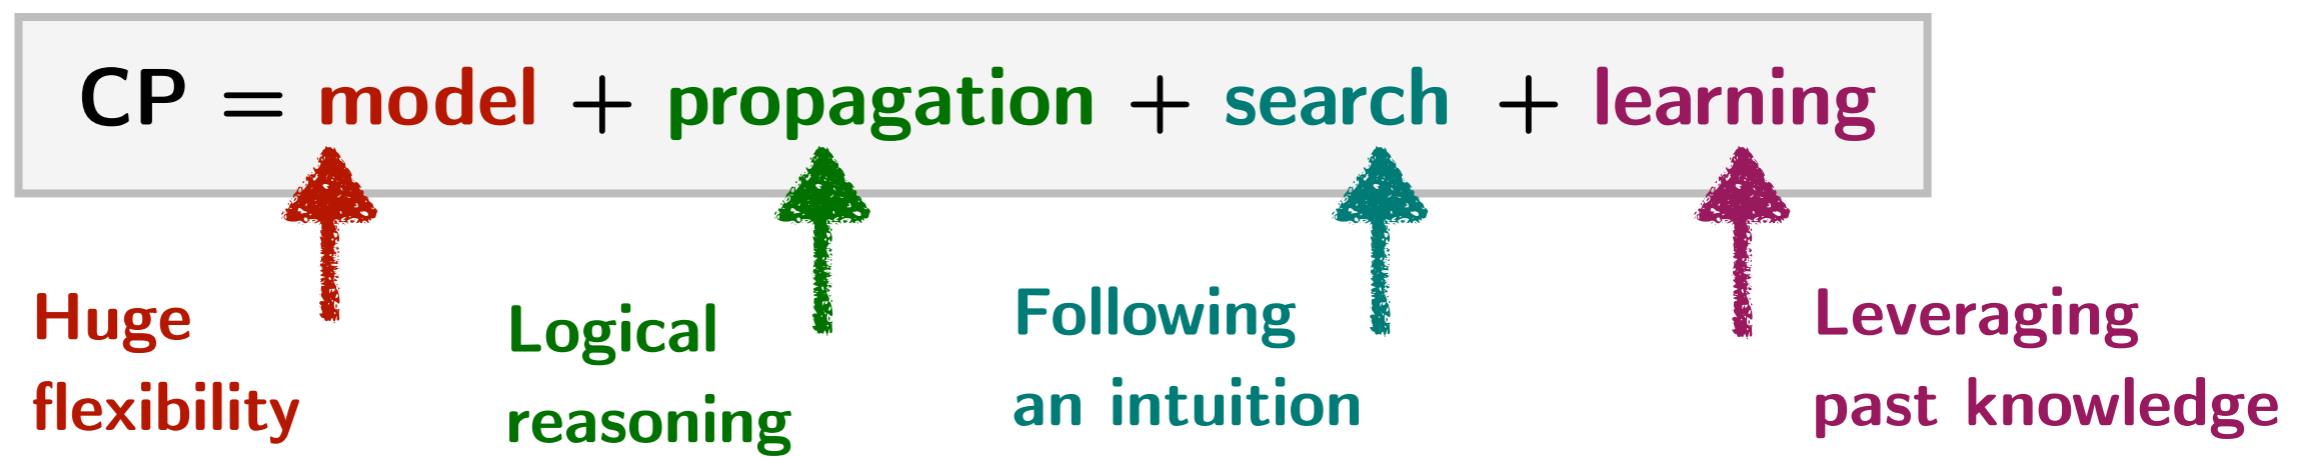
Goal: providing a better solving process for combinatorial problems

Constraint programming as a unifying framework



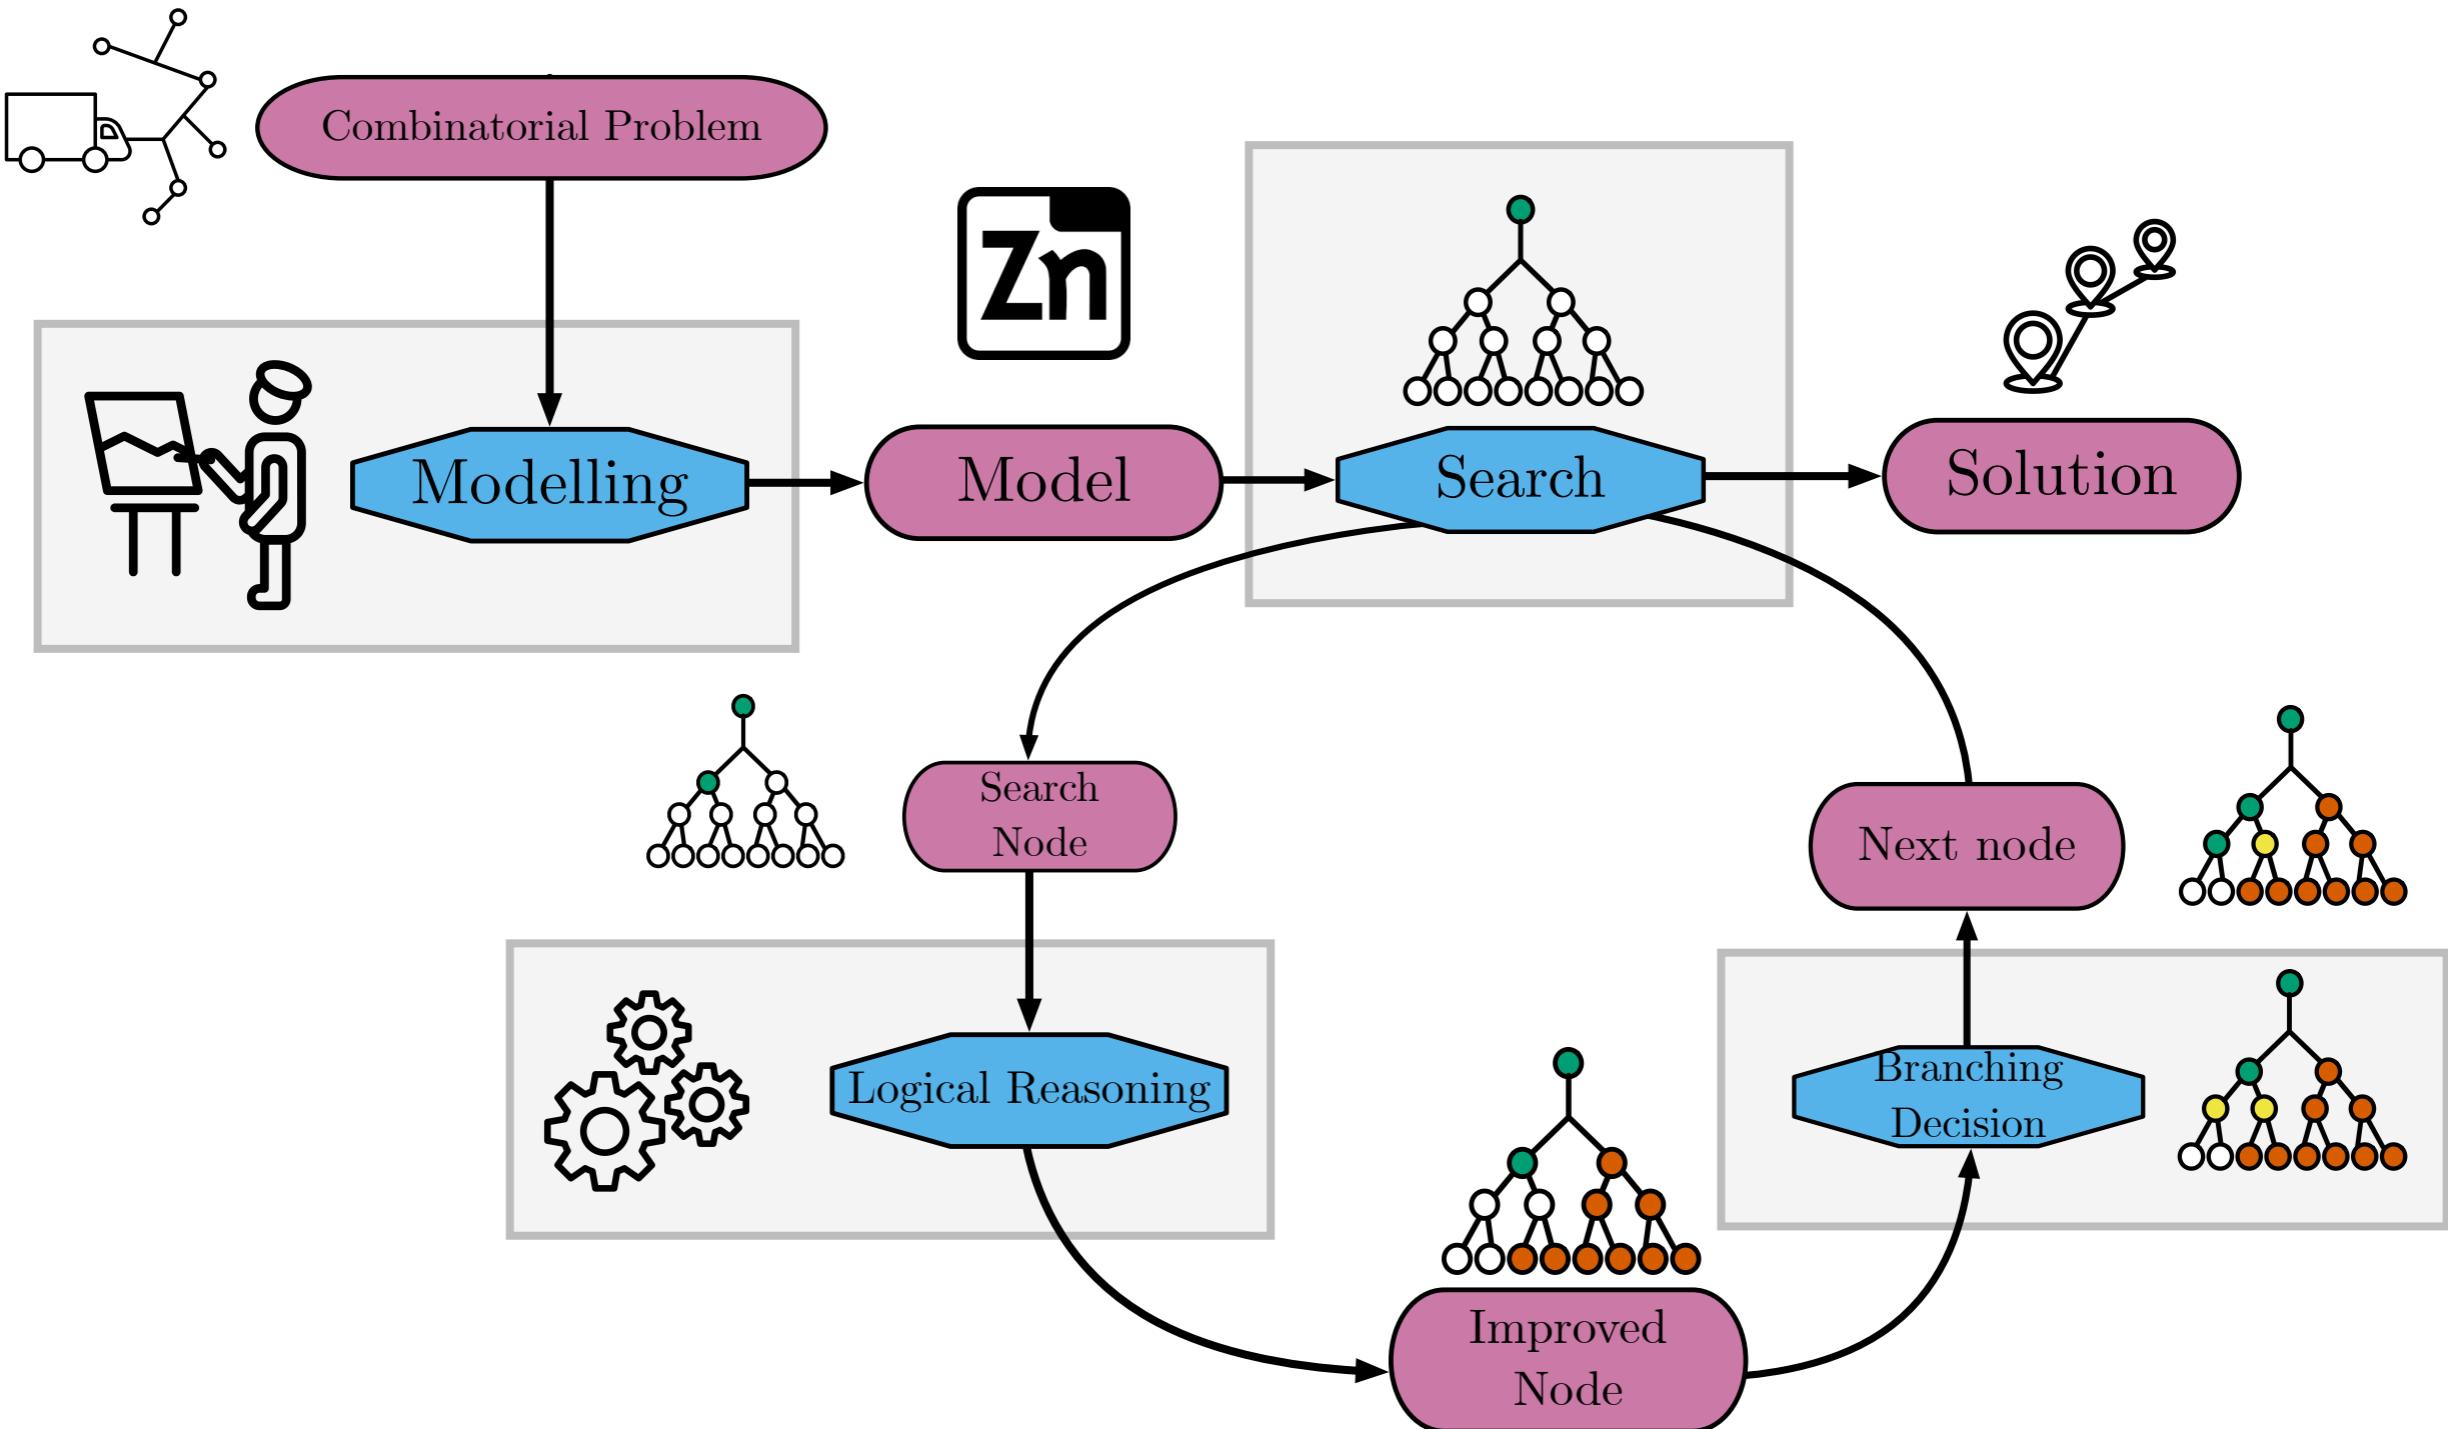
My research hypothesis

Constraint programming can be a hosting technology for building this hybrid AI



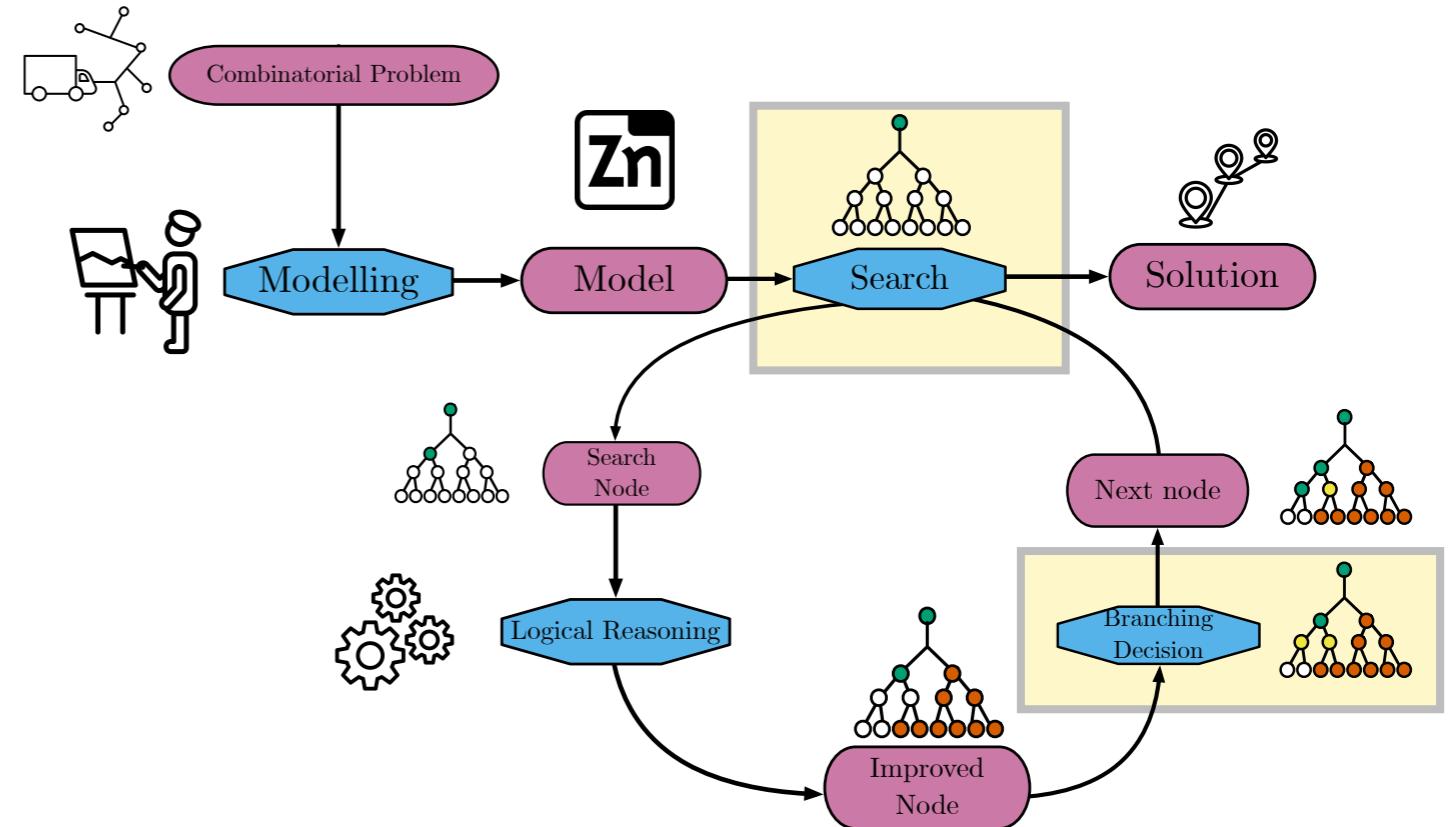
? How learning can be used to enhance the efficiency of CP solvers ?

Summary of the CP pipeline



Each decisional step has the potential to be improved with learning

Enhancing CP with learning



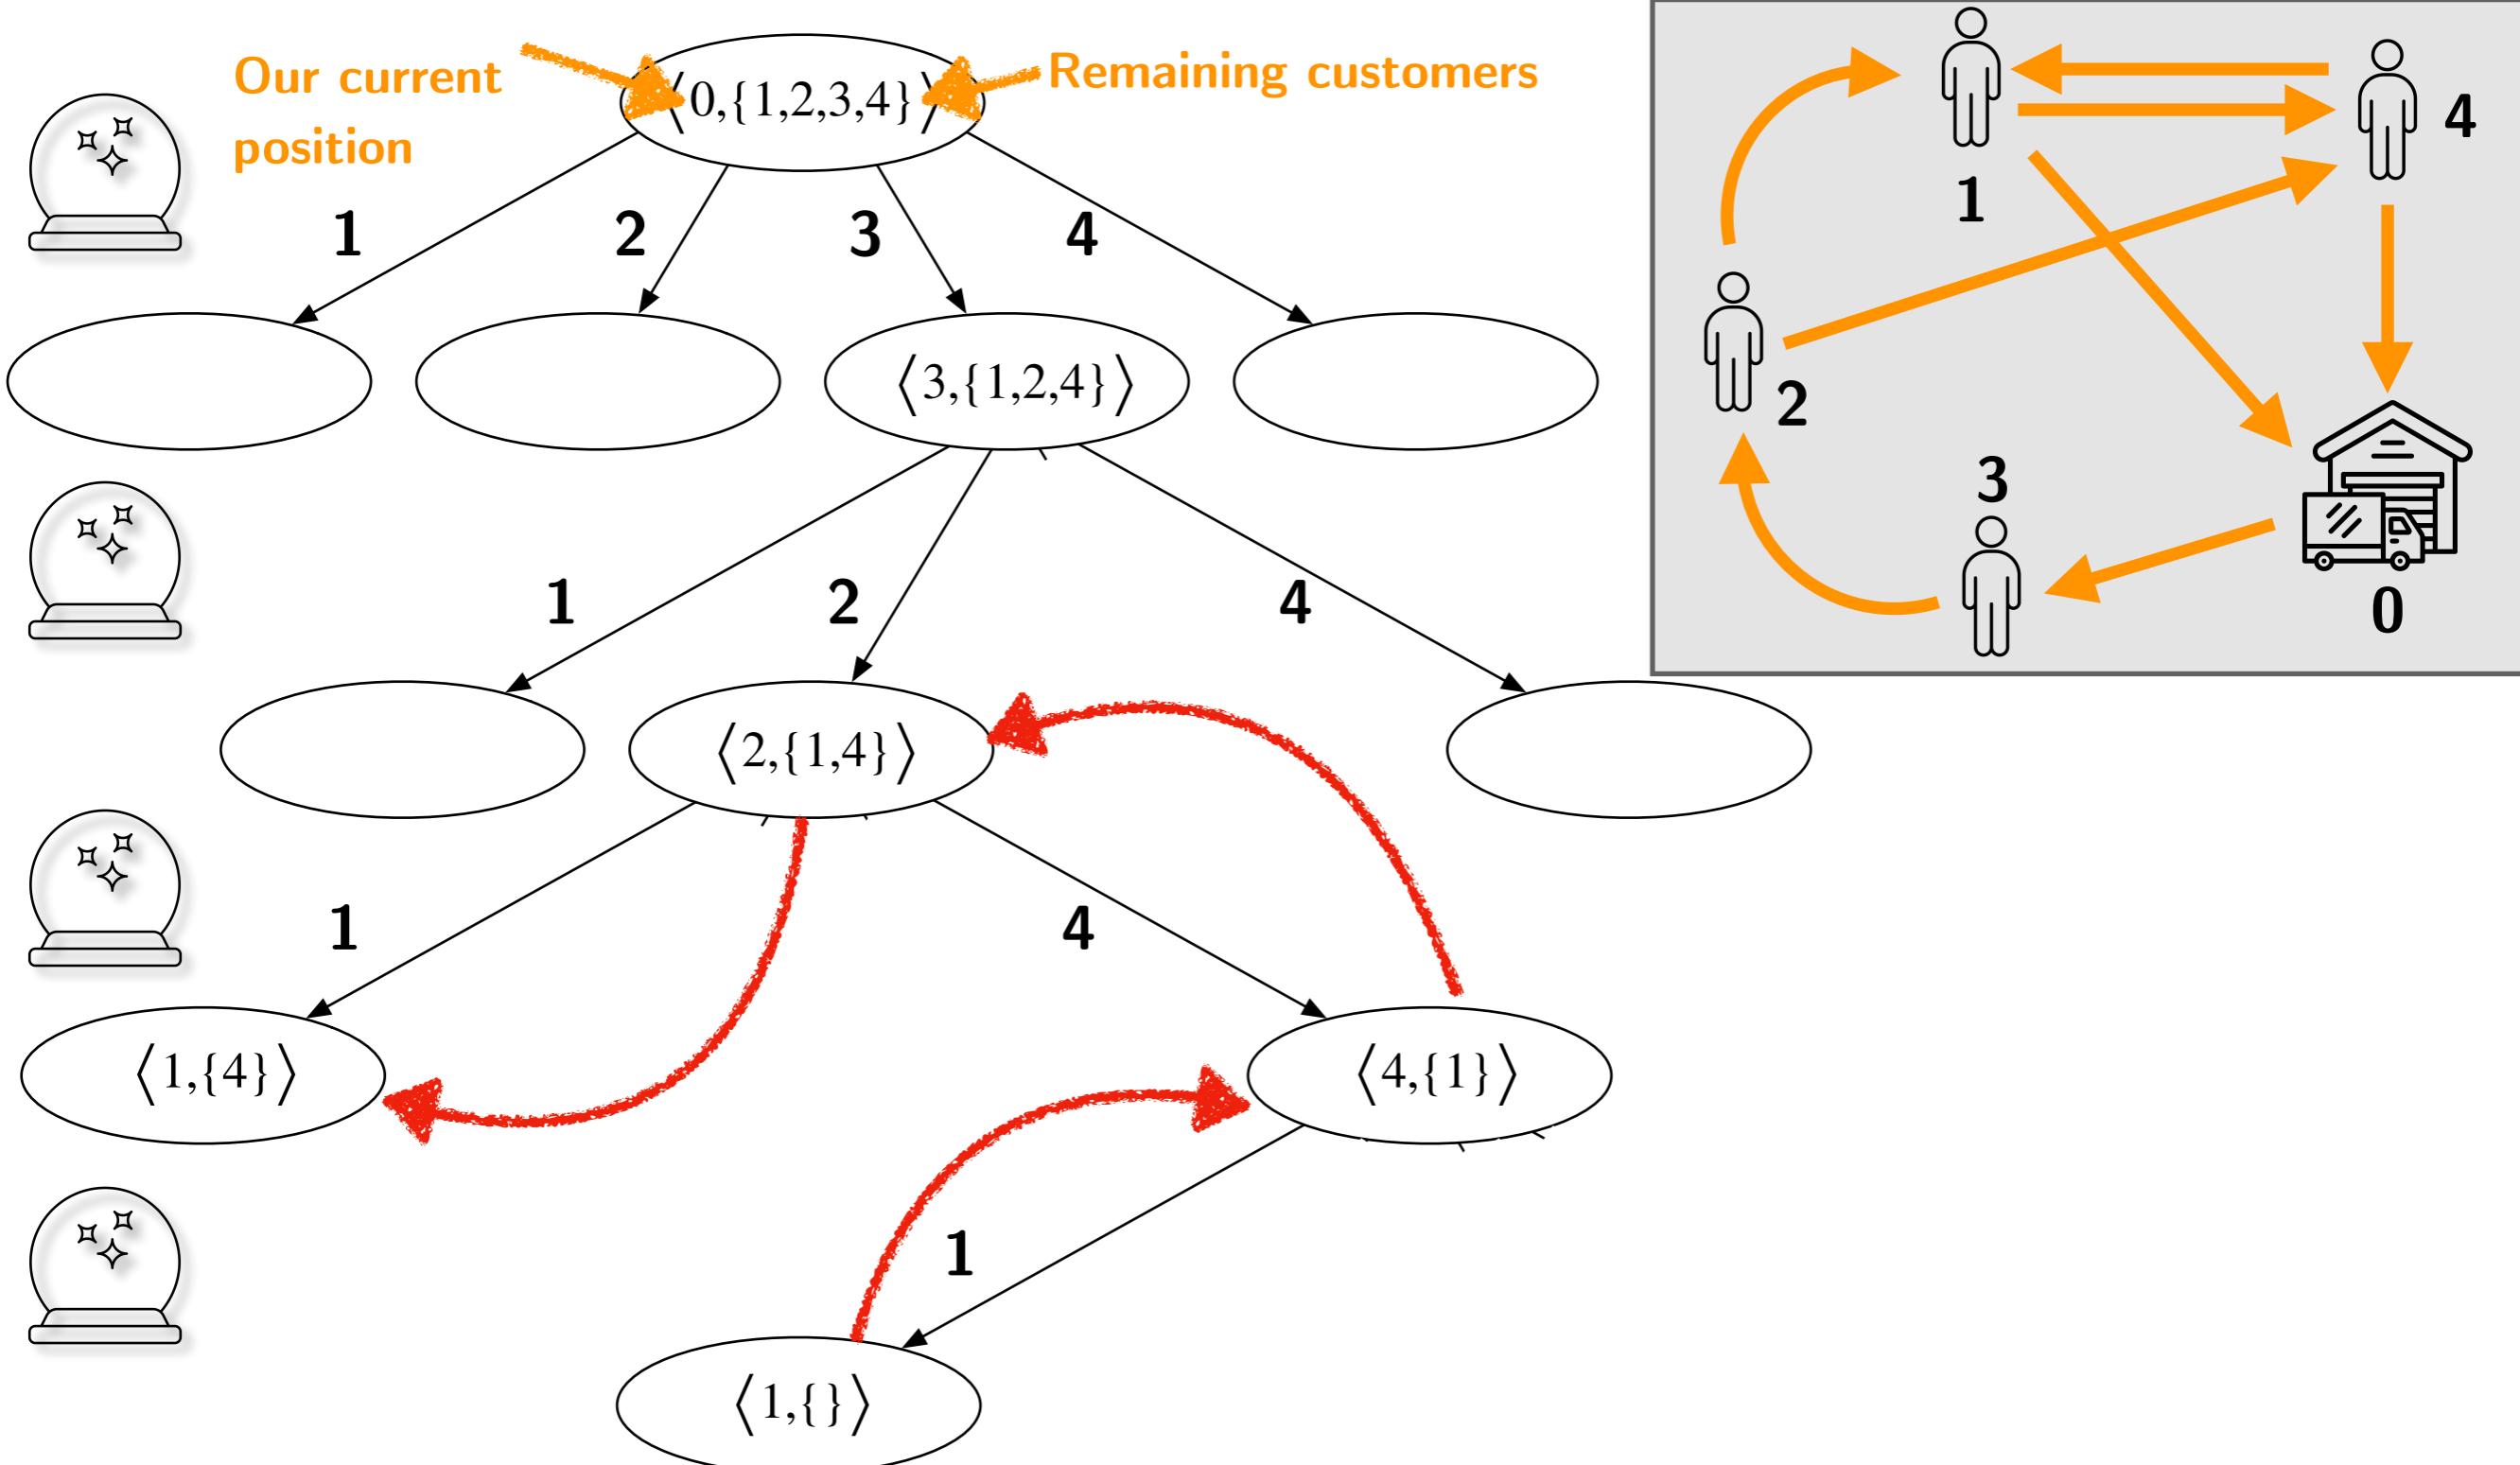
My first intuition: search is a good candidate as it heavily relies on **imperfect heuristics**

Consolidating argument: thanks to backtrack we can **recover from bad predictions**

Context: learning showed **promising results for MIP** (Gasse et al., Neurips 2019)

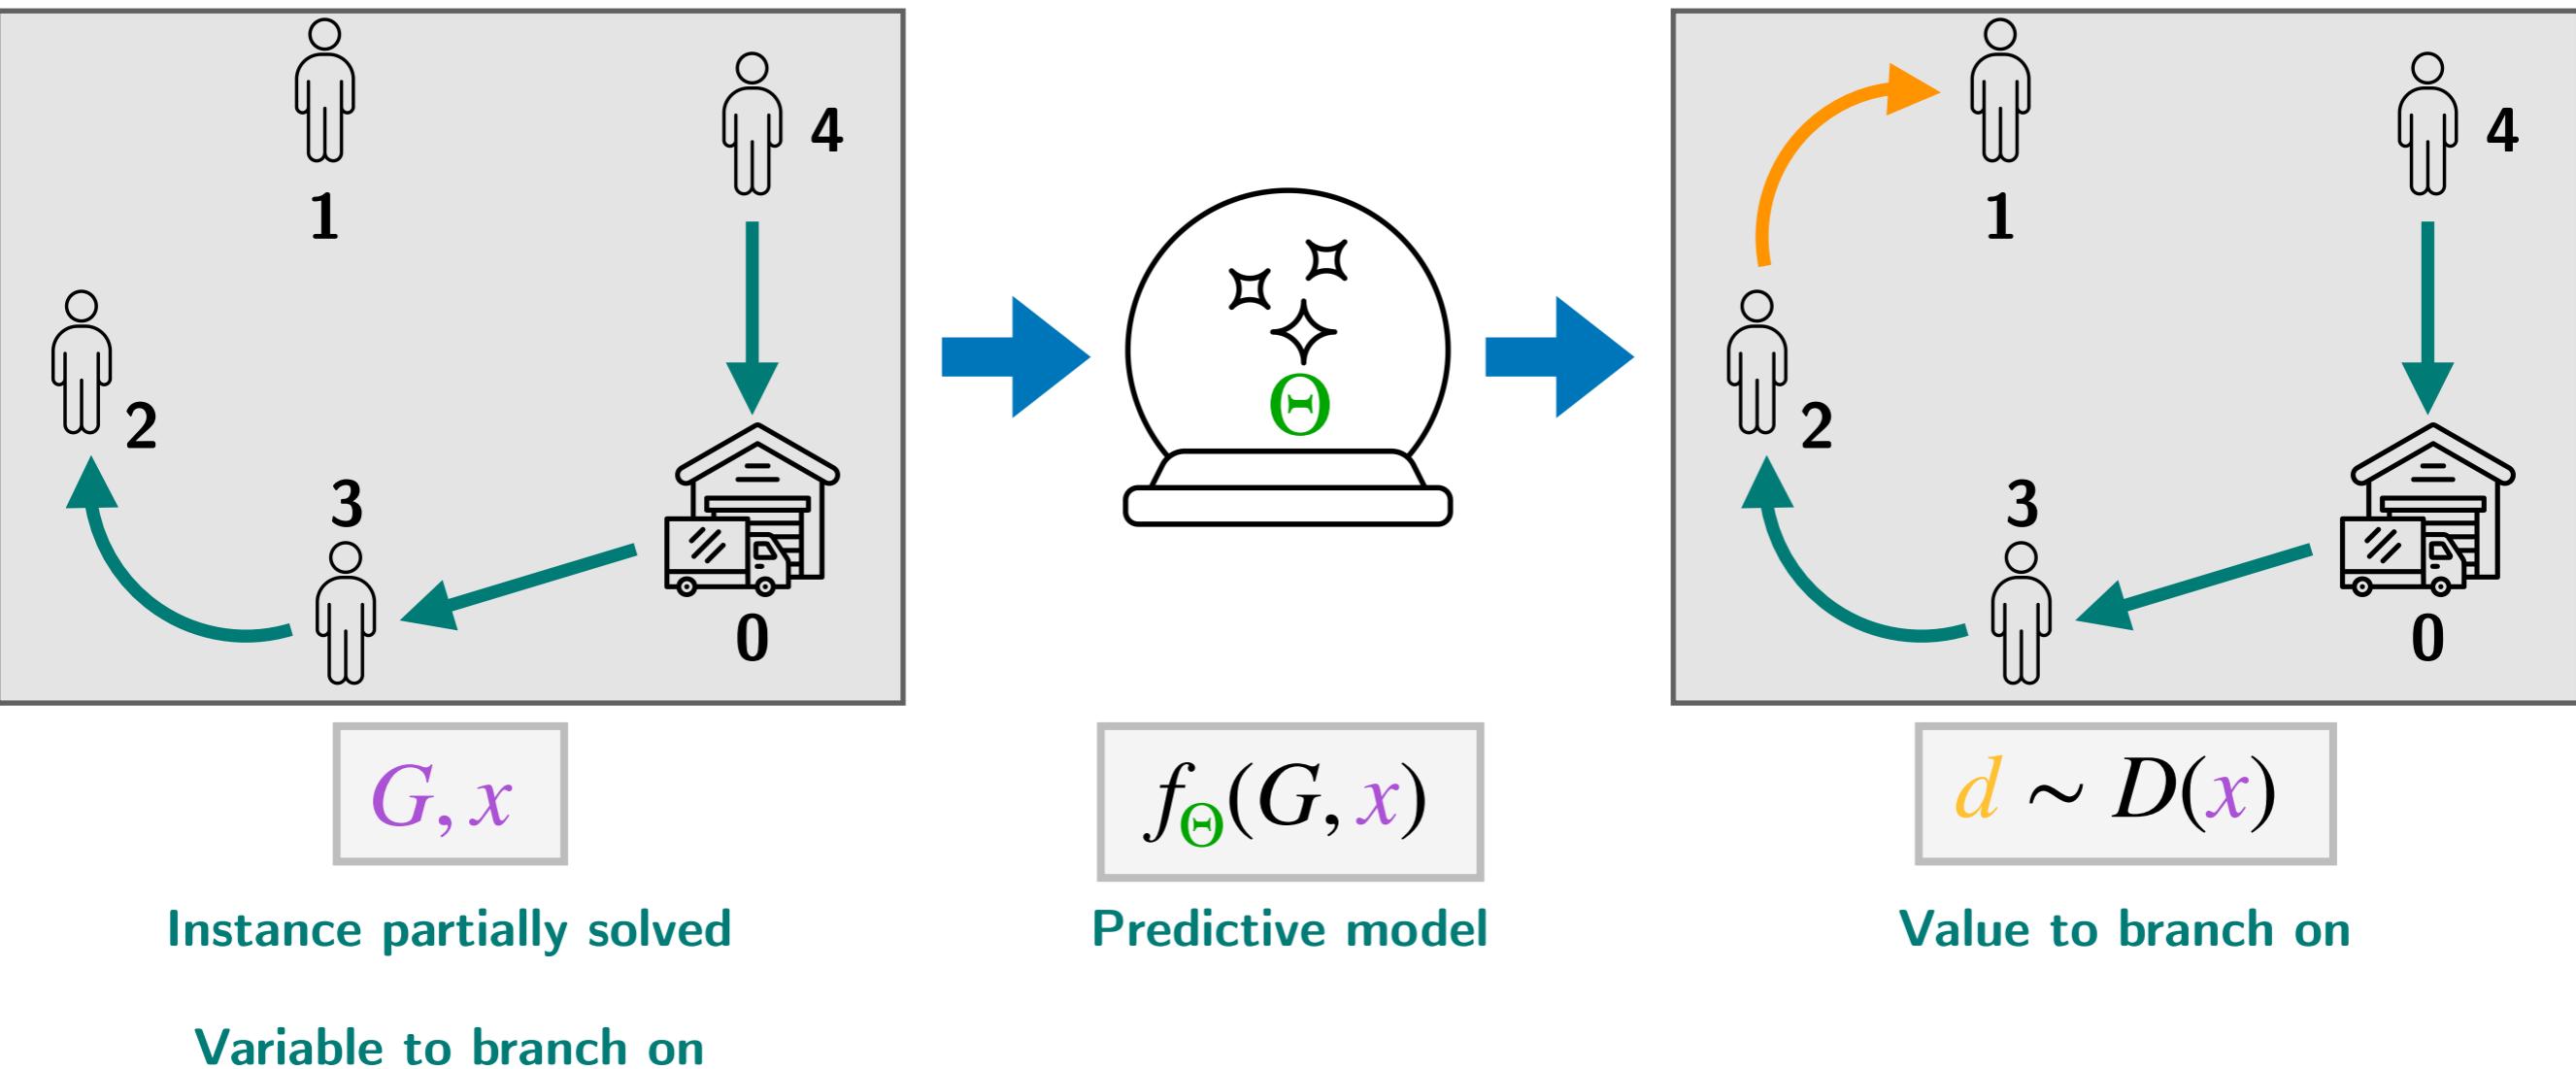
*My first research direction was to learn appropriate value-selection heuristic
(with the intent to go further by learning variable-selection heuristics)*

CP search with a learned heuristic on TSP



The ability to recover from bad decisions is fundamental for complete approaches

CP search with a learned heuristic on TSP



?

How can we build this predictive model (this magical crystal ball) ?

Challenge 1: we do not know what is the best choice (i.e., we do not have labels)

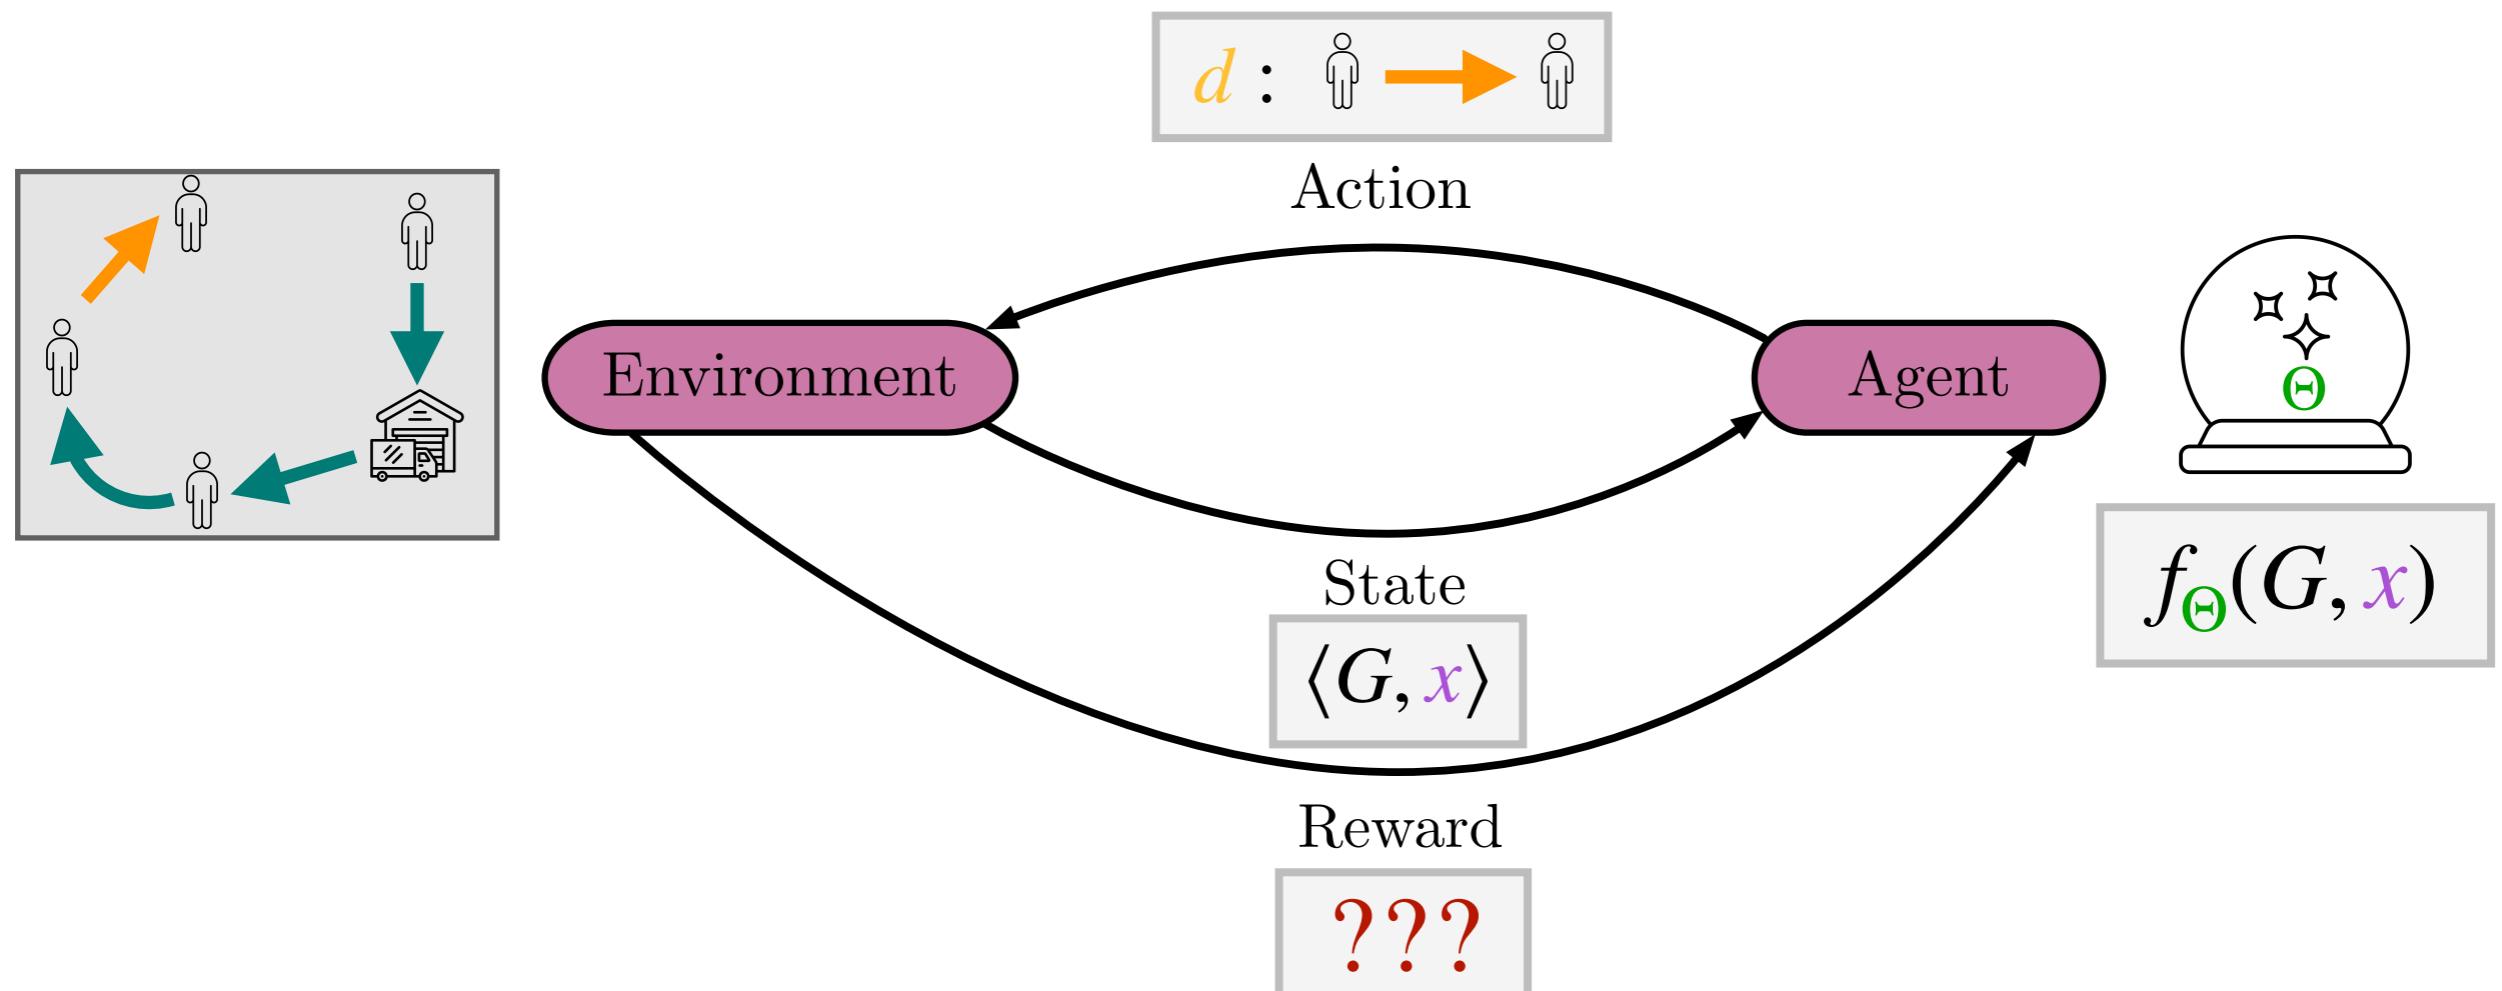
Challenge 2: difficult to represent a combinatorial problem in a proper way for learning

Tackling the first challenge (Cappart et al., AAAI 2021)

Challenge 1: we do not know what is the best choice

Consequence: tricky to rely on supervised learning

Our initial proposition: leverage reinforcement learning instead



Main idea: unveiling the connections of RL with CP through dynamic programming

Tricky part: it is not clear how to design the reward (what is a good assignment?)

Proposition of a reward function (Cappart et al., AAAI 2021)

? How to balance the fact that we want a **feasible** and **optimal** solution ?

Observation: **feasibility** is a prerequisite of **optimality**

Key idea 1: design a reward **always prioritizing finding feasible solutions**

Key idea 2: **discriminate then feasible solutions by their quality**

Reward for each action a at state s

$$= R(s, a) + \text{UB}$$

$\text{UB} > R(s, a)$

$R(s, a)$: direct increase in the objective function with assignment a

UB : strict upper bound on the optimal profit that can be reached

Intuition: **solutions deeper in the tree are always more rewarded than shallower ones**

Consequence 1: the first incentive is to go at the last depth (i.e., a feasible solution)

Consequence 2: **solutions at the same depth are rewarded by their objective value**

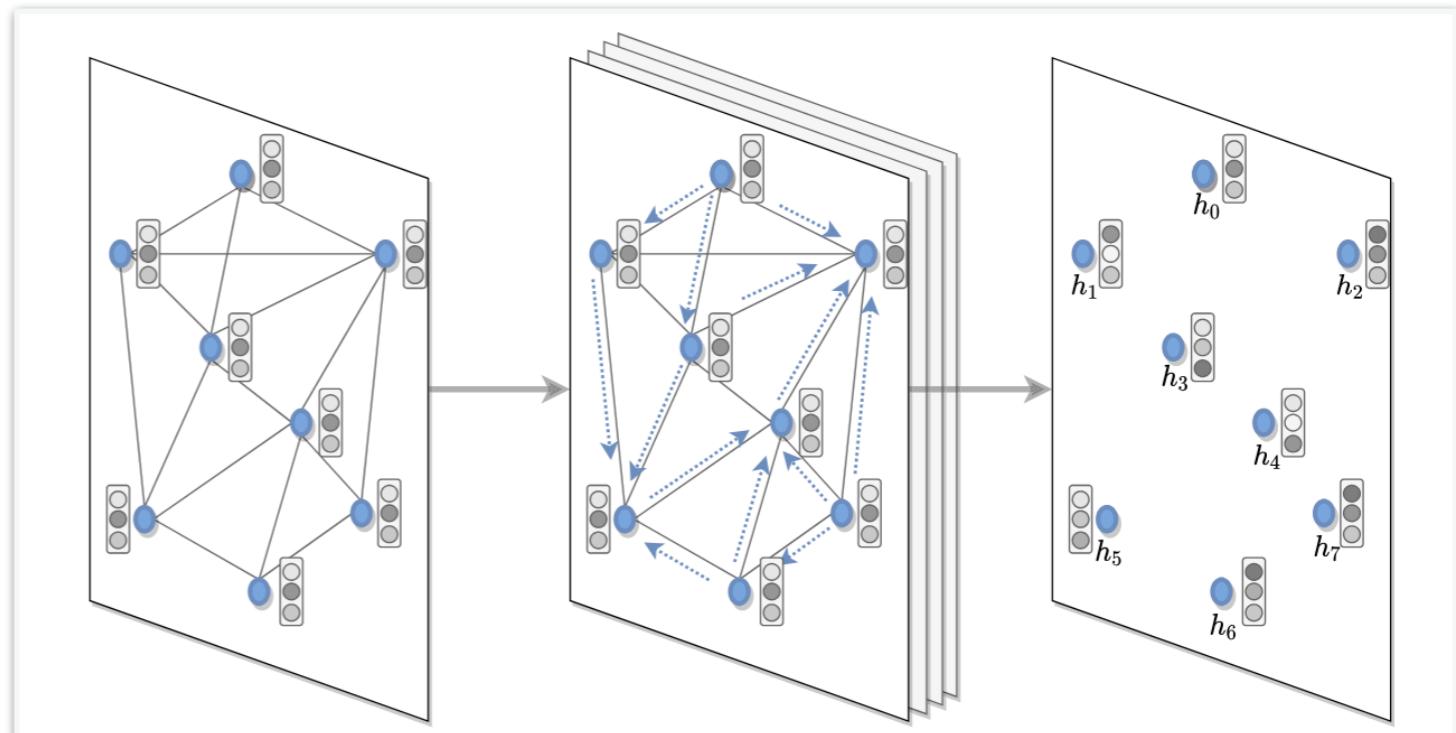
A first proof of concept (Cappart et al., AAAI 2021)

Algorithm 1: BaB-DQN Search Procedure.

```
> Pre:  $\mathcal{Q}_p$  is a COP having a DP formulation.  
> Pre:  $\mathbf{w}$  is a trained weight vector.  
 $\langle \mathbf{X}, \mathbf{D}, \mathbf{C}, \mathbf{o} \rangle := \text{CP Encoding}(\mathcal{Q}_p)$   
 $\mathcal{K} = \emptyset$   
 $\Psi := \text{BaB-search}(\langle \mathbf{X}, \mathbf{D}, \mathbf{C}, \mathbf{o} \rangle)$   
while  $\Psi$  is not completed do  
     $s := \text{encodeStateRL}(\Psi)$   
     $x := \text{takeFirstNonAssignedVar}(x)$   
    if  $s \in \mathcal{K}$  then  
         $v := \text{peek}(\mathcal{K}, s)$   
    else  
         $v := \text{argmax}_{u \in D(x)} \hat{Q}(s, u, \mathbf{w})$   
    end  
     $\mathcal{K} := \mathcal{K} \cup \{ \langle s, v \rangle \}$   
     $\text{branchAndUpdate}(\Psi, x, v)$   
end  
return bestSolution( $\Psi$ )
```



Challenge 2: how to represent a COP for learning?



Representing a problem as a **graph**, and feeding it to a **graph neural network**

Experiments: results showed that **relevant branching decisions can be learned**

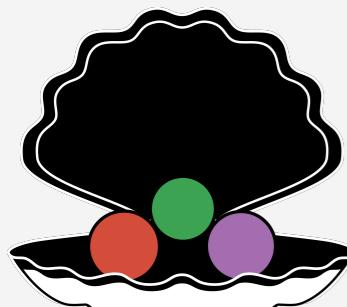
Limitation: far below state-of-the-art results and limited to relatively small instances

Technical difficulty: not-so-efficient to embed learning into existing CP solvers

These first results showed the promise in this direction and drove my research

Improvements in learning to branch

(Chalumeau, Coulon et al., CPAIOR 2021)



SeaPearl.jl

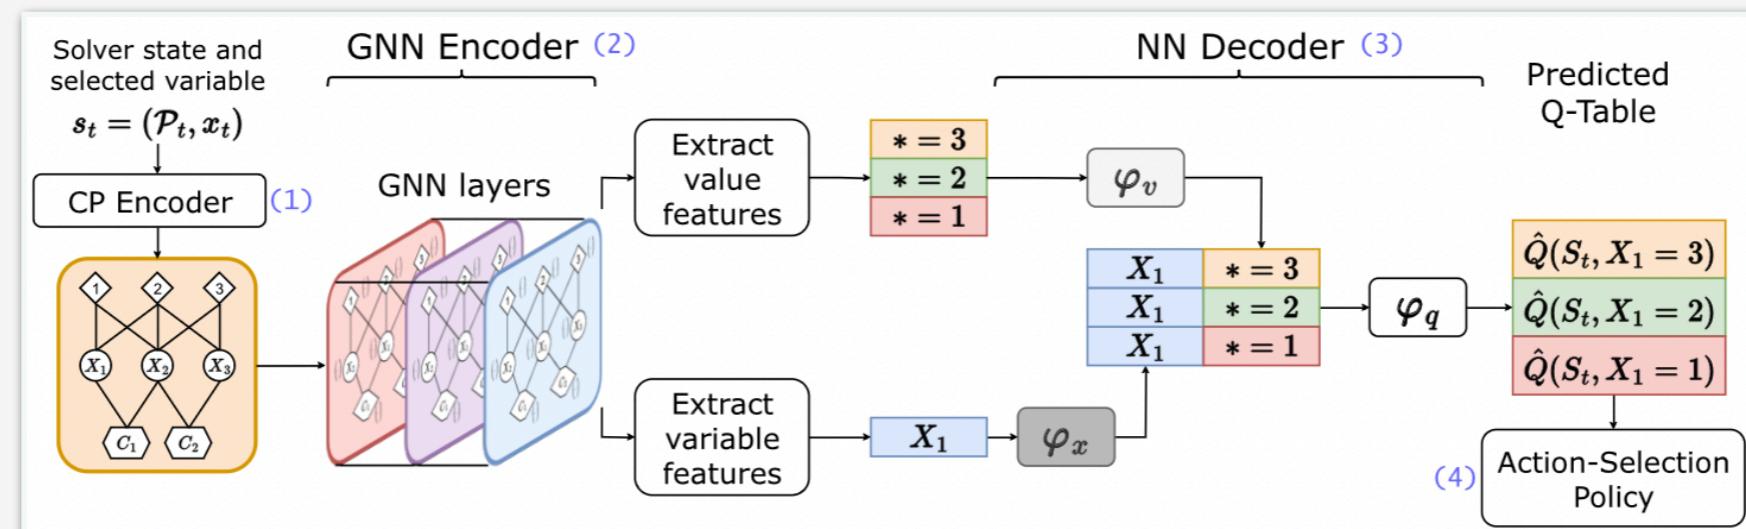
Seapearl: minimalist CP solver aiming to ease the integration of learning

Improvement: carrying out the learning inside the solver and not outside

Limitation 1: huge overhead in calling a GNN at each search node

Limitation 2: challenging to train (large resources, instability, etc.)

(Marty et al., CP 2023)



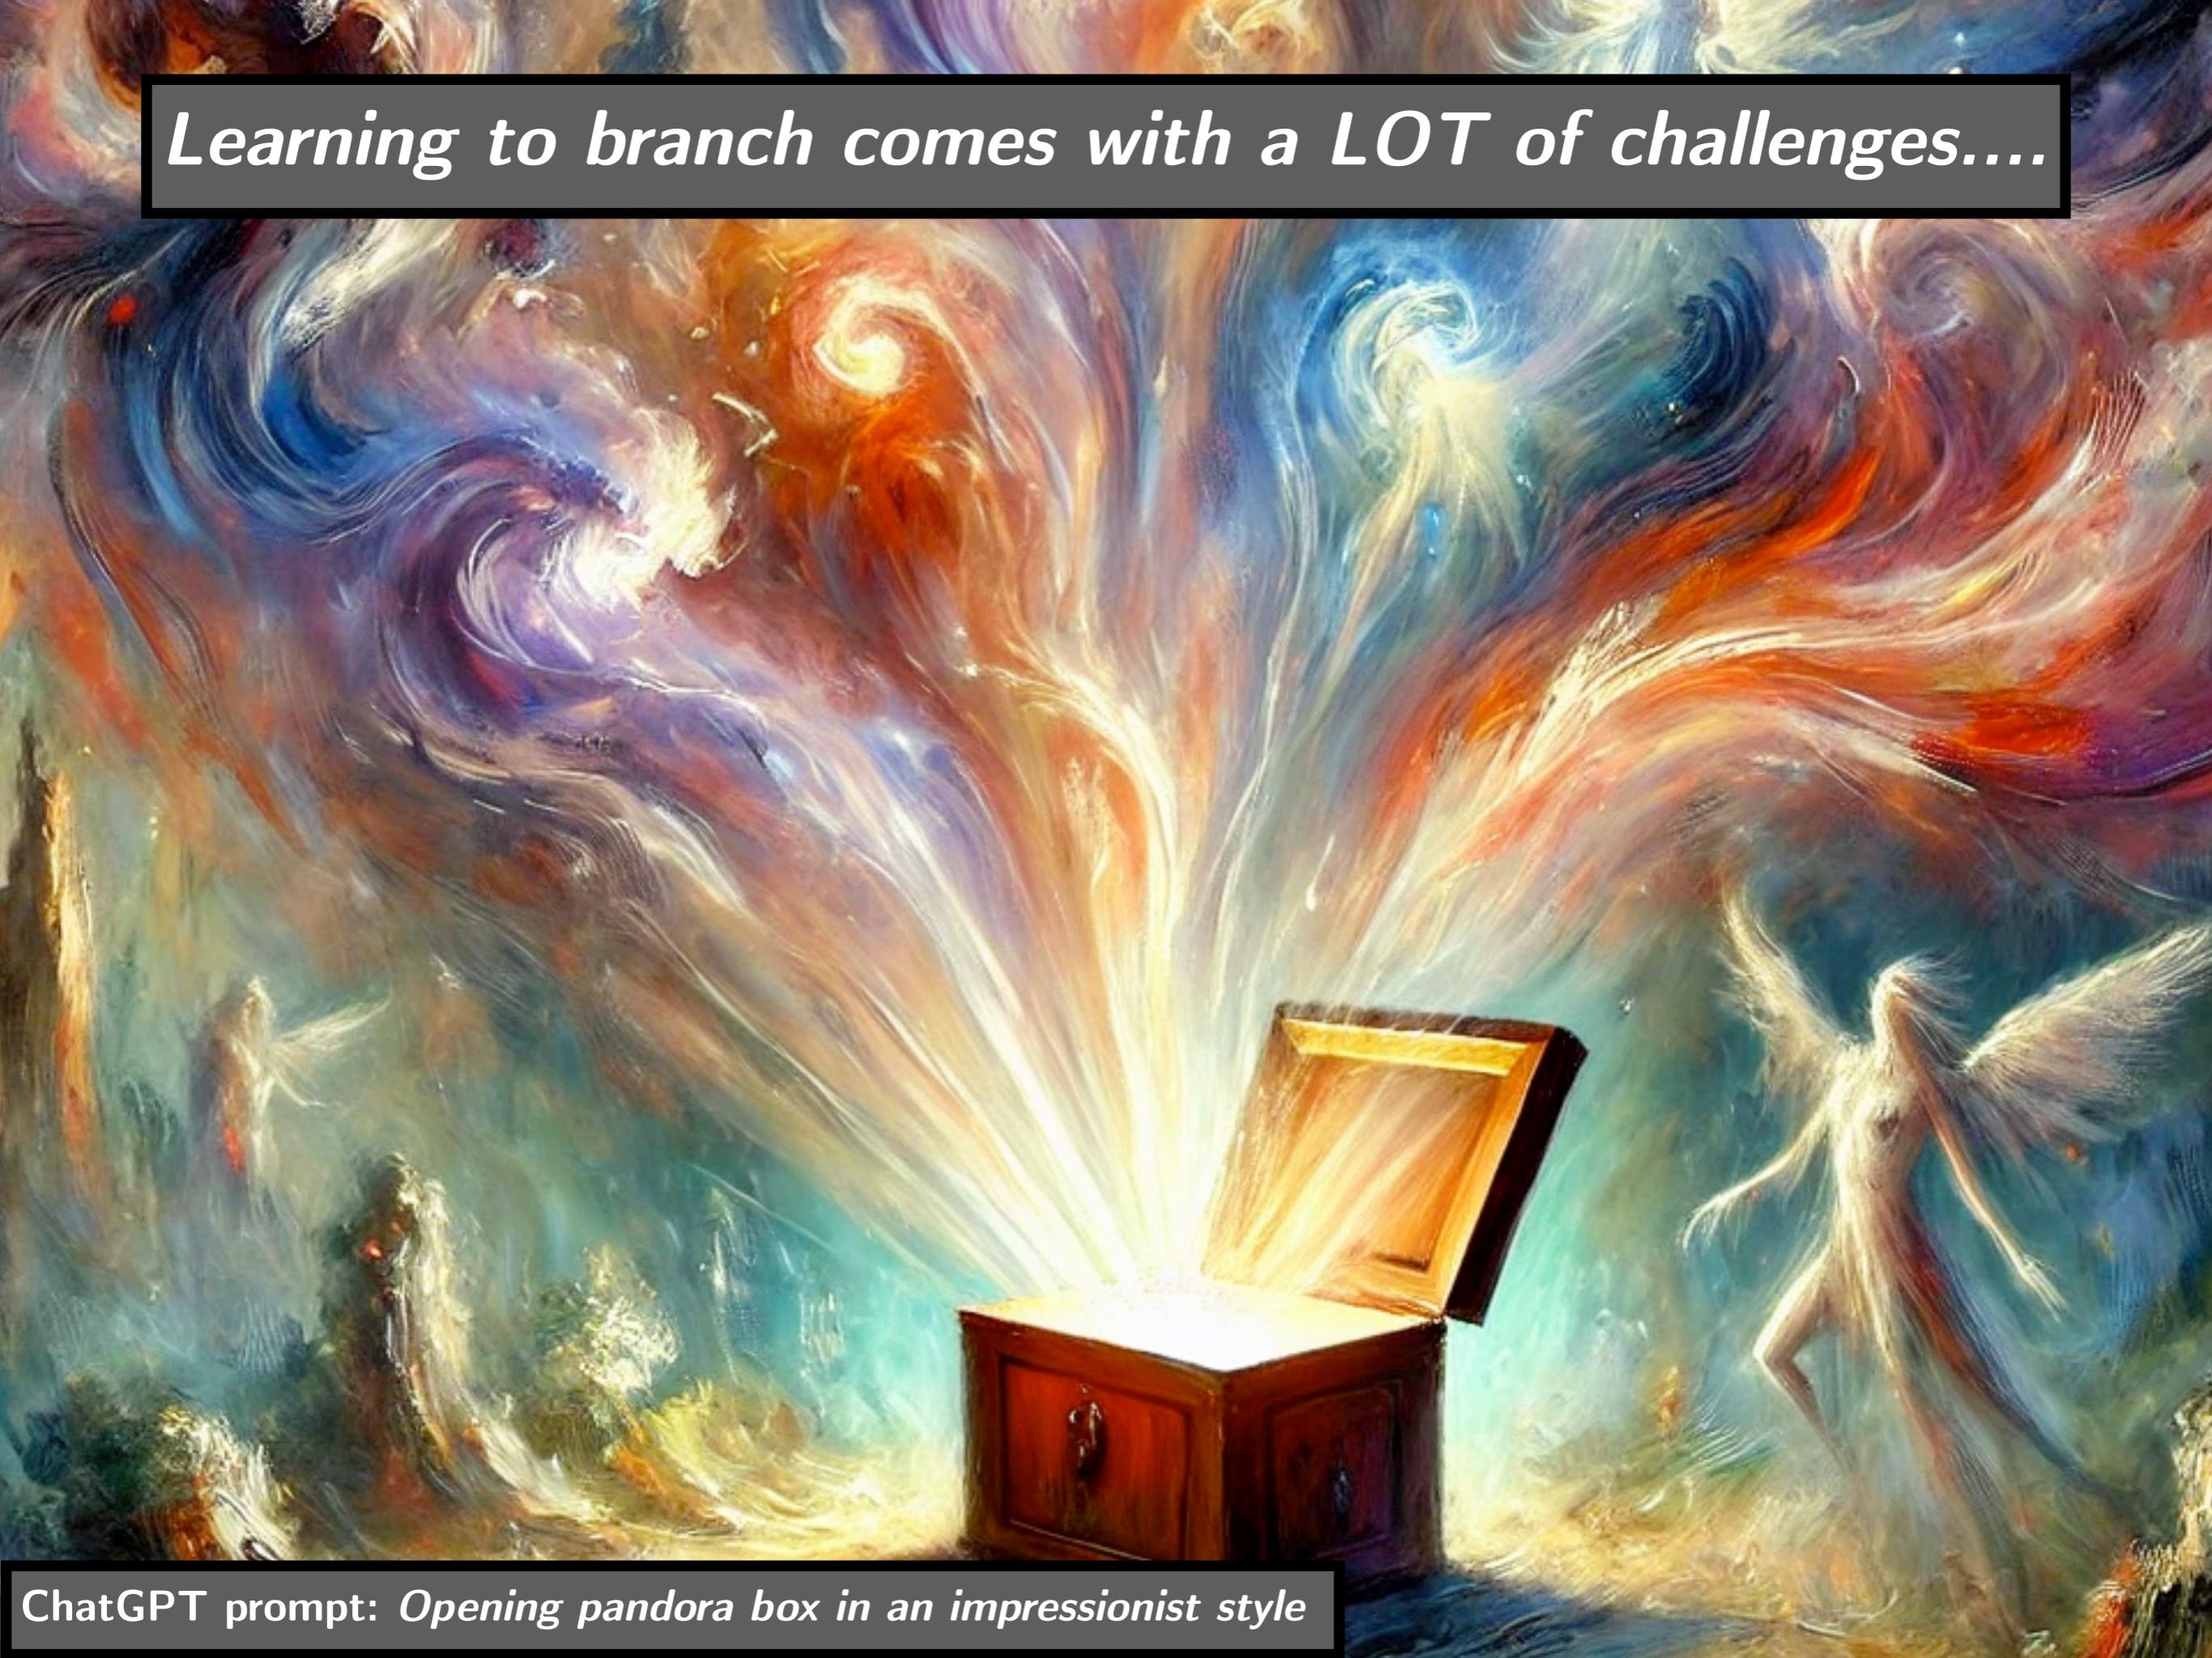
Main idea: mainly simplifying the framework!

Improvement: forget optimality proof (too difficult with value-selection heuristics)

New goal: finding quickly good solutions (redefining the reward function)

Limitation: still a huge overhead in calling often the GNN

Learning to branch comes with a LOT of challenges....



ChatGPT prompt: *Opening pandora box in an impressionist style*

My current thoughts in learning to branch

Learning branching heuristics comes with a LOT of challenges...

- (1) Overhead in calling a heavy model at each node (compared to cheap heuristics)
- (2) Subjective choices in designing the reward (feasibility, quality, optimality)
- (3) A model good at the root node may be poor at deeper levels (*distributional shift*)

Here are few advices based on what worked the best for me

- (1) Do not use learning at each branching step (too costly compared to heuristics)
- (2) Call the model only at the top of the tree (where you likely have more samples)
- (3) Proving optimality may be out-of-range (how to properly reward this?)
- (4) Focus on finding quickly good solutions (in few nodes)
- (5) Prioritize hybrid approaches (learning at the beginning, then use heuristics)

Going further with my thoughts...

Reinforcement learning may not be the best method for improving CP

Argument 1: much harder to train than supervised learning

Argument 2: indirectly require an unknown label (reward is an approximation of it)

Learning a value-selection heuristic may not be the best thing to do

Observation 1: learning is **too costly** and **unstable** to replace cheap heuristics

Observation 2: learning something does not always give practical improvements

Note: please only see this as **my personal opinion**, and not as an irrevocable truth :-)

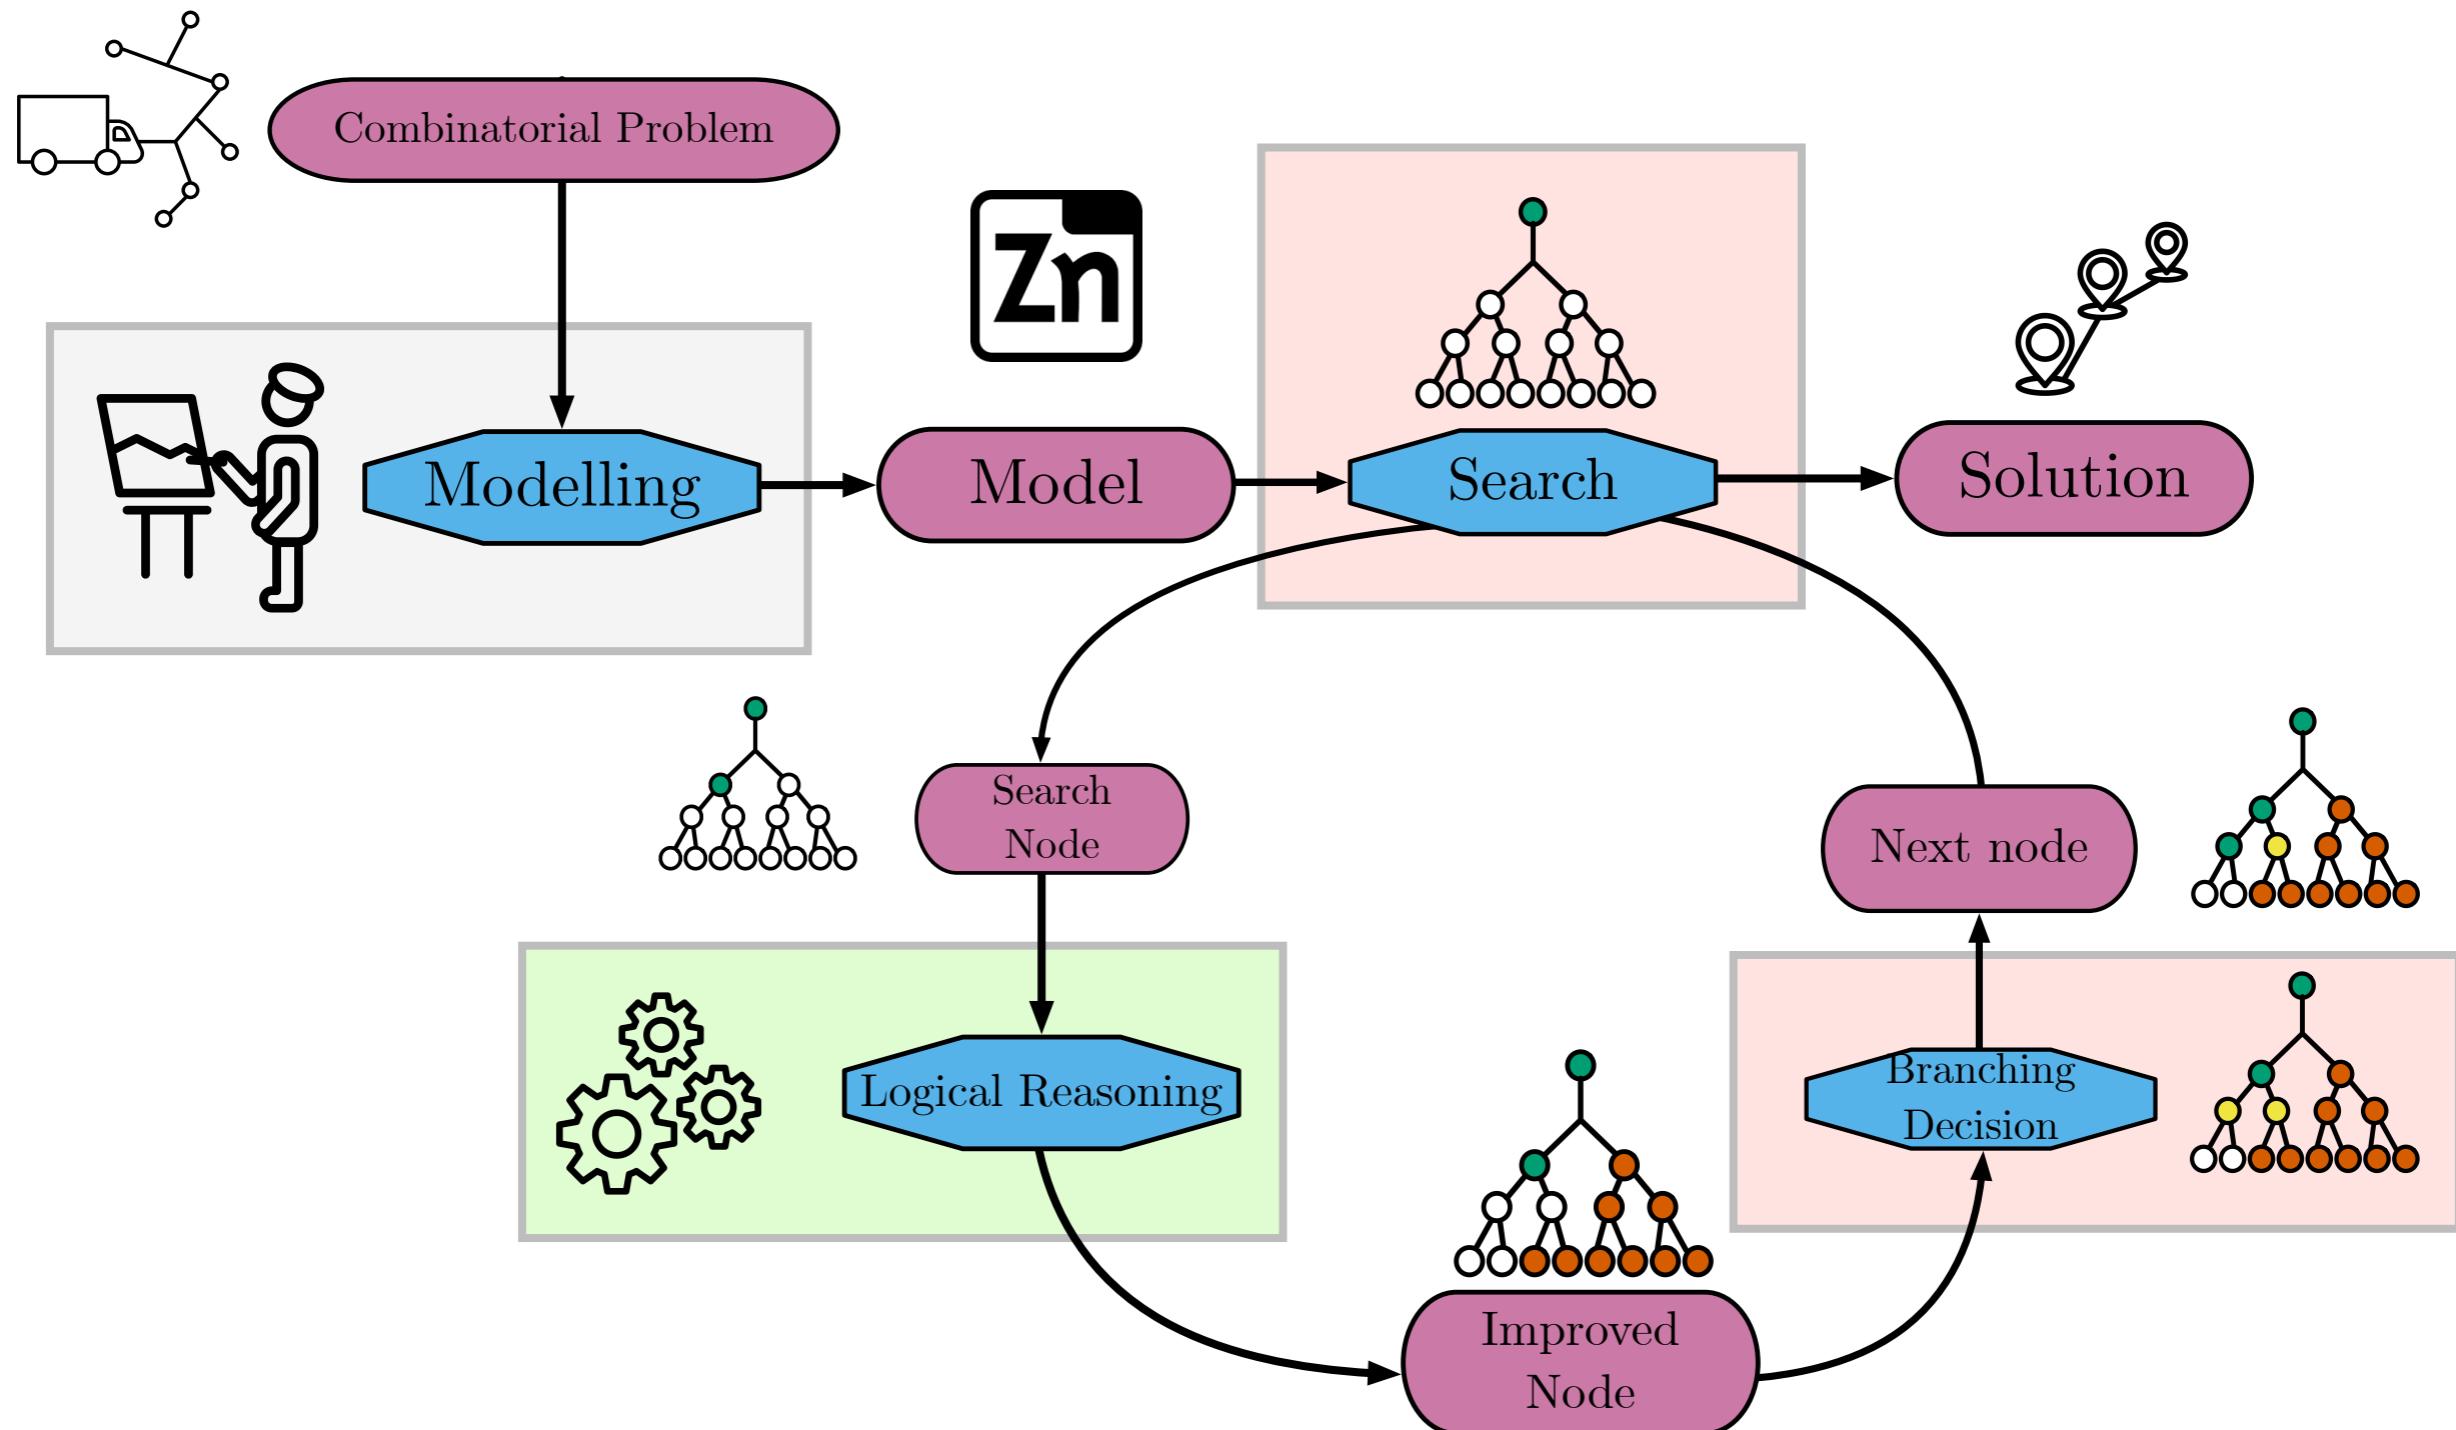


What about learning a variable-selection heuristic for CP ?

My intuition: we will probably have similar challenges than for the value-selection

Ok, but what do you propose then ?

Can we learn something else?



*My main current research direction is to learn how to prune the search space
(Said differently, I plan to improve the quality of filtering)*

Types of propagation in a CP solver



But how to do that ? Propagation is mainly algorithmic

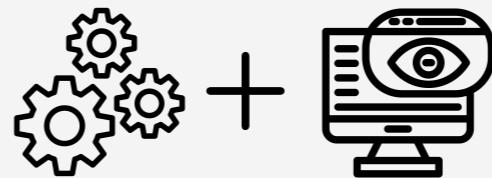
$CP = \text{model} + \text{propagation} + \text{search}$



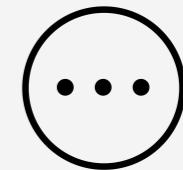
Fix-point



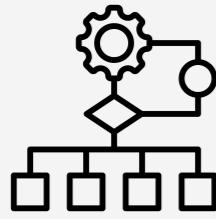
Consistency
(AC3, etc.)



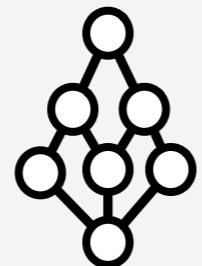
Global
constraints



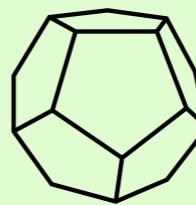
Other tools
(nogoods, etc.)



Pure
algorithmic



MDD



Cost-based
filtering

Some propagators relies on tricky-to-get information that we propose to learn

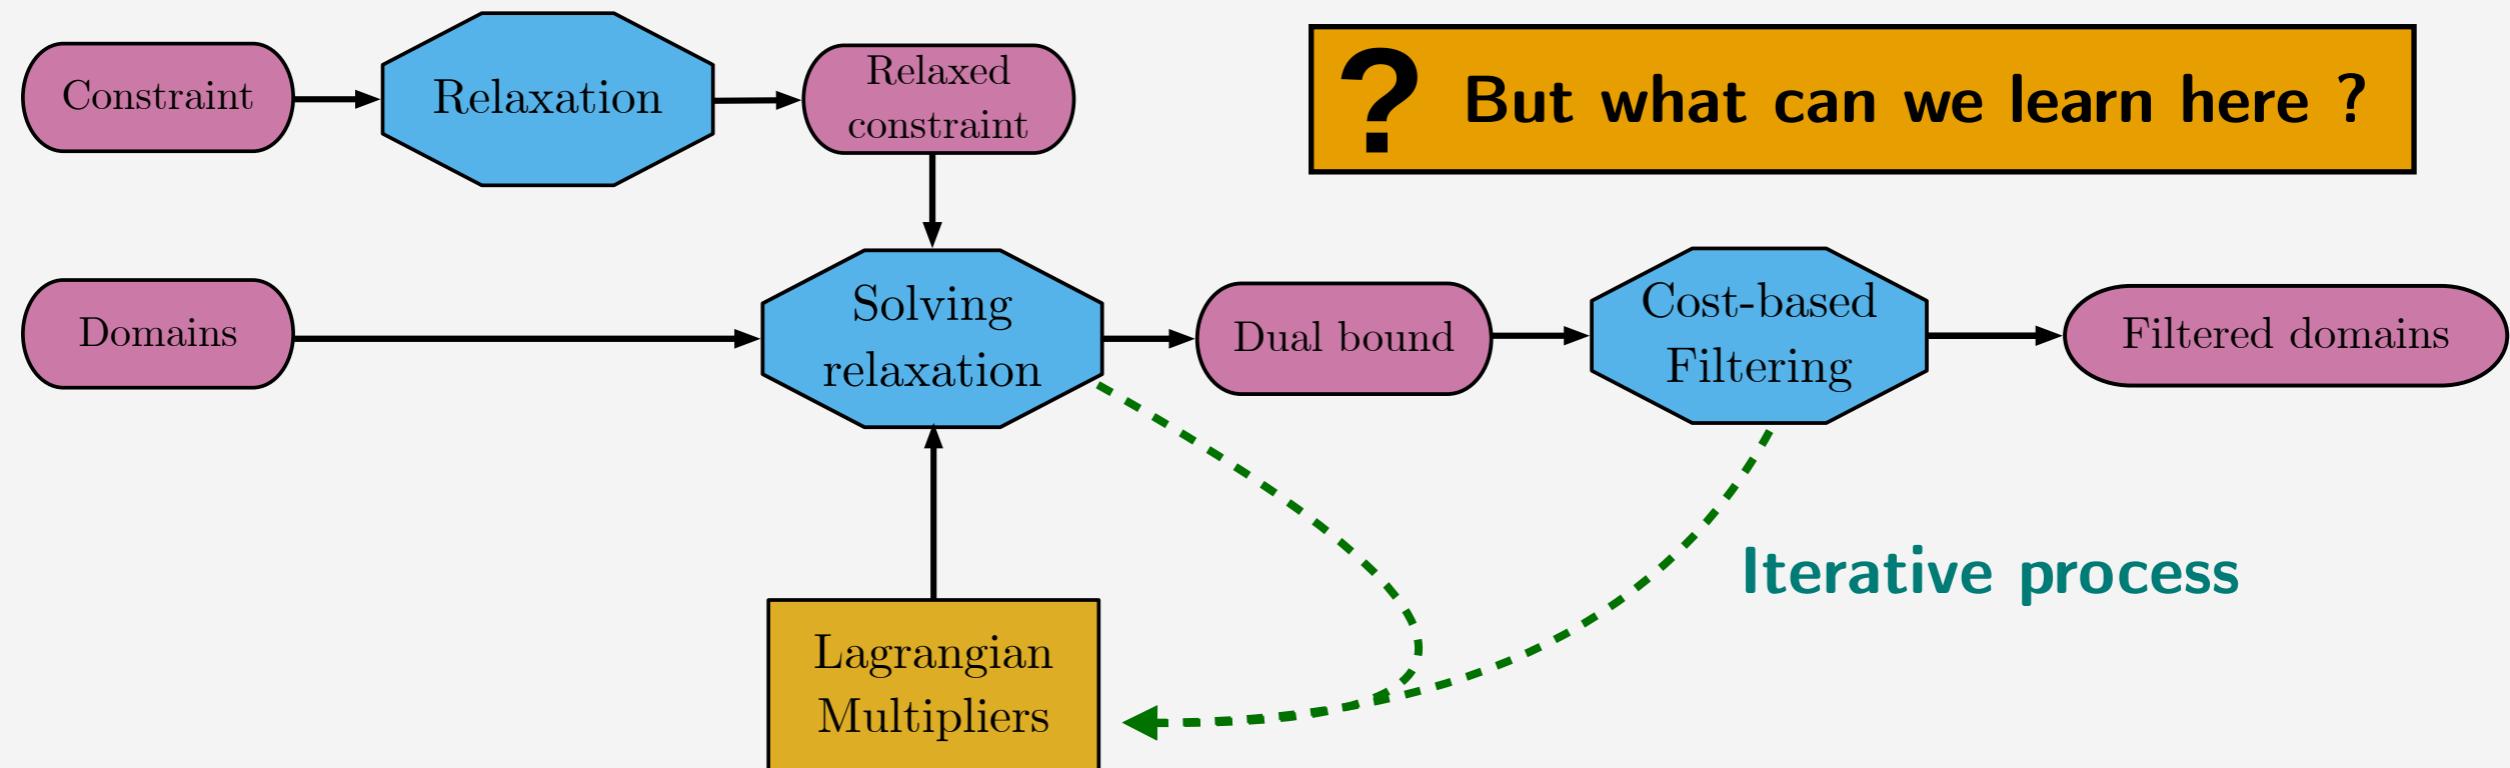
Cost-based filtering (Focacci et al., CP 1999)

Cost-based filtering leverages relaxation to improve the filtering of a constraint

Step 1: a valid relaxation is embedded into the global constraint

Step 2: the relaxed problem is solved to get a dual bound on the cost

Step 3: the bound is used to prune the search space



Trick: the quality of the relaxation depends on Lagrangian multipliers

Bad news: determining the best values of the multipliers is generally costly

We propose to obtain the multipliers through self-supervised learning

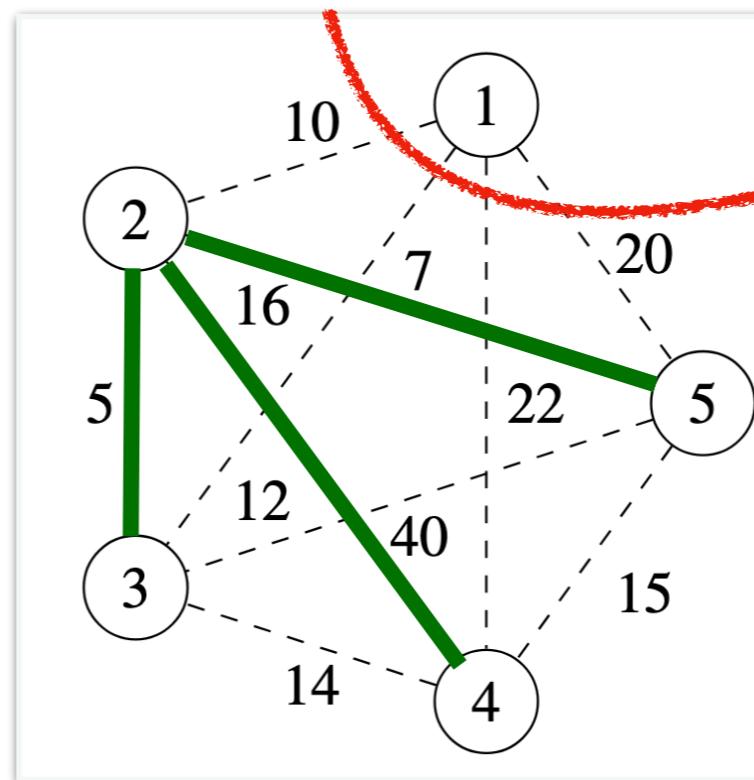
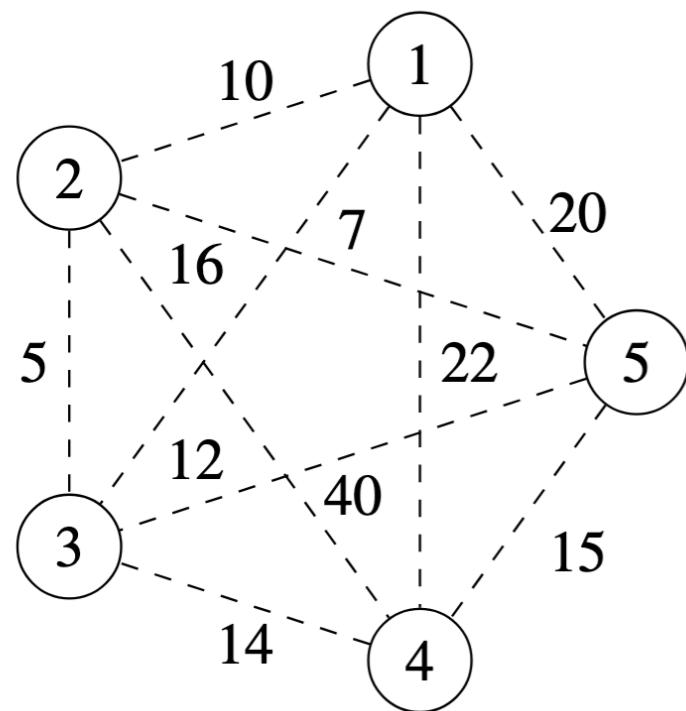
Learning Lagrangian Multipliers for the TSP

WeightedCircuit(X, G, d) : ensure that variables X form a TSP tour in G of cost $\leq d$

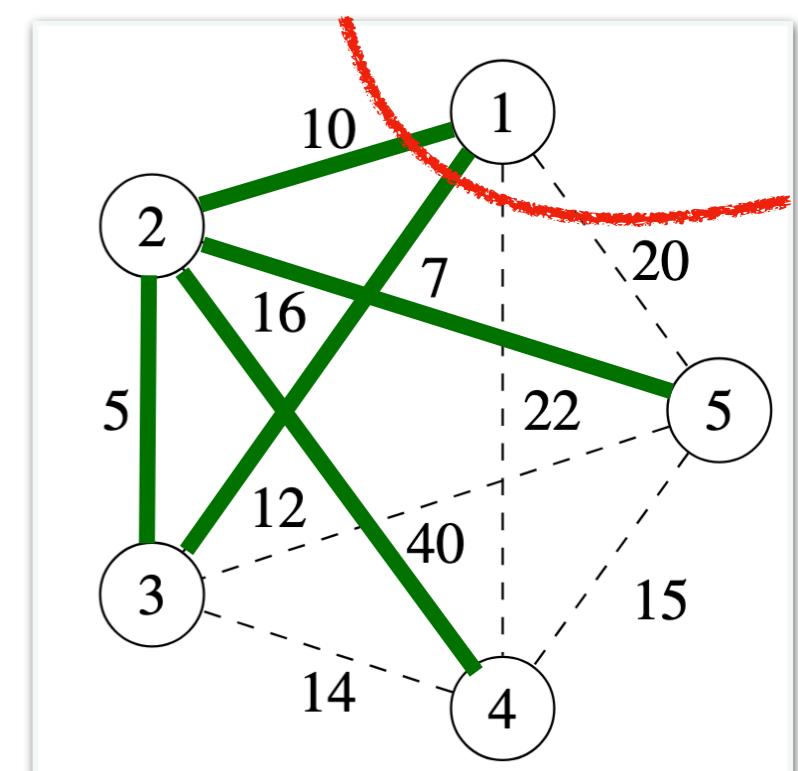
1-tree relaxation(X, G) : allows the TSP to go twice in a specific node

Step 1: ensure that variables X form a minimum spanning tree in $G \setminus \{v_1\}$

Step 2: link the remaining node v_1 to the tree with the two cheapest edges



Optimal TSP cost: 62



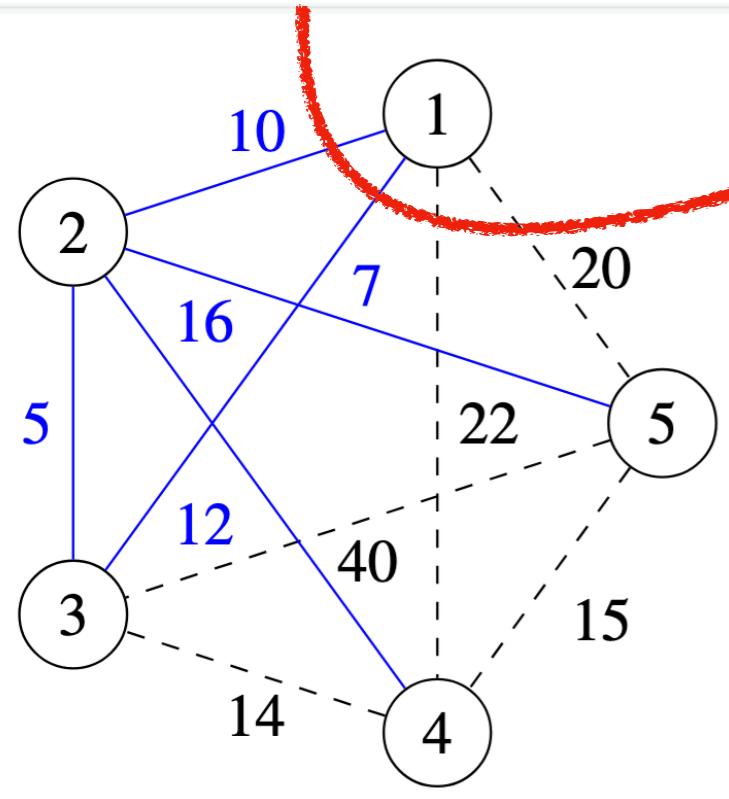
Relaxation cost: 50

We have here a example of how a specific constraint can be relaxed

Learning Lagrangian Multipliers for the TSP

?

The bound is not very tight, could we improve it ?



Lagrangian
Multipliers

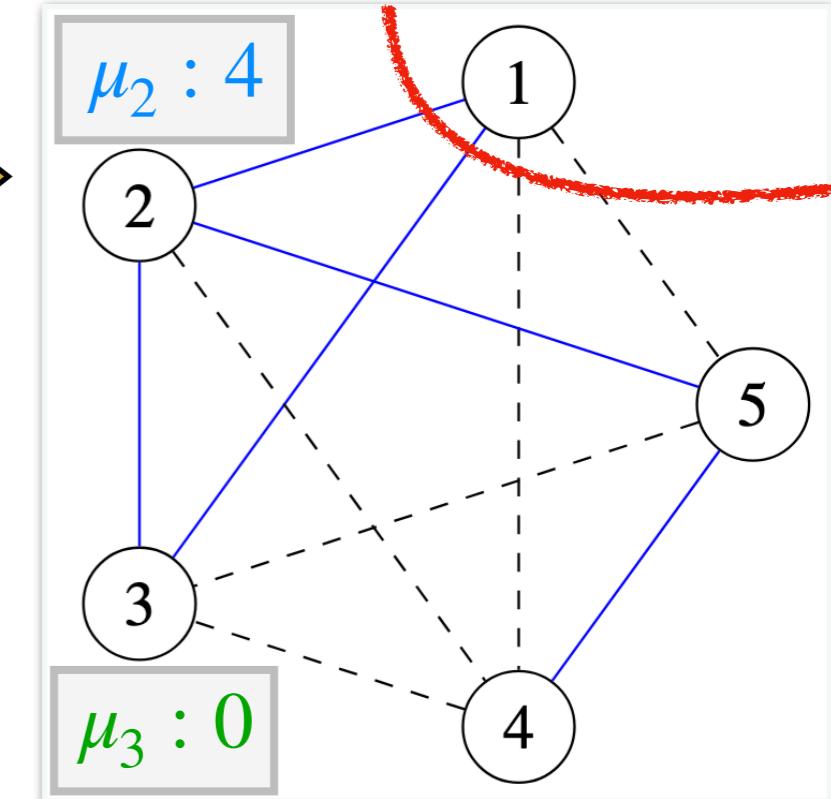
$$\langle \mu_1, \mu_2, \mu_3, \mu_4, \mu_5 \rangle$$

$$c'_{u,v} = c_{u,v} + \mu_u + \mu_v$$

$$c'_{2,3} = c_{2,3} + \mu_2 + \mu_3$$

$$9 = 5 + 4 + 0$$

Relaxation cost: 50



Relaxation cost: 59

Main idea: perturbate the cost of each edge with values called **Lagrangian multipliers**

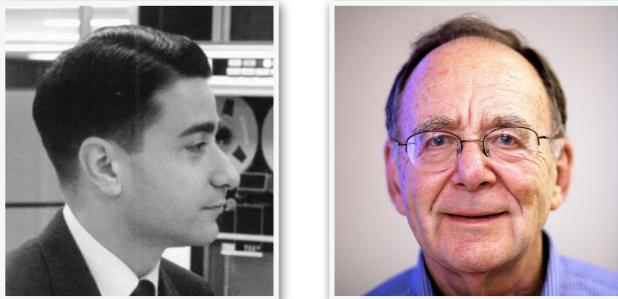
Lagrangian multiplier: value associated with each node

Perturbation: change the cost of each edge based on the multipliers of its nodes

Nice property 1: the optimal TSP tour is invariant under this perturbation

Nice property 2: better bounds generally result in a much better filtering

Learning Lagrangian Multipliers for the TSP



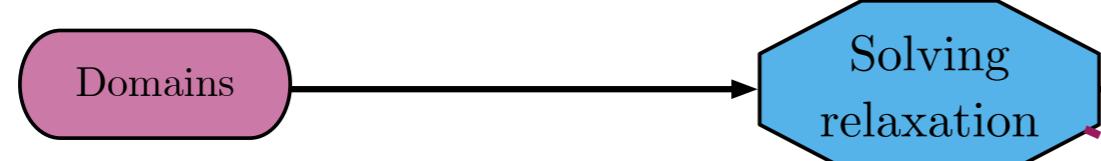
?

But how can we find good values for the multipliers ?

Held and Karp, 1970: proposed an iterative algorithm for that

WeightedCircuit

1-tree relaxation



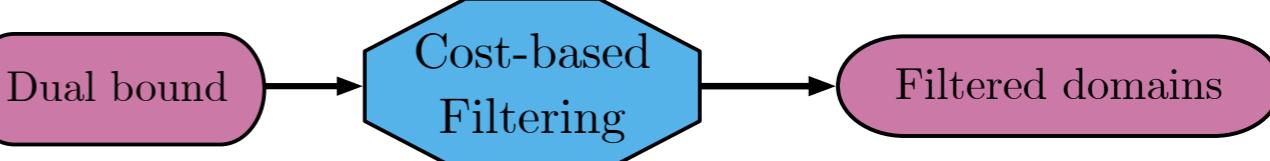
$$f_{\Theta}(G, x) : \langle \mu_1, \dots, \mu_n \rangle$$

Lagrangian
Multipliers

Constraints (2012) 17:205–233
DOI 10.1007/s10601-012-9119-x

Improved filtering for weighted circuit constraints

Pascal Benchimol · Willem-Jan van Hoeve ·
Jean-Charles Régin · Louis-Martin Rousseau ·
Michel Rueher



Held-Karp iterative process

Bad news: the iterative adjustment of the multipliers is computationally expensive

Parjadis et al. (CP 2024): use learning to initialize the multipliers

Idea: use a GNN to predict multipliers (one per node) from a TSP instance

Results: better filtering achieved on random and symmetric TSPs

Observation: replacing HK process with learning was not successful



Types of propagation in a CP solver



Can we learn a more general bounding mechanism in CP ?

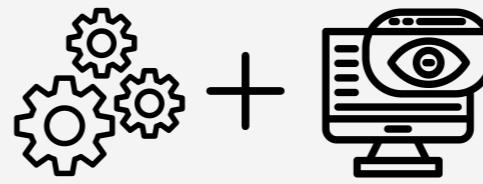
$CP = \text{model} + \text{propagation} + \text{search}$



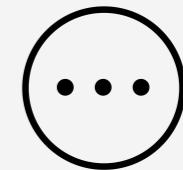
Fix-point



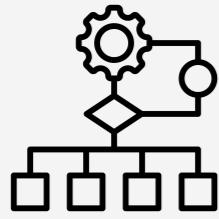
Consistency
(AC3, etc.)



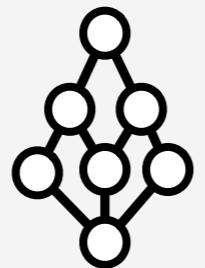
Global
constraints



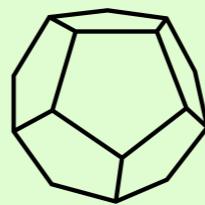
Other tools
(nogoods, etc.)



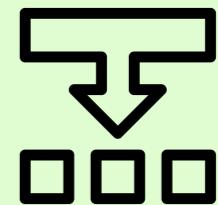
Pure
algorithmic



MDD



Cost-based
filtering



Lagrangian
decomposition

We extend the idea to learn multipliers to **CP-based Lagrangian decomposition**

Lagrangian decomposition (Guignard and Kim, 1987)

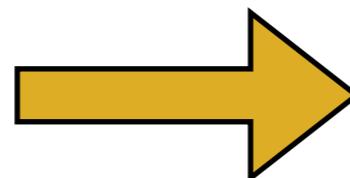
Lagrangian decomposition splits the problem into independent and easier subproblems

Step 1: each variable in each constraint is duplicated, except for the first constraint

Step 2: a constraint linking the values is added for each new variable

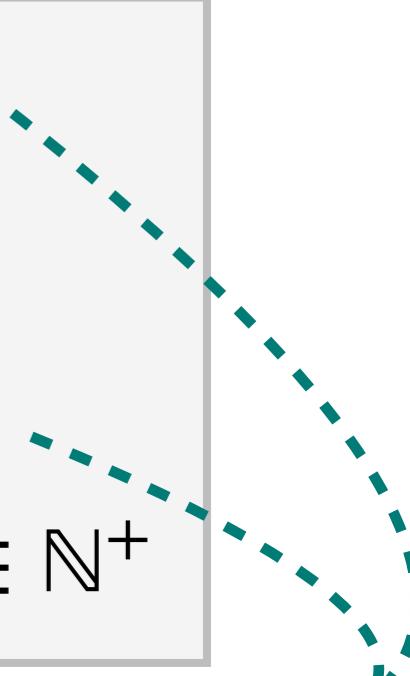
Step 3: these constraints are moved into the objective function with a penalty term

$$\begin{aligned} \max \quad & f(X_1, X_2, X_3) \\ \text{s.t.} \quad & C_1(X_1, X_2, X_3) \\ & C_2(X_2, X_3) \\ & X_1, X_2, X_3 \in \mathbb{N}^+ \end{aligned}$$



$$\begin{aligned} \max \quad & f(X_1, X_2, X_3) \\ \text{s.t.} \quad & C_1(X_1, X_2) \\ & C_2(Y_2, Y_3) \\ & Y_2 = X_2, Y_3 = X_3 \\ & X_1, X_2, Y_2, X_3, Y_3 \in \mathbb{N}^+ \end{aligned}$$

$$\max \quad f(X_1, X_2, X_3) + \boxed{\mu_2} \cdot (Y_2 - X_2) + \boxed{\mu_3} \cdot (Y_3 - X_3)$$



Again, we have Lagrangian multipliers, but in a more generic way than before

Lagrangian decomposition (Guignard and Kim, 1987)

$$\mathcal{B}(\mu_2, \mu_3) = \left\{ \begin{array}{l} \max f(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) + \mu_2 \cdot (\textcolor{violet}{X}_2 - \textcolor{red}{Y}_2) + \mu_3 \cdot (\textcolor{green}{X}_3 - \textcolor{blue}{Y}_3) \\ \text{s.t. } C_1(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) \\ \quad C_2(\textcolor{red}{Y}_2, \textcolor{blue}{Y}_3) \\ \quad X_1, \textcolor{violet}{X}_2, \textcolor{red}{Y}_2, \textcolor{green}{X}_3, \textcolor{blue}{Y}_3 \in \mathbb{N}^+ \end{array} \right.$$

Solving this relaxed problem will give a dual bound

? But, is it easy to solve ?

Observation: by construction, each constraint has its own set of variables

Consequence: each constraint can be solved independently

$$\mathcal{B}(\mu_2, \mu_3) = \left\{ \begin{array}{l} \max \left(f(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) + \mu_2 \cdot \textcolor{violet}{X}_2 + \mu_3 \cdot \textcolor{green}{X}_3 \right) \\ \text{s.t. } C_1(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) \\ \quad X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3 \in \mathbb{N}^+ \end{array} \right. + \left\{ \begin{array}{l} \max \left(-\mu_2 \cdot \textcolor{red}{Y}_2 - \mu_3 \cdot \textcolor{blue}{Y}_3 \right) \\ \text{s.t. } C_2(\textcolor{red}{Y}_2, \textcolor{blue}{Y}_3) \\ \quad \textcolor{red}{Y}_2, \textcolor{blue}{Y}_3 \in \mathbb{N}^+ \end{array} \right.$$

Given some multipliers, we can obtain a bound by solving several subproblems

Lagrangian decomposition in CP (Hà et al. CP 2015)

$$\mathcal{B}(\mu_2, \mu_3) = \begin{array}{l} \max \left(f(X_1, X_2, X_3) + \mu_2 \cdot X_2 + \mu_3 \cdot X_3 \right) \\ \text{s.t. } C_1(X_1, X_2, X_3) \\ X_1, X_2, X_3 \in \mathbb{N}^+ \end{array} + \begin{array}{l} \max \left(-\mu_2 \cdot Y_2 - \mu_3 \cdot Y_3 \right) \\ \text{s.t. } C_2(Y_2, Y_3) \\ Y_2, Y_3 \in \mathbb{N}^+ \end{array}$$



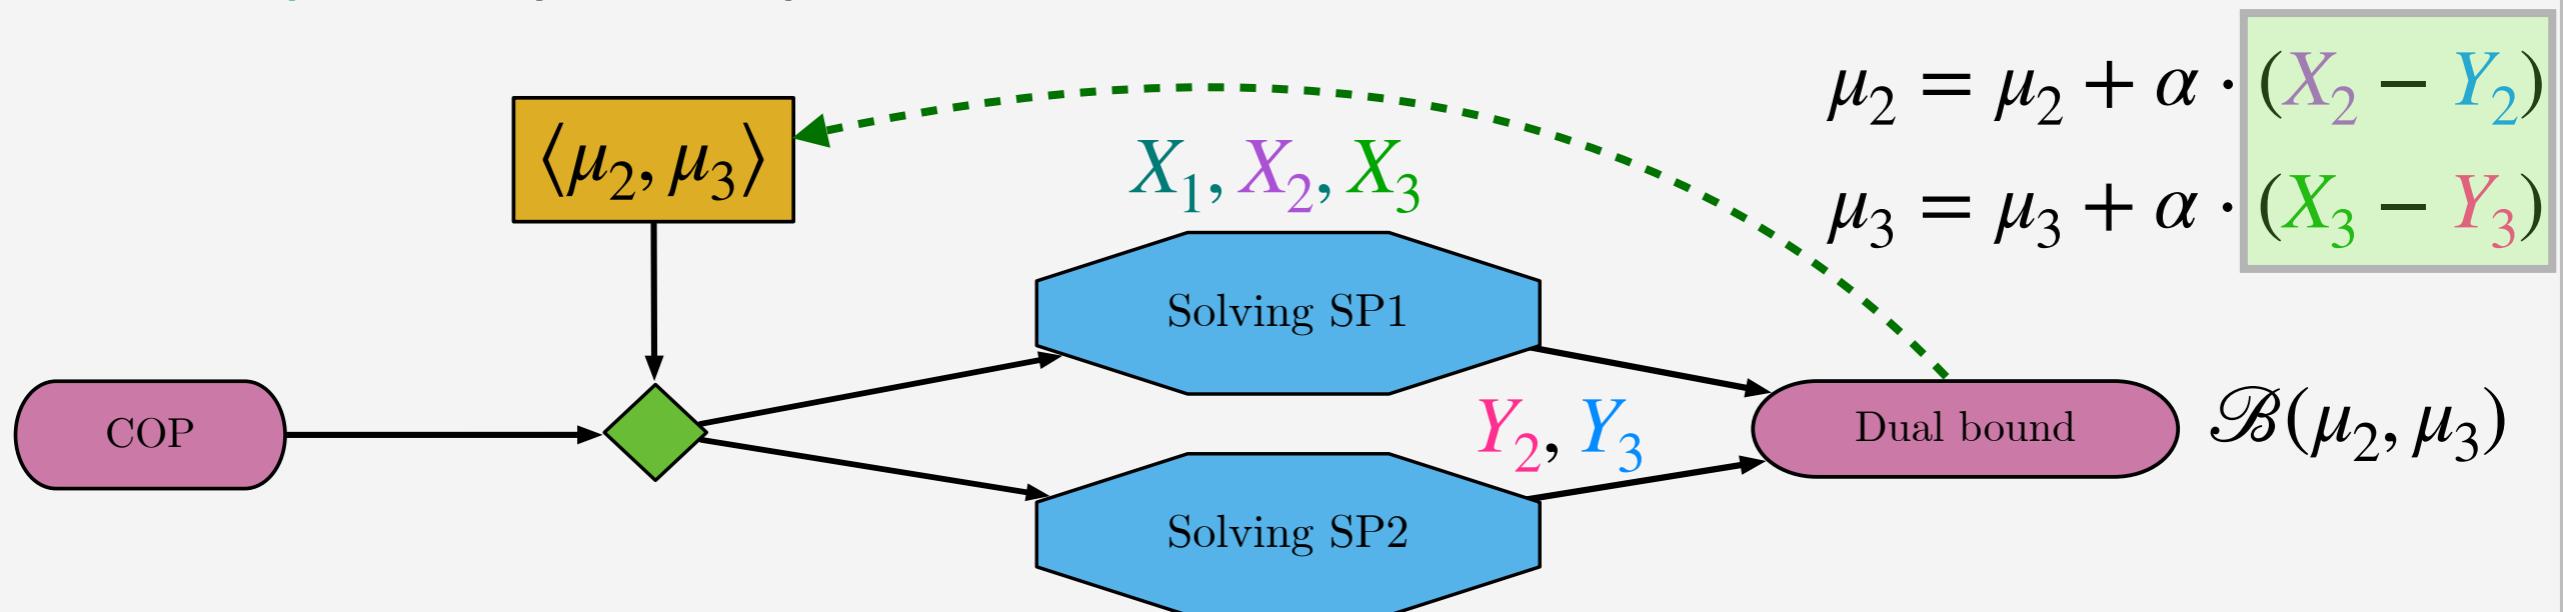
How to set the values of the multipliers ?

Initialization: we set the multipliers to an arbitrary value

Step 1: we solve all subproblems with these values (we get a dual bound)

Step 2: we update the multipliers with sub-gradient (we improve the bound)

Main loop: we repeat steps 1 and 2 for x iterations



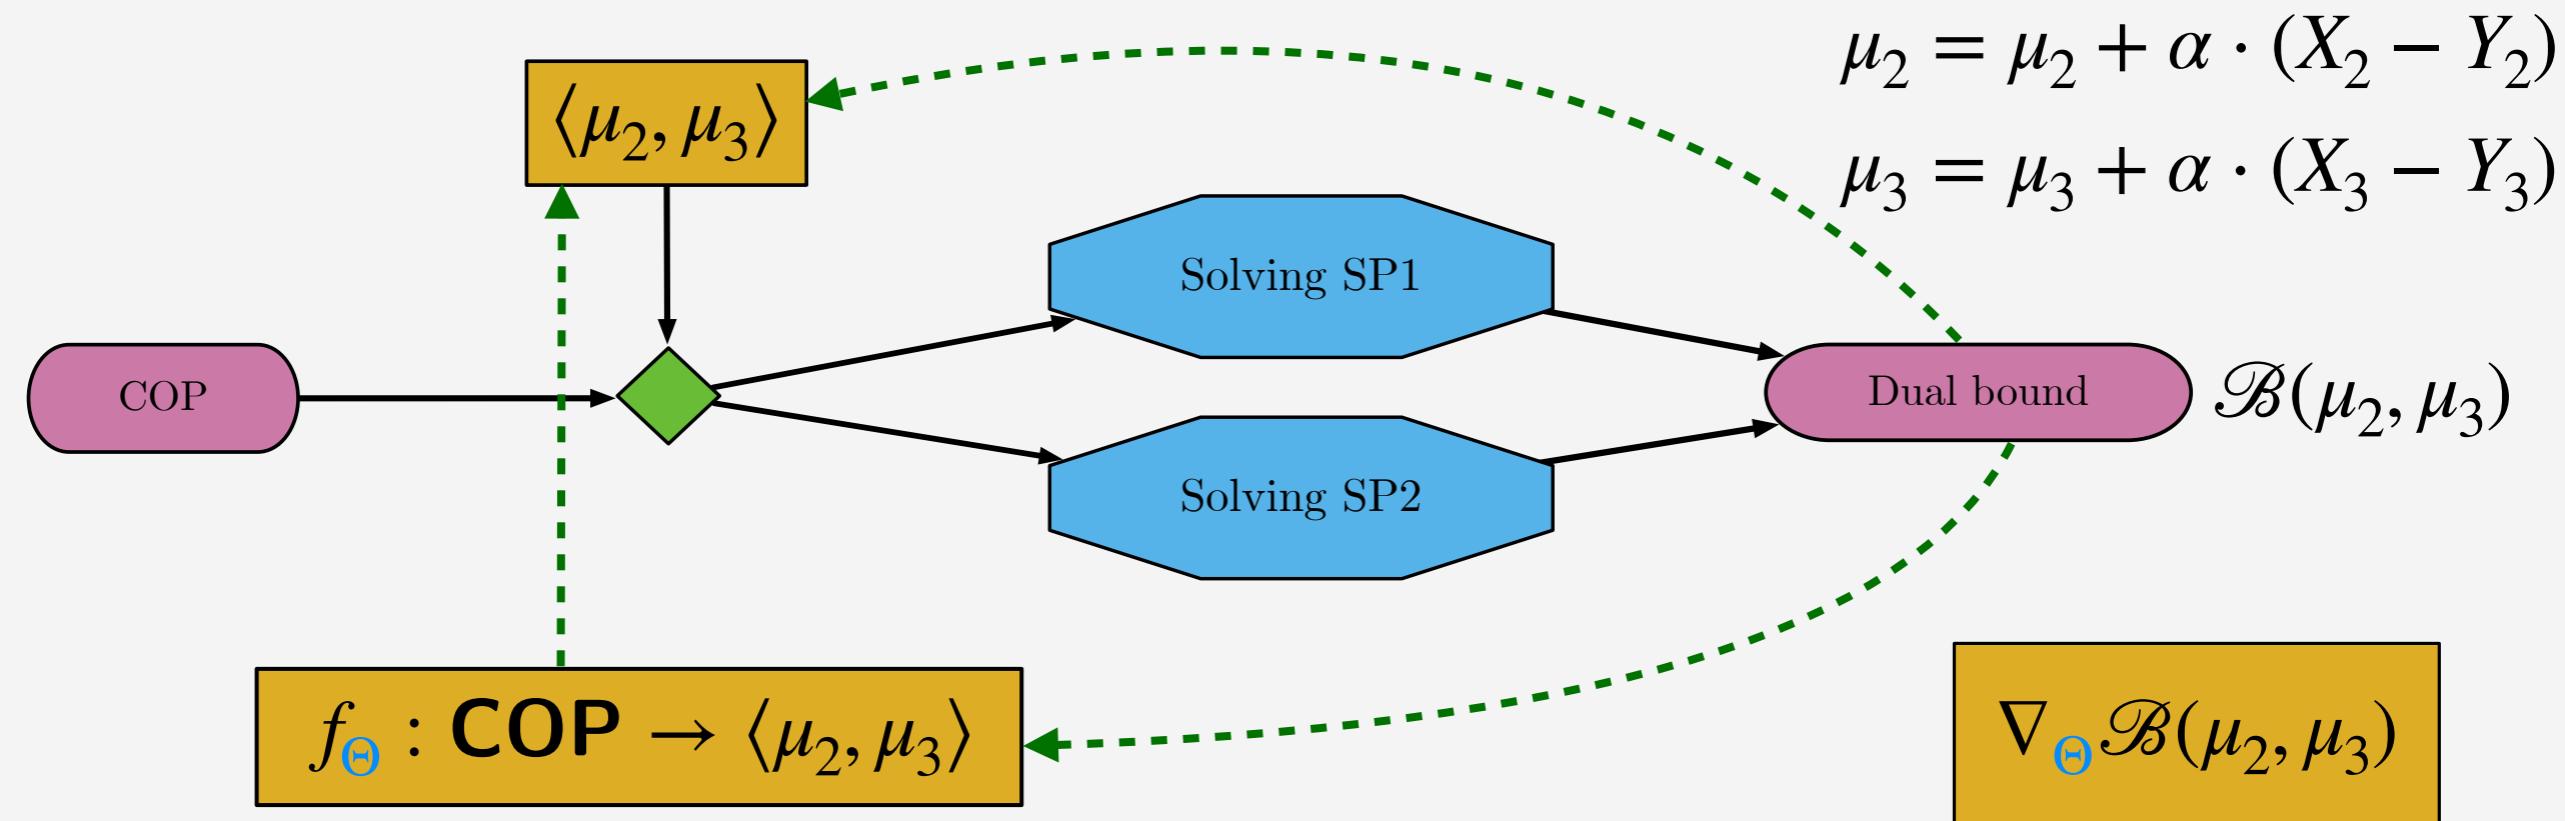
Learning multipliers for Lagrangian decomposition

This process is very costly as it requires solving few subproblems at each iteration

We propose a self-supervised learning approach to compute them

Step 1: multipliers are now obtained by a differentiable predictive model

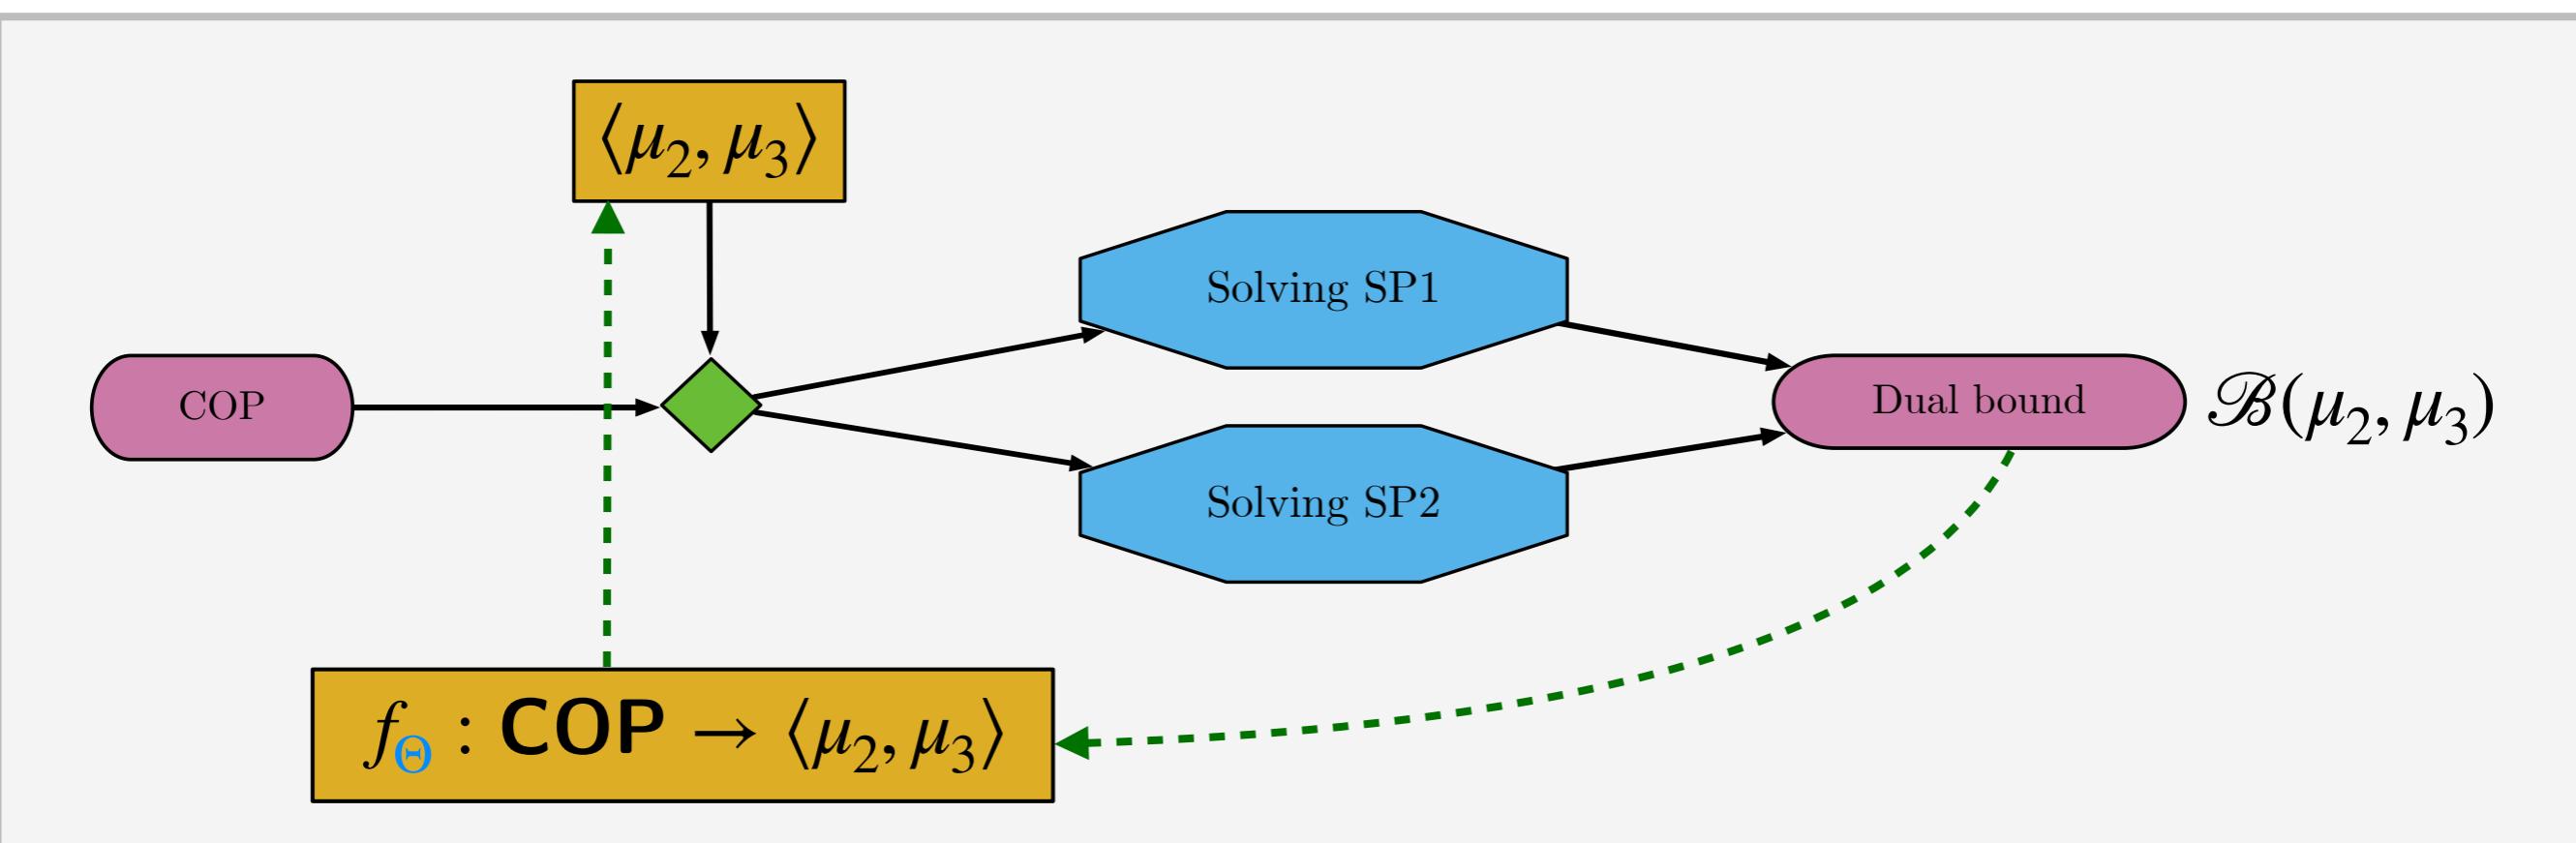
Step 2: the model is trained end-to-end by differentiating the bound



Intuition: the optimisation process is carried out offline during a training phase

Learning multipliers for Lagrangian decomposition

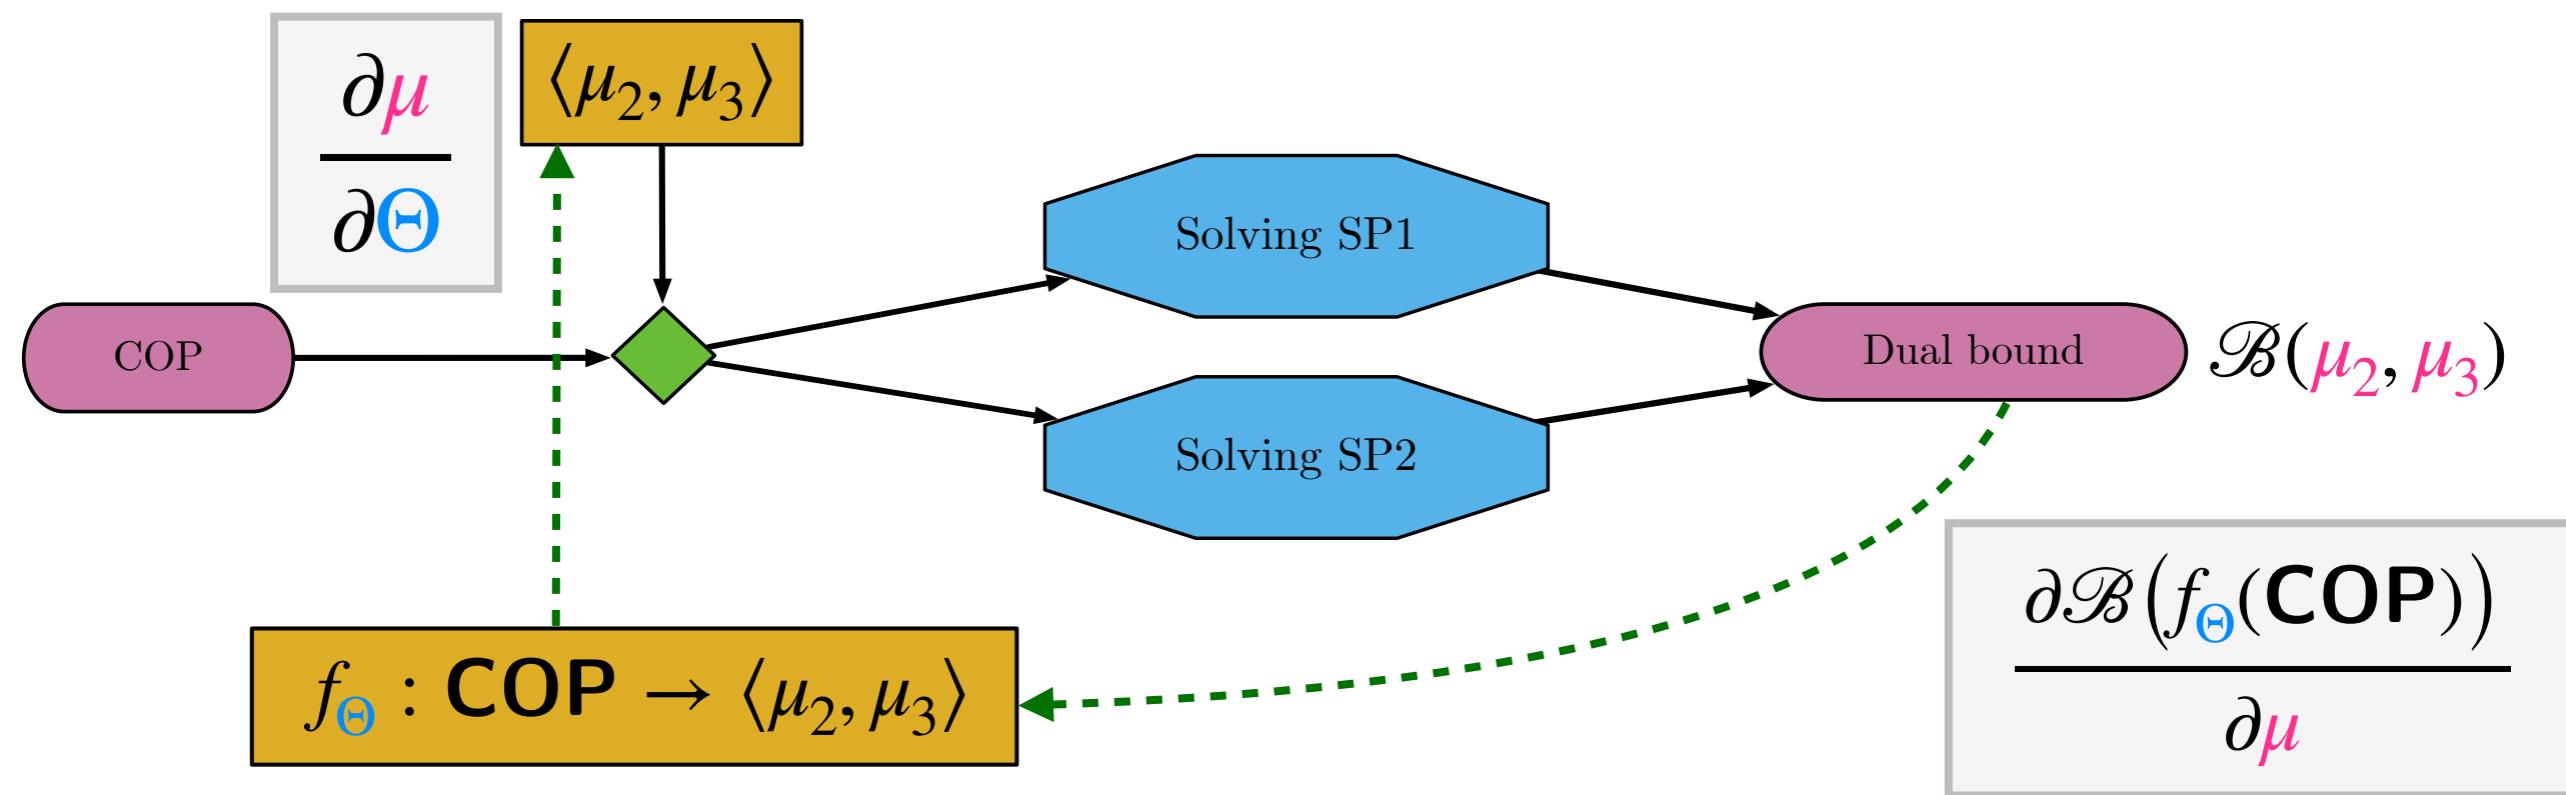
$$\mathcal{B}(\mu_2, \mu_3) = \begin{array}{l} \max \left(f(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) + \mu_2 \cdot \textcolor{violet}{X}_2 + \mu_3 \cdot \textcolor{green}{X}_3 \right) \\ \text{s.t. } C_1(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) \\ X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3 \in \mathbb{N}^+ \end{array} + \begin{array}{l} \max \left(-\mu_2 \cdot \textcolor{magenta}{Y}_2 - \mu_3 \cdot \textcolor{blue}{Y}_3 \right) \\ \text{s.t. } C_2(\textcolor{magenta}{Y}_2, \textcolor{blue}{Y}_3) \\ \textcolor{magenta}{Y}_2, \textcolor{blue}{Y}_3 \in \mathbb{N}^+ \end{array}$$



Gradient: $\nabla_{\Theta} \mathcal{B}(\mu)$

? How to compute the gradient of this bound ?

Learning multipliers for Lagrangian decomposition



$$\nabla_{\Theta} \mathcal{B}(\mu) = \frac{\partial \mathcal{B}(f_\Theta(\text{COP}))}{\partial \mu} \times \frac{\partial \mu}{\partial \Theta}$$

$$= (X - Y) \times \frac{\partial \mu}{\partial \Theta}$$

Step 1: we use the **chain-rule** to uncover dependencies

Step 2: right-term is a simple **backpropagation** in the predictive model

Step 3: left-term reuses the **initial sub-gradient expression**

Training: differentiating on training instances (no need to have labels)

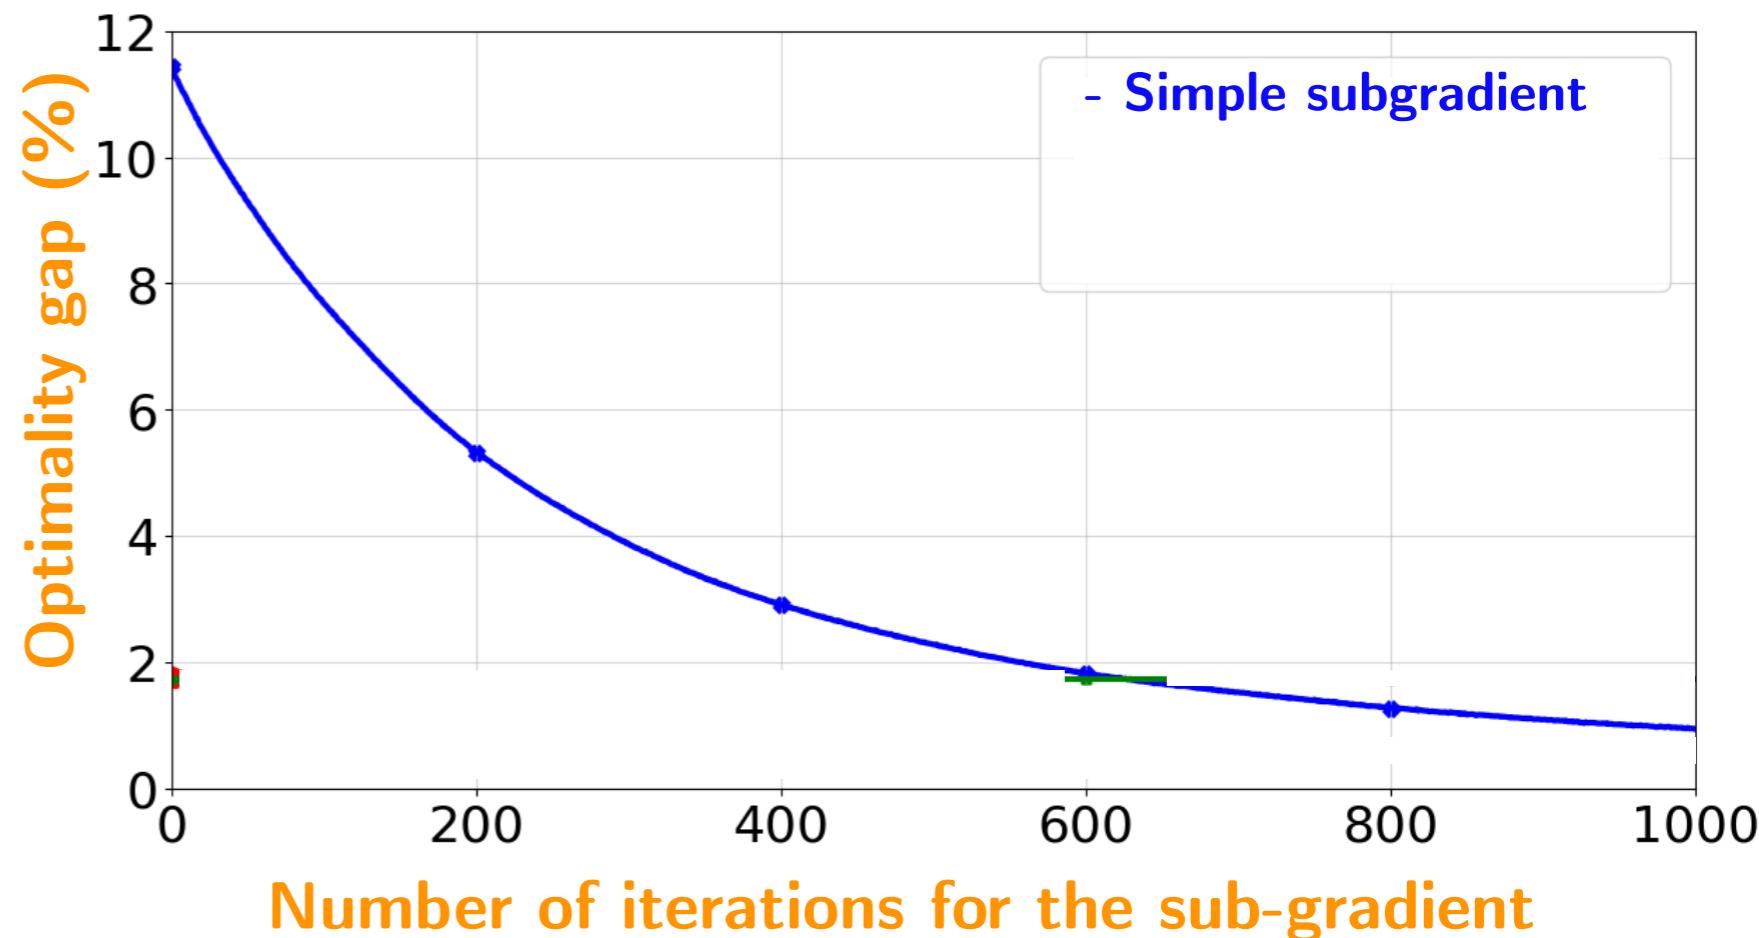
Qualitative results

Multidimensional knapsack



Size: 100 items

Dimension: 5 constraints



Metric: quality of the **dual bound** obtained at the root node

Observation 1: simple subgradient requires a lot of iterations to find good bounds

Observation 2: learning alone manages to get directly good bounds

Observation 3: learning to initialize sub-gradient quickly gives better bounds

Our results showed both the interest of learning alone or with sub-gradient

My current thoughts in learning to bound

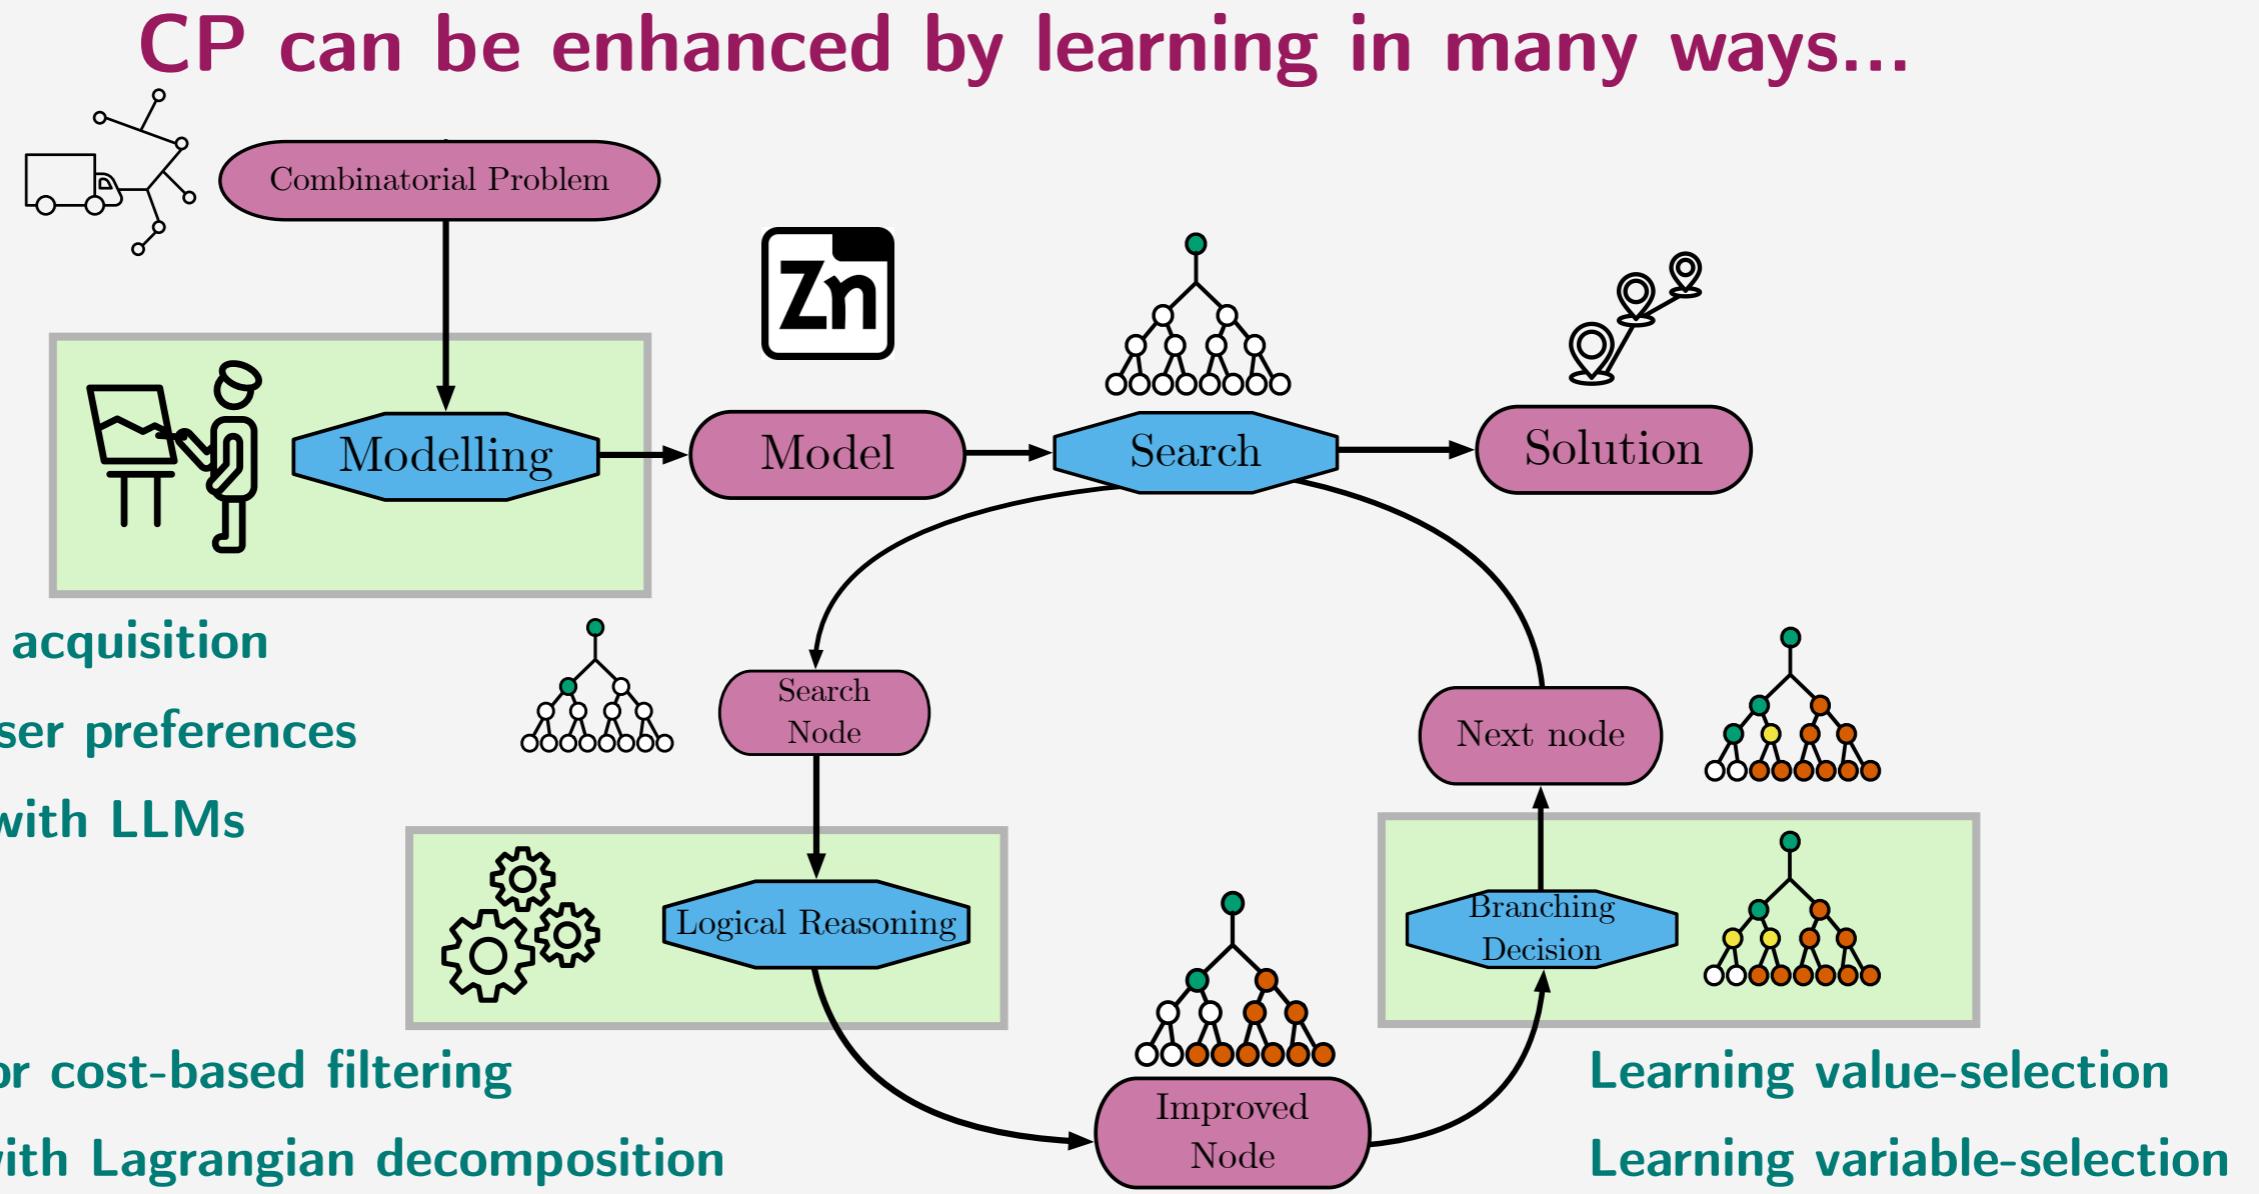
I think that learning to bound in CP is a promising direction

- (1) The lack of generic relaxation is a shortcoming of CP (compared to MIP solvers)
- (2) Lagrangian relaxation (or decomposition) ensures the validity of the bound
- (3) Learning can fully replace the subgradient (less filtering but cheaper)
- (4) Learning can also only initialize subgradient (better filtering but more expensive)
- (5) Results are encouraging in terms of execution time (more than for branching)

But there are still questions and challenges to address...

- (1) Which predictive model should we use? (is GNN a right choice?)
- (2) Can we have a generic representation of any COP (Boisvert et al., CPAIOR 2024)
- (3) How to handle problems with only few representative instances?
- (4) There are still the subproblems that need to be solved efficiently

Conclusion with my final notes



My take-home messages

- (1) Learning something meaningful does not always result in improved performances
- (2) Combining learning with existing tools mitigate the shortcomings of learning
- (3) Promising research direction, but a lot of challenges to handle for a practical use

Many thanks!!



Reference list (work we carried out)

Learning to branch

- ❖ Cappart, Q., Moisan, T., Rousseau, L. M., Prémont-Schwarz, I., & Cire, A. A.
Combining reinforcement learning and constraint programming for combinatorial optimization. [\[AAAI 2021\]](#)
- ❖ Chalumeau, F., Coulon, I., Cappart, Q., & Rousseau, L. M.
Seapearl: A constraint programming solver guided by reinforcement learning. [\[CPAIOR 2021\]](#)
- ❖ Marty, T., François, T., Tessier, P., Gautier, L., Rousseau, L. M., & Cappart, Q.
Learning a Generic Value-Selection Heuristic Inside a Constraint Programming Solver. [\[CP 2023 - distinguished paper\]](#)

Learning to bound

- ❖ Cappart, Q., Bergman, D., Rousseau, L. M., Prémont-Schwarz, I., & Parjadis, A.
Improving variable orderings of approximate decision diagrams using reinforcement learning. [\[IJOC 2022\]](#)
- ❖ Parjadis, A., Cappart, Q., Dilkina, B., Ferber, A., & Rousseau, L. M.
Learning Lagrangian Multipliers for the Travelling Salesman Problem. [\[CP 2024 - Best ML paper award\]](#)
- ❖ Dabert, D., Bessa, S., Bourgeat, M., Rousseau, L.M., Cappart, Q.
Learning Valid Dual Bounds in Constraint Programming [\[ArXivPreprint 2024\]](#)

Learning to model and graph neural networks

- ❖ Cappart, Q., Chételat, D., Khalil, E. B., Lodi, A., Morris, C., & Veličković, P.
Combinatorial optimization and reasoning with graph neural networks. [\[IJCAI 2021, JMLR 2023\]](#)
- ❖ Boisvert, L., Verhaeghe, H., & Cappart, Q.
Towards a Generic Representation of Combinatorial Problems for Learning-Based Approaches. [\[CPAIOR 2024\]](#)
- ❖ Barral, H., Gaha, M., Dems, A., Côté, A., Nguewouo, F., & Cappart, Q.
Acquiring Constraints for a Non-linear Transmission Maintenance Scheduling Problem. [\[CPAIOR 2024\]](#)

Reference list (other related works)

Learning to branch (also in MIP)

- ❖ Khalil, E., Le Bodic, P., Song, L., Nemhauser, G., & Dilkina, B.
Learning to branch in mixed integer programming. [AAAI 2016]
- ❖ Gasse, M., Chételat, D., Ferroni, N., Charlin, L., & Lodi, A.
Exact combinatorial optimization with graph convolutional neural networks. [NeurIPS 2019]
- ❖ Song, W., Cao, Z., Zhang, J., Xu, C., & Lim, A.
Learning variable ordering heuristics for solving constraint satisfaction problems. [EAAA 2022]

Foundations of our work on learning to prune

- ❖ Held, M., & Karp, R. The traveling-salesman problem and minimum spanning trees [Operations Research 1970]
- ❖ Guignard, M., & Kim, S. Lagrangean decomposition: A model yielding stronger Lagrangean bounds. [Math. Prog. 1987]
- ❖ Focacci, F., Lodi, A., & Milano, M. Cost-based domain filtering [CP 1999]
- ❖ Bergman, D., Cire, A. A., & van Hoeve, W. J. Improved constraint propagation via lagrangian decomposition. [CP 2015]
- ❖ Hà, M. H., Quimper, C. G., & Rousseau, L. M. General bounding mechanism for constraint programs. [CP 2015]

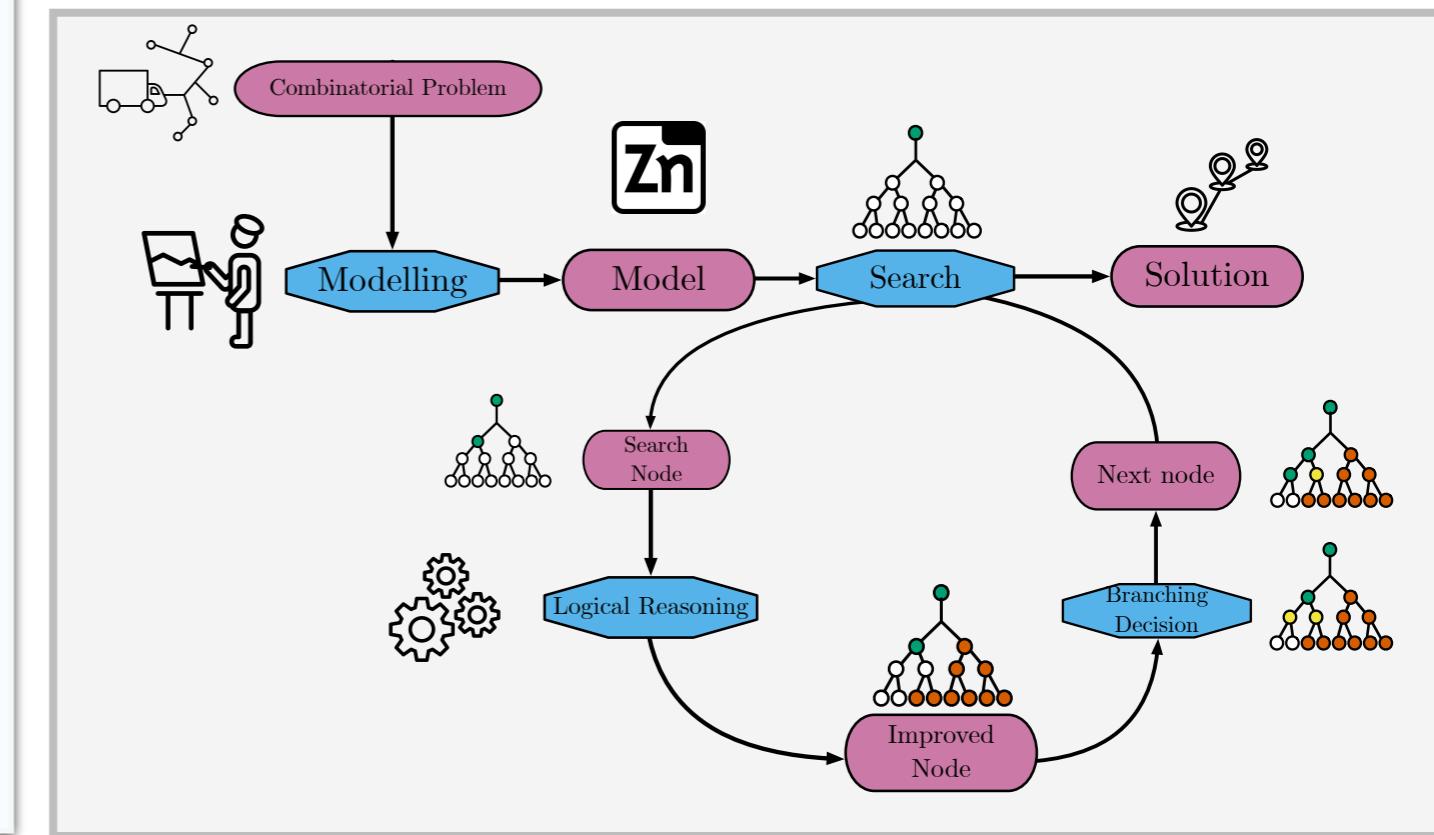
Learning dual bounds (only outside CP?)

- ❖ Deng, Y., Kong, S., Liu, C., & An, B.
Deep attentive belief propagation: Integrating reasoning and learning for solving constraint optimization problems. [NeurIPS 2022]
- ❖ Abbas, A., & Swoboda, P. Doge-train: Discrete optimization on gpu with end-to-end training. [AAAI 2024]

Please reach me if you know other works using learning to get dual bounds :-)



Research group: corail-research.github.io
Personal page: qcappart.github.io
Slides: QR code (or via my personal page)
Email: quentin.cappart@polymtl.ca



POLYTECHNIQUE
MONTRÉAL

Mila

CORAIL



Combinatorial Optimization and
Reasoning in
Artificial Intelligence
Laboratory



ACP
Quentin Cappart