

# Enhancing constraint programming with machine learning

Current challenges and future opportunities

CP 2024 - Girona



POLYTECHNIQUE  
MONTRÉAL

CORAIL




Combinatorial Optimization and  
Reasoning in  
Artificial Intelligence  
Laboratory



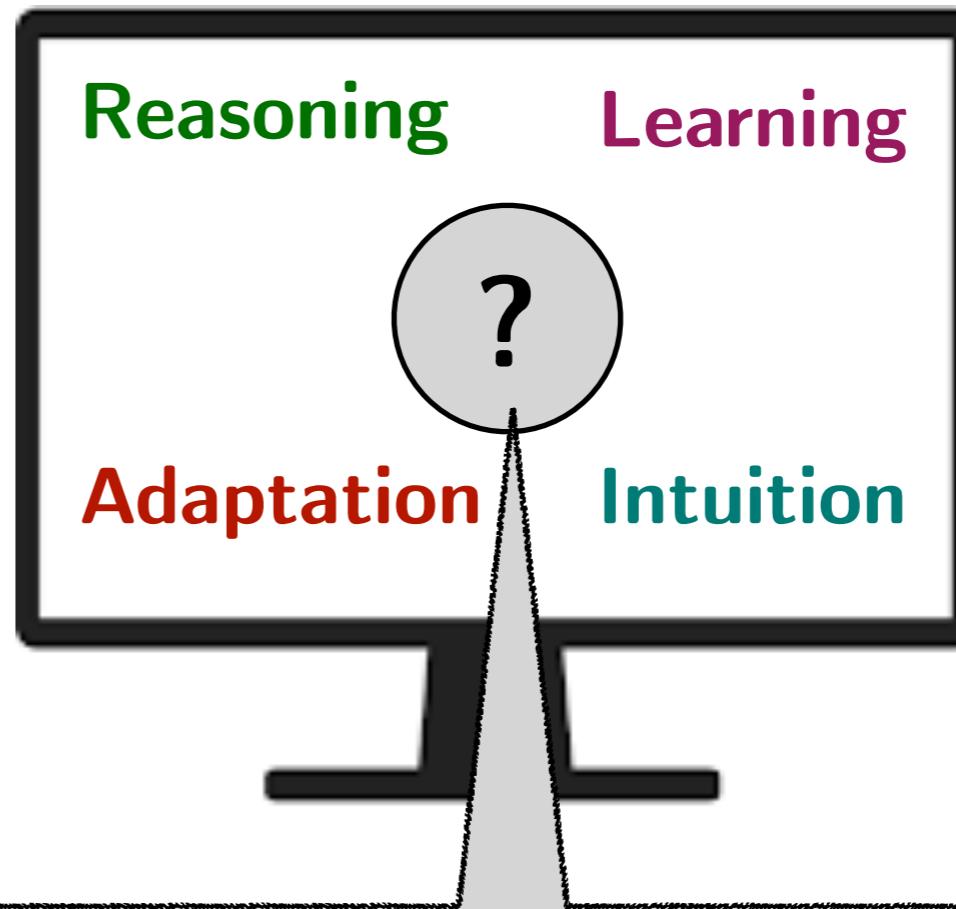
ACP  
Quentin Cappart

# Human intelligence versus artificial intelligence

## Human intelligence



## Artificial intelligence

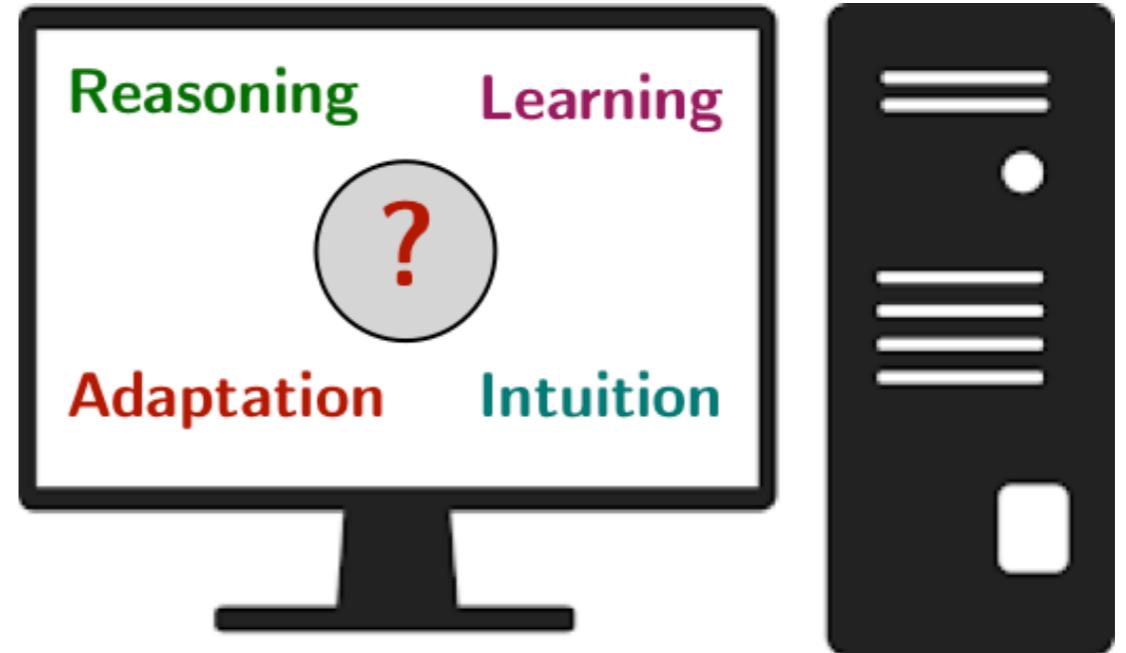


This connection is not yet established

**Long-term research plan: building an AI with these connections**


**Goal: providing a better solving process for combinatorial problems**

# Constraint programming as a unifying framework




## My research hypothesis

*Constraint programming can be a hosting technology for building this hybrid AI*



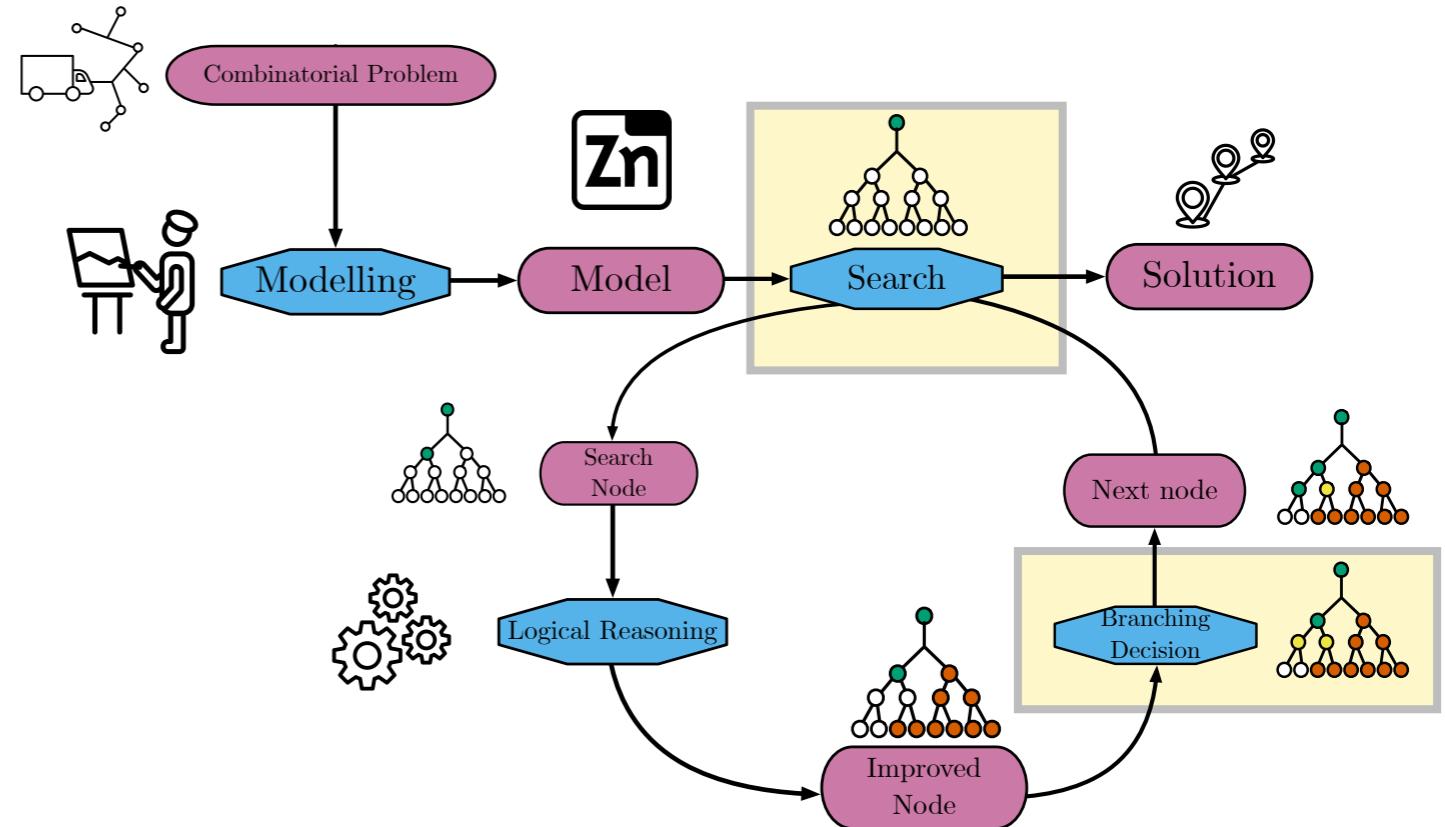
? How learning can be used to enhance the efficiency of CP solvers ?

# Summary of the CP pipeline



**Each decisional step has the potential to be improved with learning**

# Enhancing CP with learning



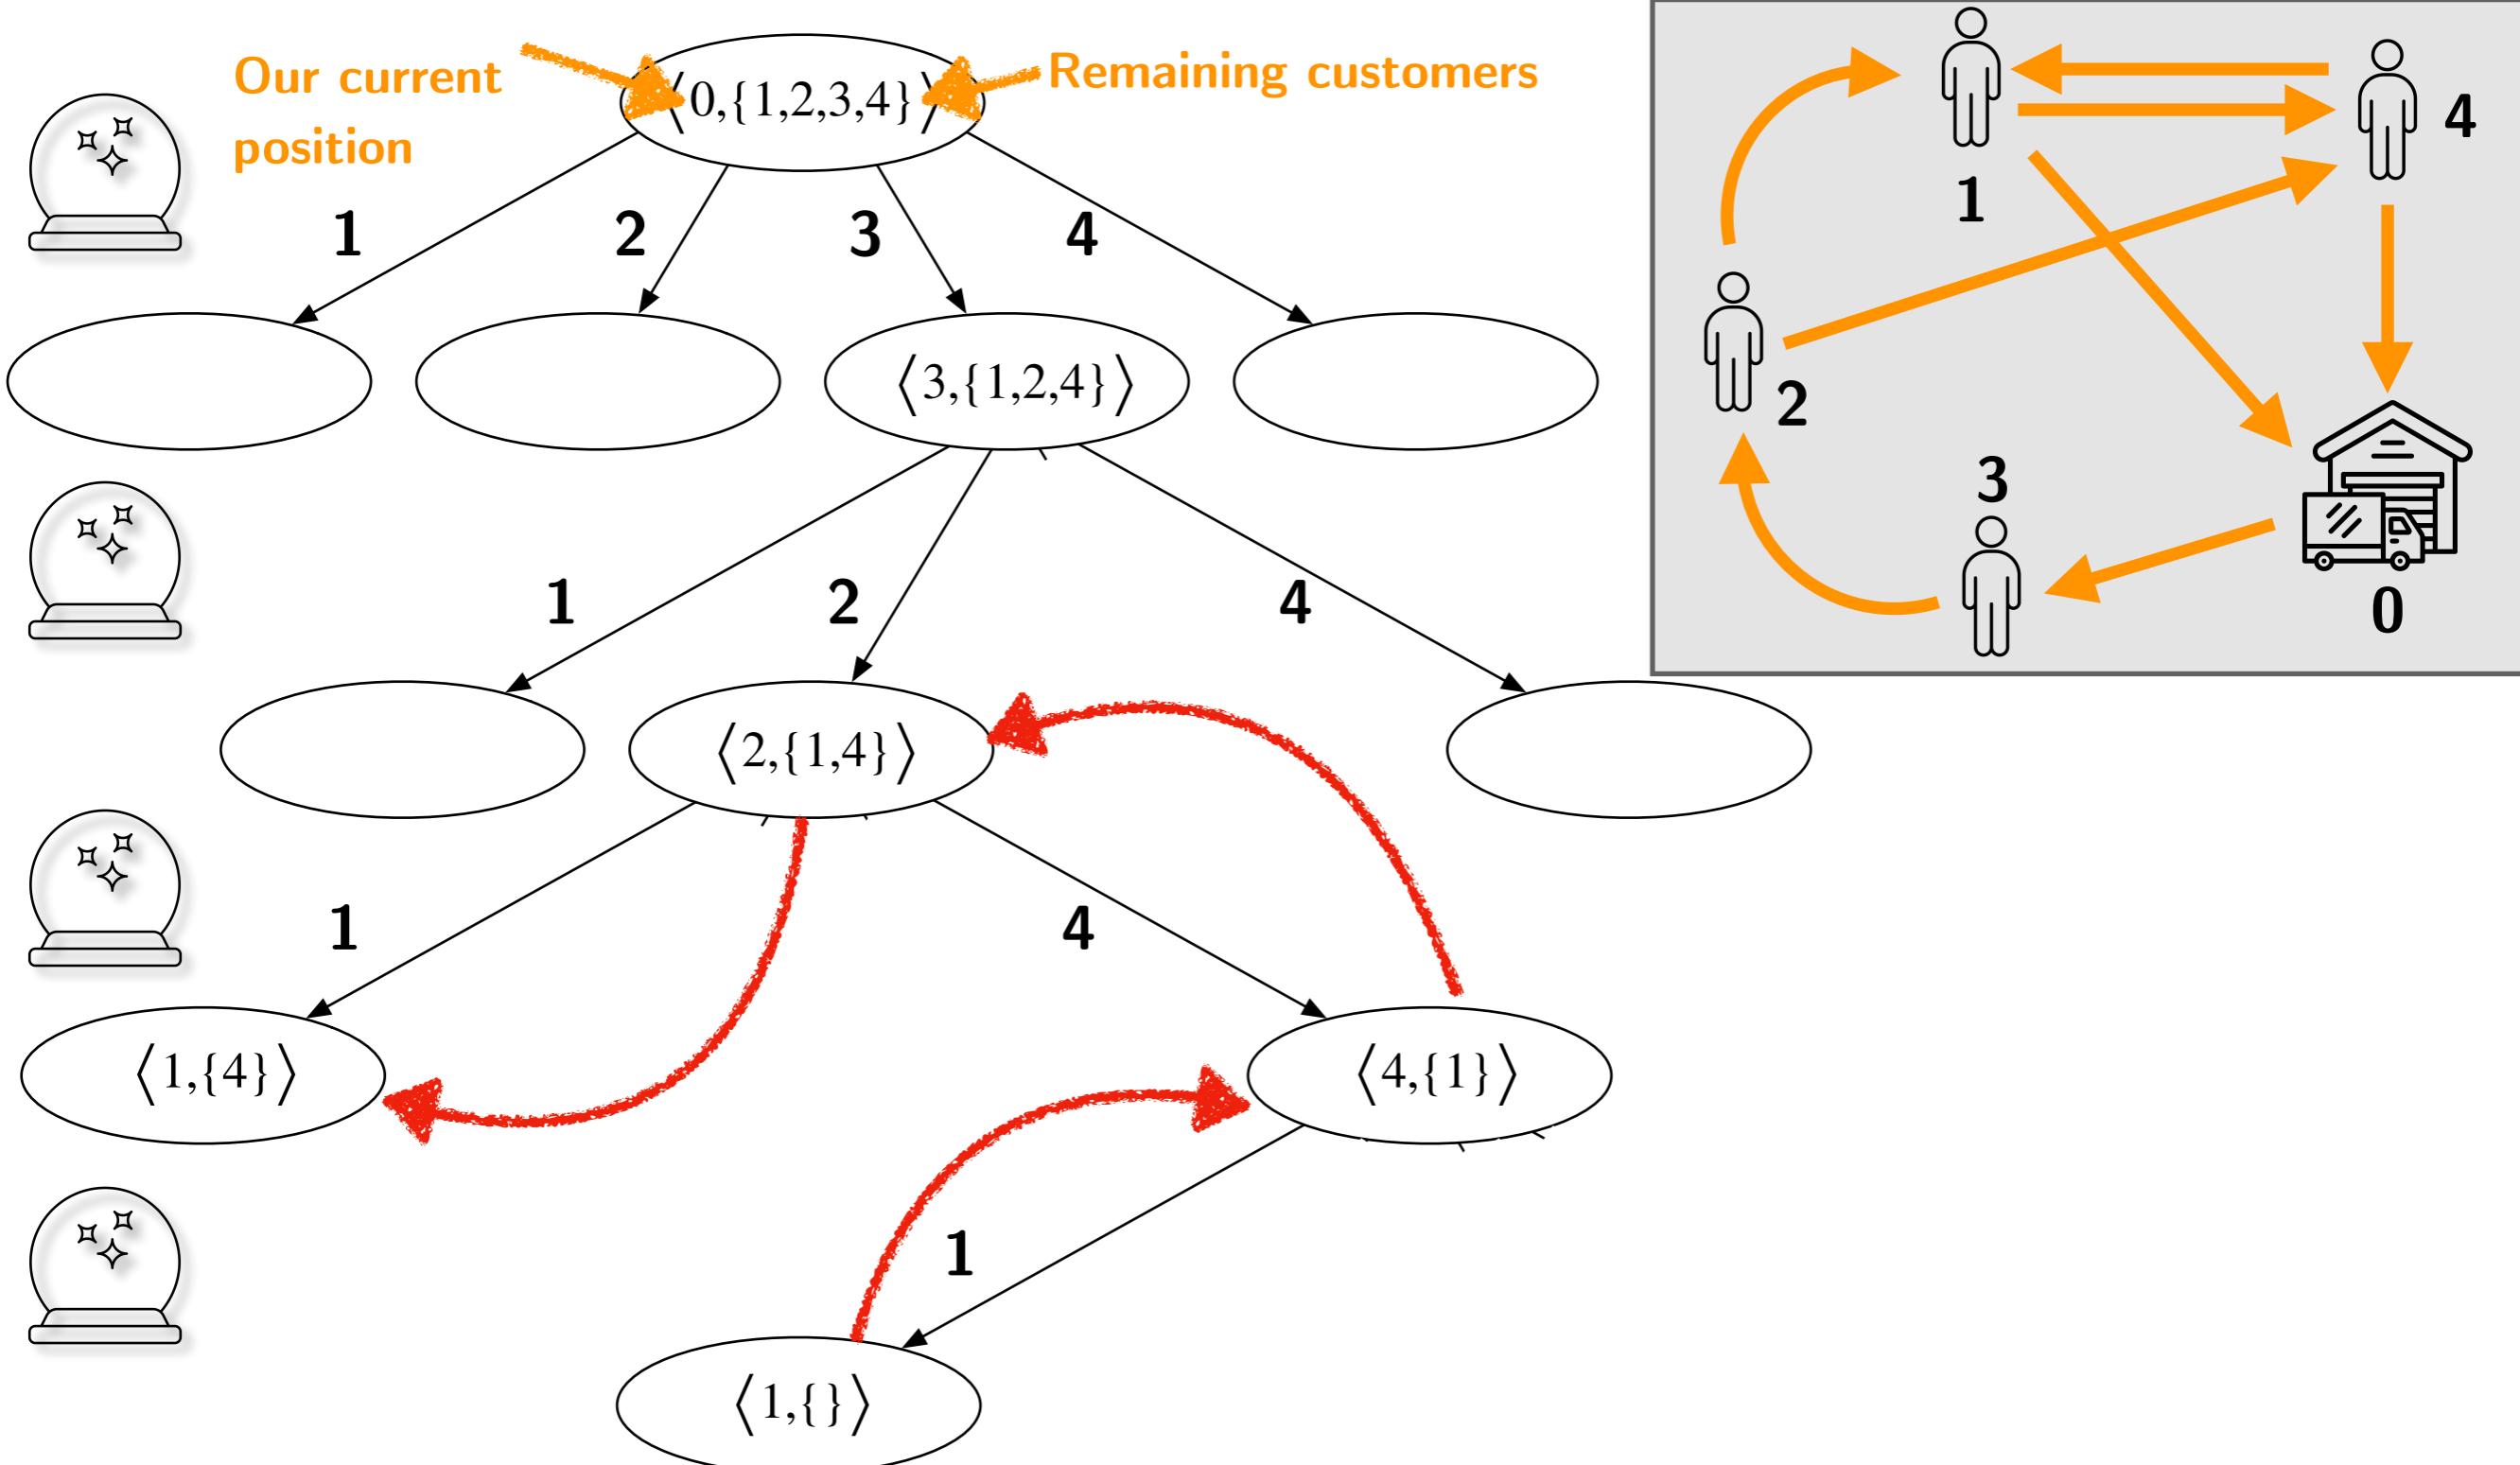
**My first intuition:** search is a good candidate as it heavily relies on **imperfect heuristics**

**Consolidating argument:** thanks to backtrack we can **recover from bad predictions**

**Context:** learning showed **promising results for MIP** (Gasse et al., Neurips 2019)

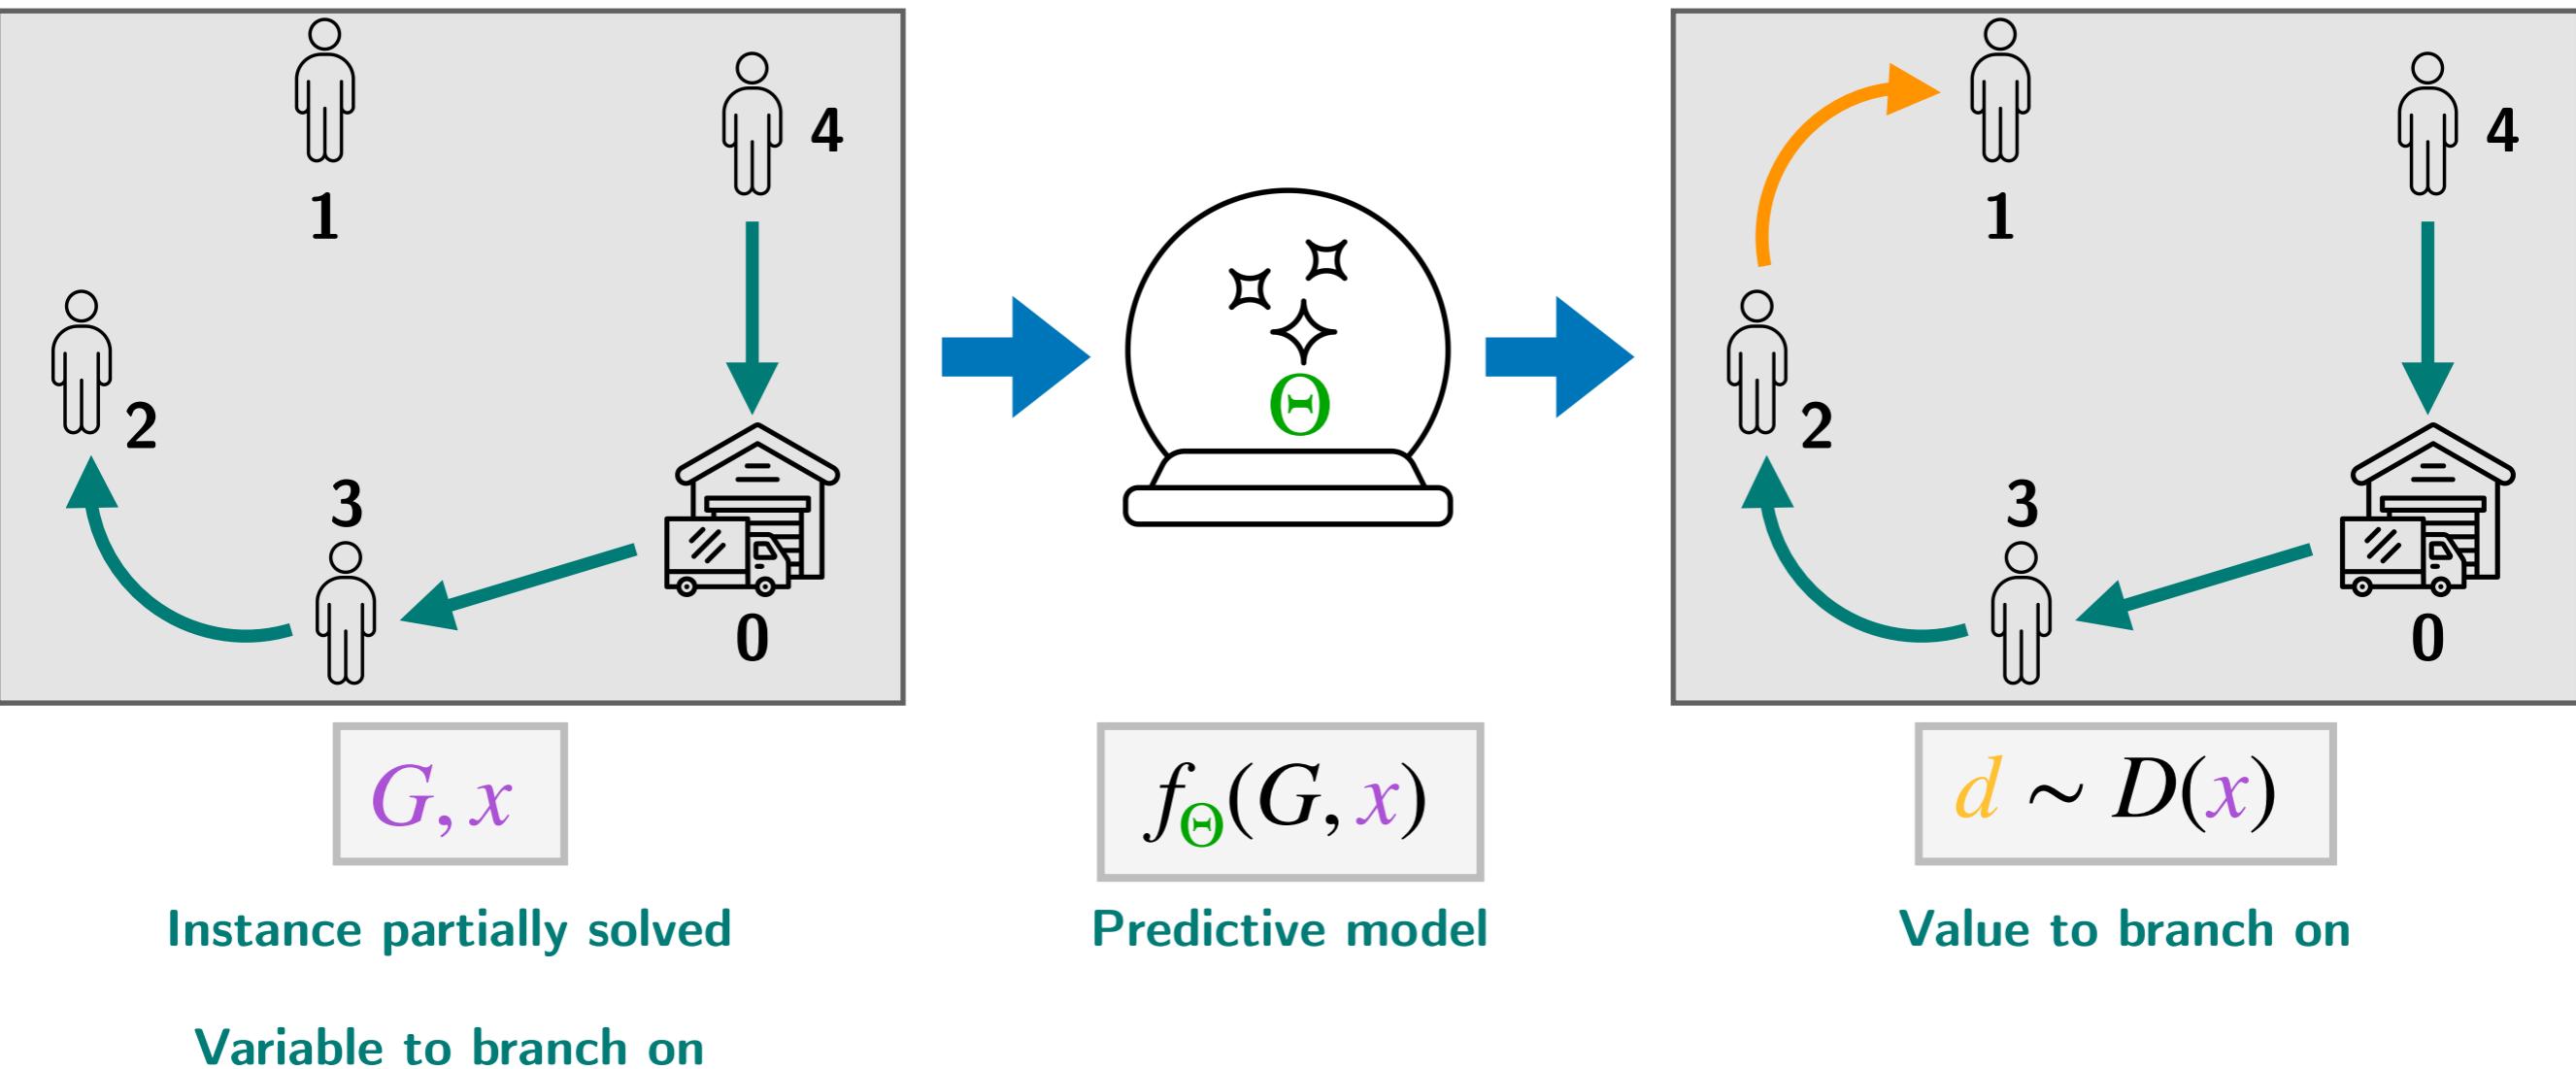
*My first research direction was to learn appropriate value-selection heuristic  
(with the intent to go further by learning variable-selection heuristics)*

# CP search with a learned heuristic on TSP



*The ability to recover from bad decisions is fundamental for complete approaches*

# CP search with a learned heuristic on TSP



?

How can we build this predictive model (this magical crystal ball) ?

**Challenge 1:** we do not know what is the best choice (i.e., we do not have labels)

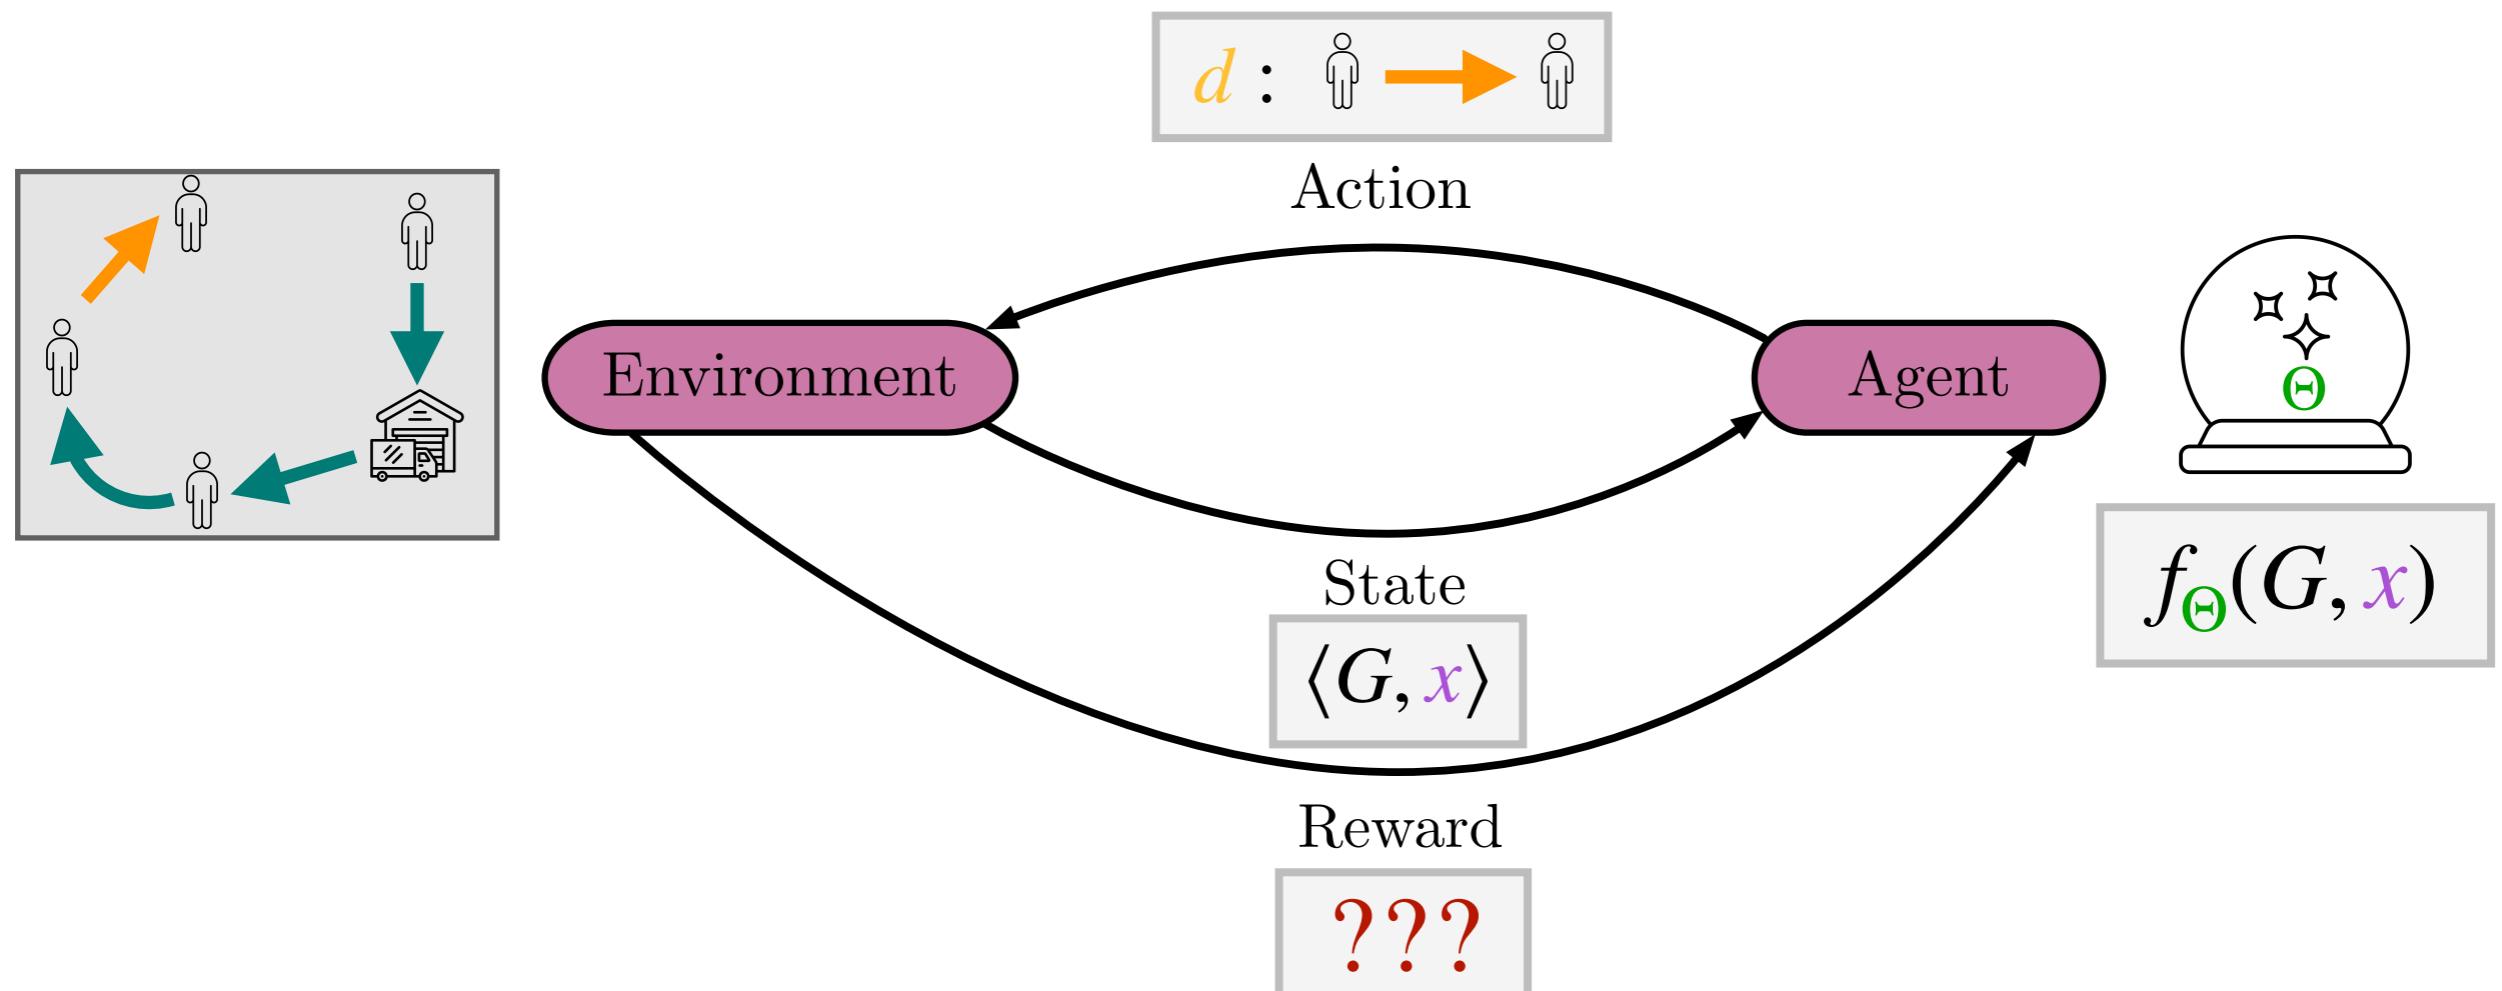
**Challenge 2:** difficult to represent a combinatorial problem in a proper way for learning

# Tackling the first challenge (Cappart et al., AAAI 2021)

**Challenge 1:** we do not know what is the best choice

**Consequence:** tricky to rely on supervised learning

**Our initial proposition:** leverage reinforcement learning instead



**Main idea:** unveiling the connections of RL with CP through dynamic programming

**Tricky part:** it is not clear how to design the reward (what is a good assignment?)

# Proposition of a reward function (Cappart et al., AAAI 2021)

? How to balance the fact that we want a **feasible** and **optimal** solution ?

**Observation:** **feasibility** is a prerequisite of **optimality**

**Key idea 1:** design a reward **always prioritizing finding feasible solutions**

**Key idea 2:** **discriminate then feasible solutions by their quality**

Reward for each action  $a$  at state  $s$

$$= R(s, a) + \text{UB}$$

$\text{UB} > R(s, a)$

$R(s, a)$  : direct increase in the objective function with assignment  $a$

$\text{UB}$  : strict upper bound on the optimal profit that can be reached

**Intuition:** **solutions deeper in the tree are always more rewarded than shallower ones**

**Consequence 1:** the first incentive is to go at the last depth (i.e., a feasible solution)

**Consequence 2:** **solutions at the same depth are rewarded by their objective value**

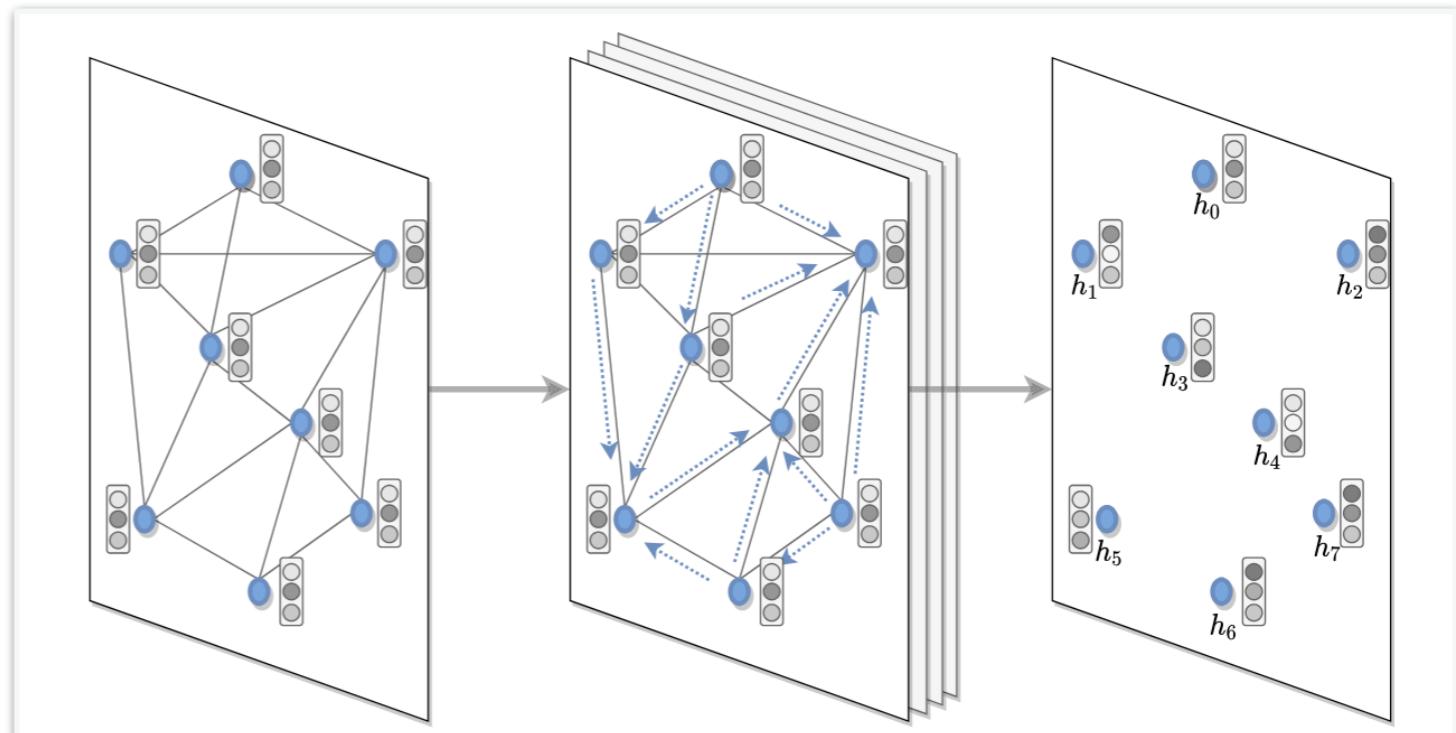
# A first proof of concept (Cappart et al., AAAI 2021)

## Algorithm 1: BaB-DQN Search Procedure.

```
> Pre:  $\mathcal{Q}_p$  is a COP having a DP formulation.  
> Pre:  $\mathbf{w}$  is a trained weight vector.  
 $\langle \mathbf{X}, \mathbf{D}, \mathbf{C}, \mathbf{o} \rangle := \text{CP Encoding}(\mathcal{Q}_p)$   
 $\mathcal{K} = \emptyset$   
 $\Psi := \text{BaB-search}(\langle \mathbf{X}, \mathbf{D}, \mathbf{C}, \mathbf{o} \rangle)$   
while  $\Psi$  is not completed do  
     $s := \text{encodeStateRL}(\Psi)$   
     $x := \text{takeFirstNonAssignedVar}(x)$   
    if  $s \in \mathcal{K}$  then  
         $v := \text{peek}(\mathcal{K}, s)$   
    else  
         $v := \text{argmax}_{u \in D(x)} \hat{Q}(s, u, \mathbf{w})$   
    end  
     $\mathcal{K} := \mathcal{K} \cup \{ \langle s, v \rangle \}$   
     $\text{branchAndUpdate}(\Psi, x, v)$   
end  
return bestSolution( $\Psi$ )
```



## Challenge 2: how to represent a COP for learning?



Representing a problem as a **graph**, and feeding it to a **graph neural network**

**Experiments:** results showed that **relevant branching decisions can be learned**


**Limitation:** far below state-of-the-art results and limited to relatively small instances

**Technical difficulty:** not-so-efficient to embed learning into existing CP solvers

*These first results showed the promise in this direction and drove my research*

# Improvements in learning to branch

(Chalumeau, Coulon et al., CPAIOR 2021)



SeaPearl.jl


**Seapearl:** minimalist CP solver aiming to ease the integration of learning

**Improvement:** carrying out the learning inside the solver and not outside

**Limitation 1:** huge overhead in calling a GNN at each search node

**Limitation 2:** challenging to train (large resources, instability, etc.)

(Marty et al., CP 2023)



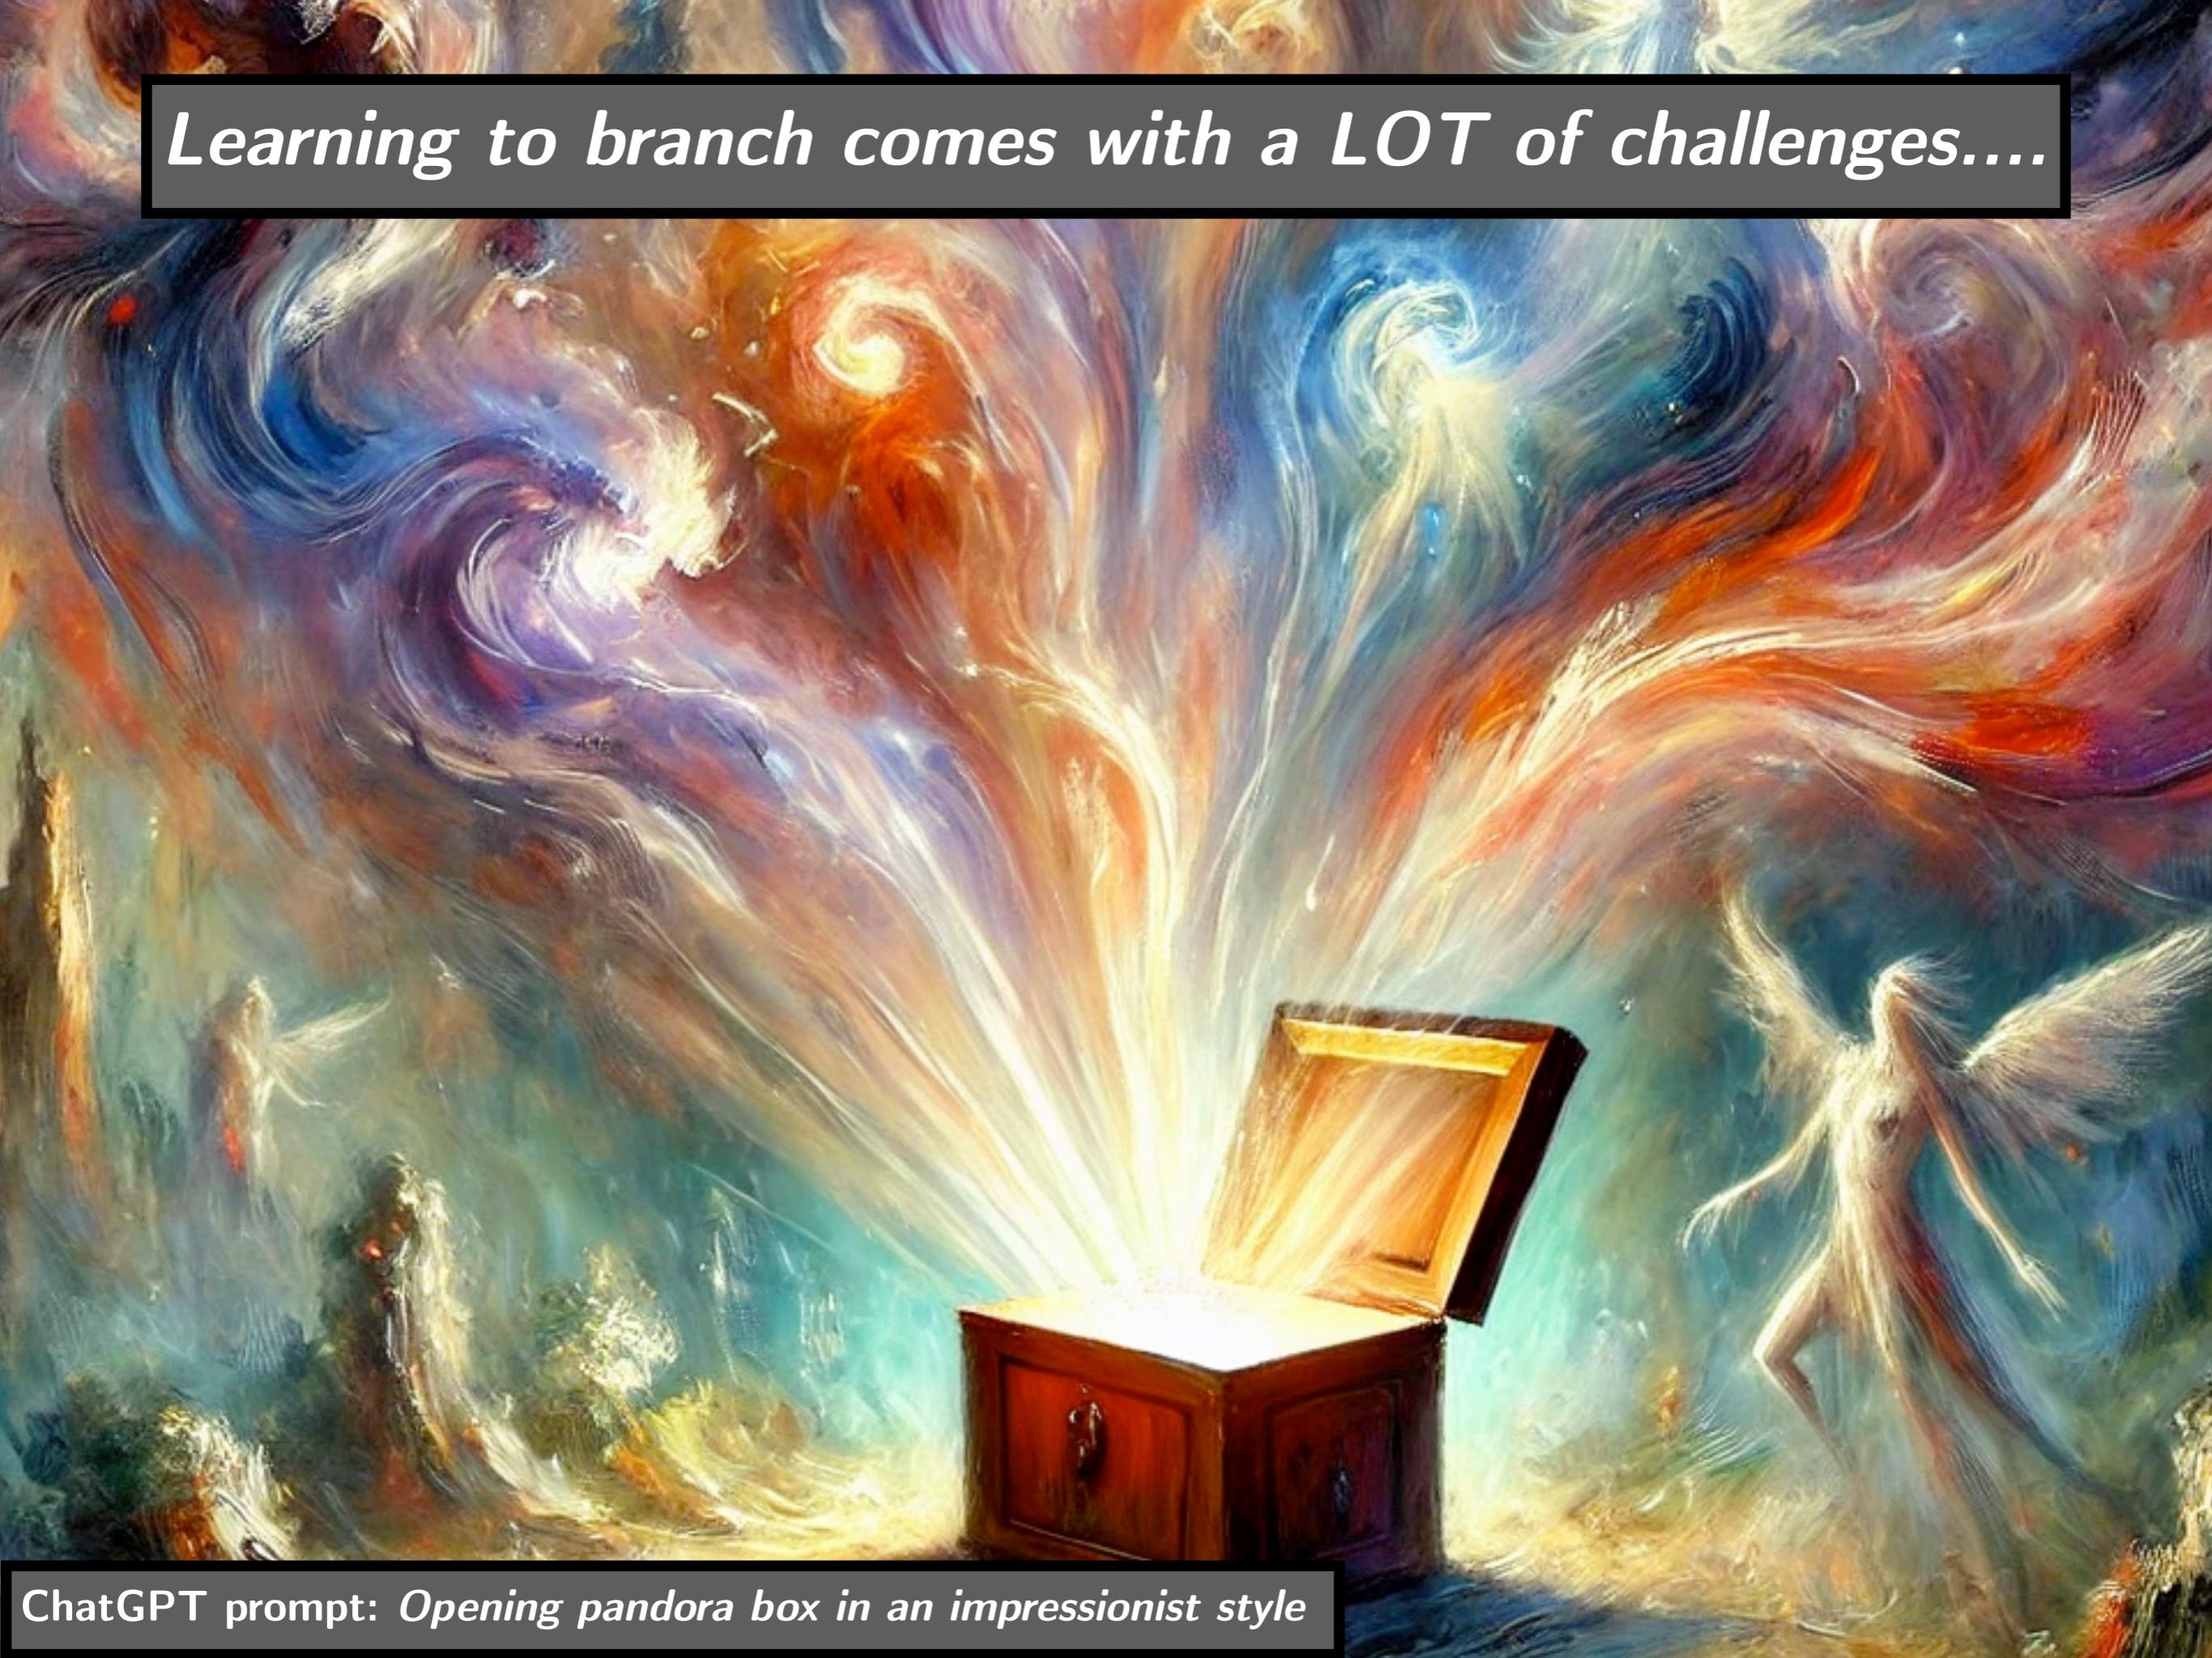
Main idea: mainly simplifying the framework!

Improvement: forget optimality proof (too difficult with value-selection heuristics)

New goal: finding quickly good solutions (redefining the reward function)

Limitation: still a huge overhead in calling often the GNN

*Learning to branch comes with a LOT of challenges....*



ChatGPT prompt: *Opening pandora box in an impressionist style*

# My current thoughts in learning to branch

Learning branching heuristics comes with a LOT of challenges...

- (1) Overhead in calling a heavy model at each node (compared to cheap heuristics)
- (2) Subjective choices in designing the reward (feasibility, quality, optimality)
- (3) A model good at the root node may be poor at deeper levels (*distributional shift*)

Here are few advices based on what worked the best for me

- (1) Do not use learning at each branching step (too costly compared to heuristics)
- (2) Call the model only at the top of the tree (where you likely have more samples)
- (3) Proving optimality may be out-of-range (how to properly reward this?)
- (4) Focus on finding quickly good solutions (in few nodes)
- (5) Prioritize hybrid approaches (learning at the beginning, then use heuristics)

# Going further with my thoughts...

*Reinforcement learning may not be the best method for improving CP*

**Argument 1:** much harder to train than supervised learning

**Argument 2:** indirectly require an unknown label (reward is an approximation of it)

*Learning a value-selection heuristic may not be the best thing to do*

**Observation 1:** learning is **too costly** and **unstable** to replace cheap heuristics

**Observation 2:** learning something does not always give practical improvements

**Note:** please only see this as **my personal opinion**, and not as an irrevocable truth :-)




What about learning a variable-selection heuristic for CP ?

**My intuition:** we will probably have similar challenges than for the value-selection

*Ok, but what do you propose then ?*

# Can we learn something else?



*My main current research direction is to learn how to prune the search space  
(Said differently, I plan to improve the quality of filtering)*

# Types of propagation in a CP solver



But how to do that ? Propagation is mainly algorithmic


$CP = \text{model} + \text{propagation} + \text{search}$



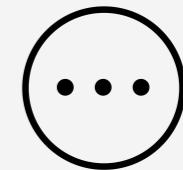
Fix-point



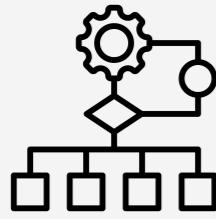
Consistency  
(AC3, etc.)




Global  
constraints



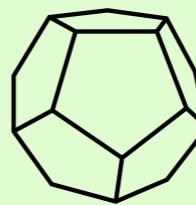
Other tools  
(nogoods, etc.)



Pure  
algorithmic



MDD



Cost-based  
filtering

*Some propagators relies on tricky-to-get information that we propose to learn*


# Cost-based filtering (Focacci et al., CP 1999)

Cost-based filtering leverages relaxation to improve the filtering of a constraint

Step 1: a valid relaxation is embedded into the global constraint

Step 2: the relaxed problem is solved to get a dual bound on the cost

Step 3: the bound is used to prune the search space



Trick: the quality of the relaxation depends on Lagrangian multipliers

Bad news: determining the best values of the multipliers is generally costly

We propose to obtain the multipliers through self-supervised learning



# Learning Lagrangian Multipliers for the TSP

**WeightedCircuit( $X, G, d$ )** : ensure that variables  $X$  form a TSP tour in  $G$  of cost  $\leq d$


**1-tree relaxation( $X, G$ )** : allows the TSP to go twice in a specific node

**Step 1:** ensure that variables  $X$  form a minimum spanning tree in  $G \setminus \{v_1\}$

**Step 2:** link the remaining node  $v_1$  to the tree with the two cheapest edges



Optimal TSP cost: 62




Relaxation cost: 50

We have here a example of how a specific constraint can be relaxed

# Learning Lagrangian Multipliers for the TSP

?

The bound is not very tight, could we improve it ?



Relaxation cost: 50


Lagrangian  
Multipliers

$$\langle \mu_1, \mu_2, \mu_3, \mu_4, \mu_5 \rangle$$

$$c'_{u,v} = c_{u,v} + \mu_u + \mu_v$$

$$c'_{2,3} = c_{2,3} + \mu_2 + \mu_3$$

$$9 = 5 + 4 + 0$$



Relaxation cost: 59

Main idea: perturbate the cost of each edge with values called **Lagrangian multipliers**

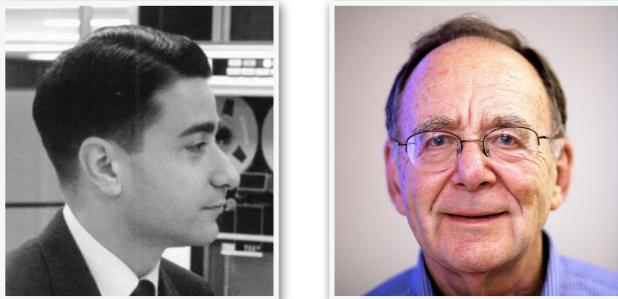
**Lagrangian multiplier:** value associated with each node

**Perturbation:** change the cost of each edge based on the multipliers of its nodes

**Nice property 1:** the optimal TSP tour is invariant under this perturbation

**Nice property 2:** better bounds generally result in a much better filtering

# Learning Lagrangian Multipliers for the TSP




?

But how can we find good values for the multipliers ?

Held and Karp, 1970: proposed an iterative algorithm for that

WeightedCircuit

1-tree relaxation




$$f_{\Theta}(G, x) : \langle \mu_1, \dots, \mu_n \rangle$$

Lagrangian  
Multipliers

Constraints (2012) 17:205–233  
DOI 10.1007/s10601-012-9119-x

Improved filtering for weighted circuit constraints

Pascal Benchimol · Willem-Jan van Hoeve ·  
Jean-Charles Régin · Louis-Martin Rousseau ·  
Michel Rueher



Held-Karp iterative process

Bad news: the iterative adjustment of the multipliers is computationally expensive

Parjadis et al. (CP 2024): use learning to initialize the multipliers

Idea: use a GNN to predict multipliers (one per node) from a TSP instance

Results: better filtering achieved on random and symmetric TSPs

Observation: replacing HK process with learning was not successful



# Types of propagation in a CP solver



Can we learn a more general bounding mechanism in CP ?


$CP = \text{model} + \text{propagation} + \text{search}$



Fix-point



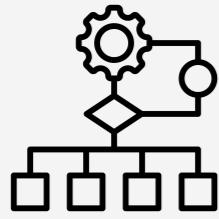
Consistency  
(AC3, etc.)



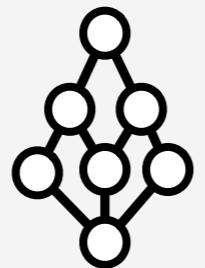
Global  
constraints



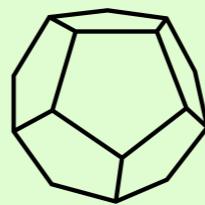
Other tools  
(nogoods, etc.)




Pure  
algorithmic



MDD



Cost-based  
filtering



Lagrangian  
decomposition

We extend the idea to learn multipliers to **CP-based Lagrangian decomposition**

# Lagrangian decomposition (Guignard and Kim, 1987)


Lagrangian decomposition splits the problem into independent and easier subproblems

Step 1: each variable in each constraint is duplicated, except for the first constraint

Step 2: a constraint linking the values is added for each new variable

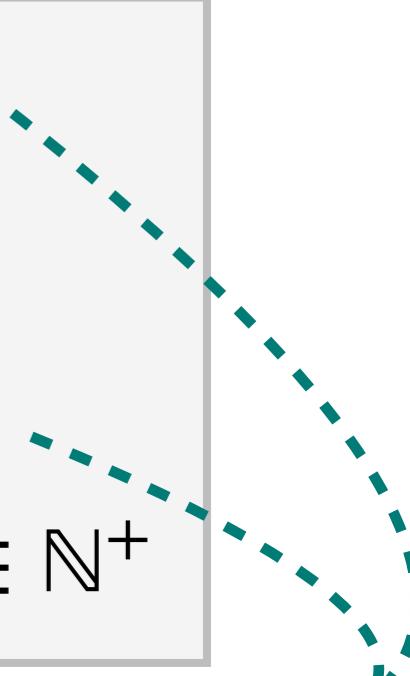
Step 3: these constraints are moved into the objective function with a penalty term

$$\begin{aligned} \max \quad & f(X_1, X_2, X_3) \\ \text{s.t.} \quad & C_1(X_1, X_2, X_3) \\ & C_2(X_2, X_3) \\ & X_1, X_2, X_3 \in \mathbb{N}^+ \end{aligned}$$



$$\begin{aligned} \max \quad & f(X_1, X_2, X_3) \\ \text{s.t.} \quad & C_1(X_1, X_2) \\ & C_2(Y_2, Y_3) \\ & Y_2 = X_2, Y_3 = X_3 \\ & X_1, X_2, Y_2, X_3, Y_3 \in \mathbb{N}^+ \end{aligned}$$

$$\max \quad f(X_1, X_2, X_3) + \boxed{\mu_2} \cdot (Y_2 - X_2) + \boxed{\mu_3} \cdot (Y_3 - X_3)$$



Again, we have Lagrangian multipliers, but in a more generic way than before

# Lagrangian decomposition (Guignard and Kim, 1987)

$$\mathcal{B}(\mu_2, \mu_3) = \begin{cases} \max & f(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) + \mu_2 \cdot (\textcolor{violet}{X}_2 - \textcolor{red}{Y}_2) + \mu_3 \cdot (\textcolor{green}{X}_3 - \textcolor{blue}{Y}_3) \\ \text{s.t.} & C_1(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) \\ & C_2(\textcolor{red}{Y}_2, \textcolor{blue}{Y}_3) \\ & X_1, \textcolor{violet}{X}_2, \textcolor{red}{Y}_2, \textcolor{green}{X}_3, \textcolor{blue}{Y}_3 \in \mathbb{N}^+ \end{cases}$$

Solving this relaxed problem will give a dual bound

? But, is it easy to solve ?

**Observation:** by construction, each constraint has its own set of variables

**Consequence:** each constraint can be solved independently

$$\mathcal{B}(\mu_2, \mu_3) = \begin{array}{l} \max \left( f(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) + \mu_2 \cdot \textcolor{violet}{X}_2 + \mu_3 \cdot \textcolor{green}{X}_3 \right) \\ \text{s.t. } C_1(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) \\ X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3 \in \mathbb{N}^+ \end{array} + \begin{array}{l} \max \left( -\mu_2 \cdot \textcolor{red}{Y}_2 - \mu_3 \cdot \textcolor{blue}{Y}_3 \right) \\ \text{s.t. } C_2(\textcolor{red}{Y}_2, \textcolor{blue}{Y}_3) \\ \textcolor{red}{Y}_2, \textcolor{blue}{Y}_3 \in \mathbb{N}^+ \end{array}$$

Given some multipliers, we can obtain a bound by solving several subproblems

# Lagrangian decomposition in CP (Hà et al. CP 2015)

$$\mathcal{B}(\mu_2, \mu_3) = \begin{array}{l} \max \left( f(X_1, X_2, X_3) + \mu_2 \cdot X_2 + \mu_3 \cdot X_3 \right) \\ \text{s.t. } C_1(X_1, X_2, X_3) \\ X_1, X_2, X_3 \in \mathbb{N}^+ \end{array} + \begin{array}{l} \max \left( -\mu_2 \cdot Y_2 - \mu_3 \cdot Y_3 \right) \\ \text{s.t. } C_2(Y_2, Y_3) \\ Y_2, Y_3 \in \mathbb{N}^+ \end{array}$$



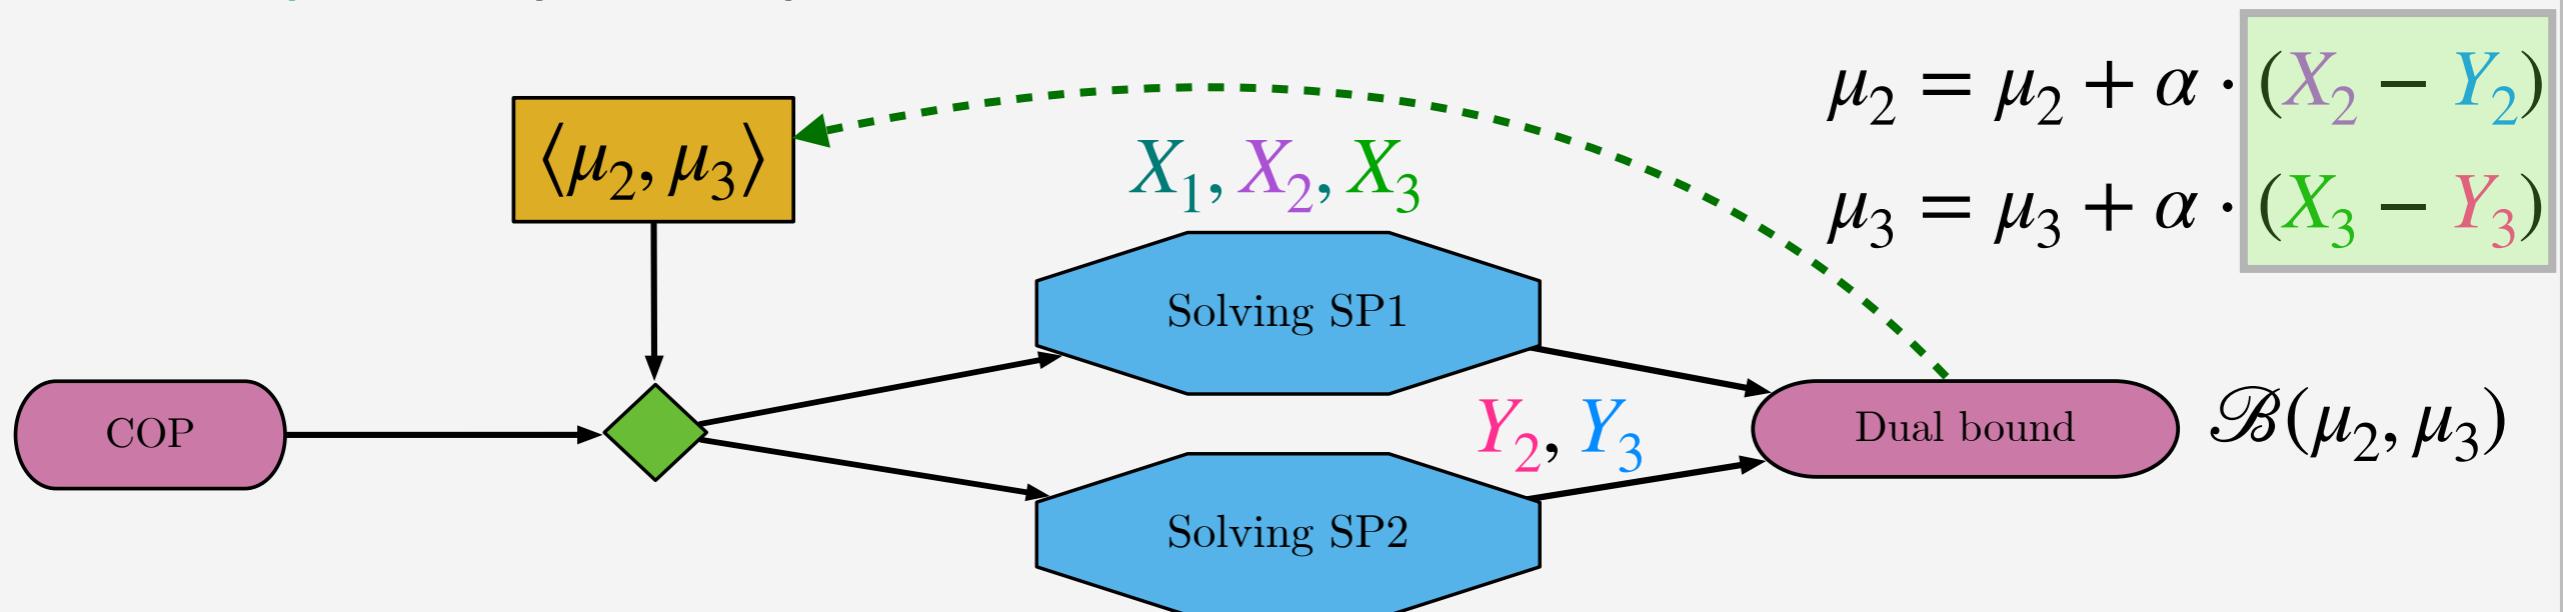
How to set the values of the multipliers ?

**Initialization:** we set the multipliers to an arbitrary value

**Step 1:** we solve all subproblems with these values (we get a dual bound)

**Step 2:** we update the multipliers with sub-gradient (we improve the bound)

**Main loop:** we repeat steps 1 and 2 for x iterations




# Learning multipliers for Lagrangian decomposition

This process is very costly as it requires solving few subproblems at each iteration

We propose a self-supervised learning approach to compute them

Step 1: multipliers are now obtained by a differentiable predictive model

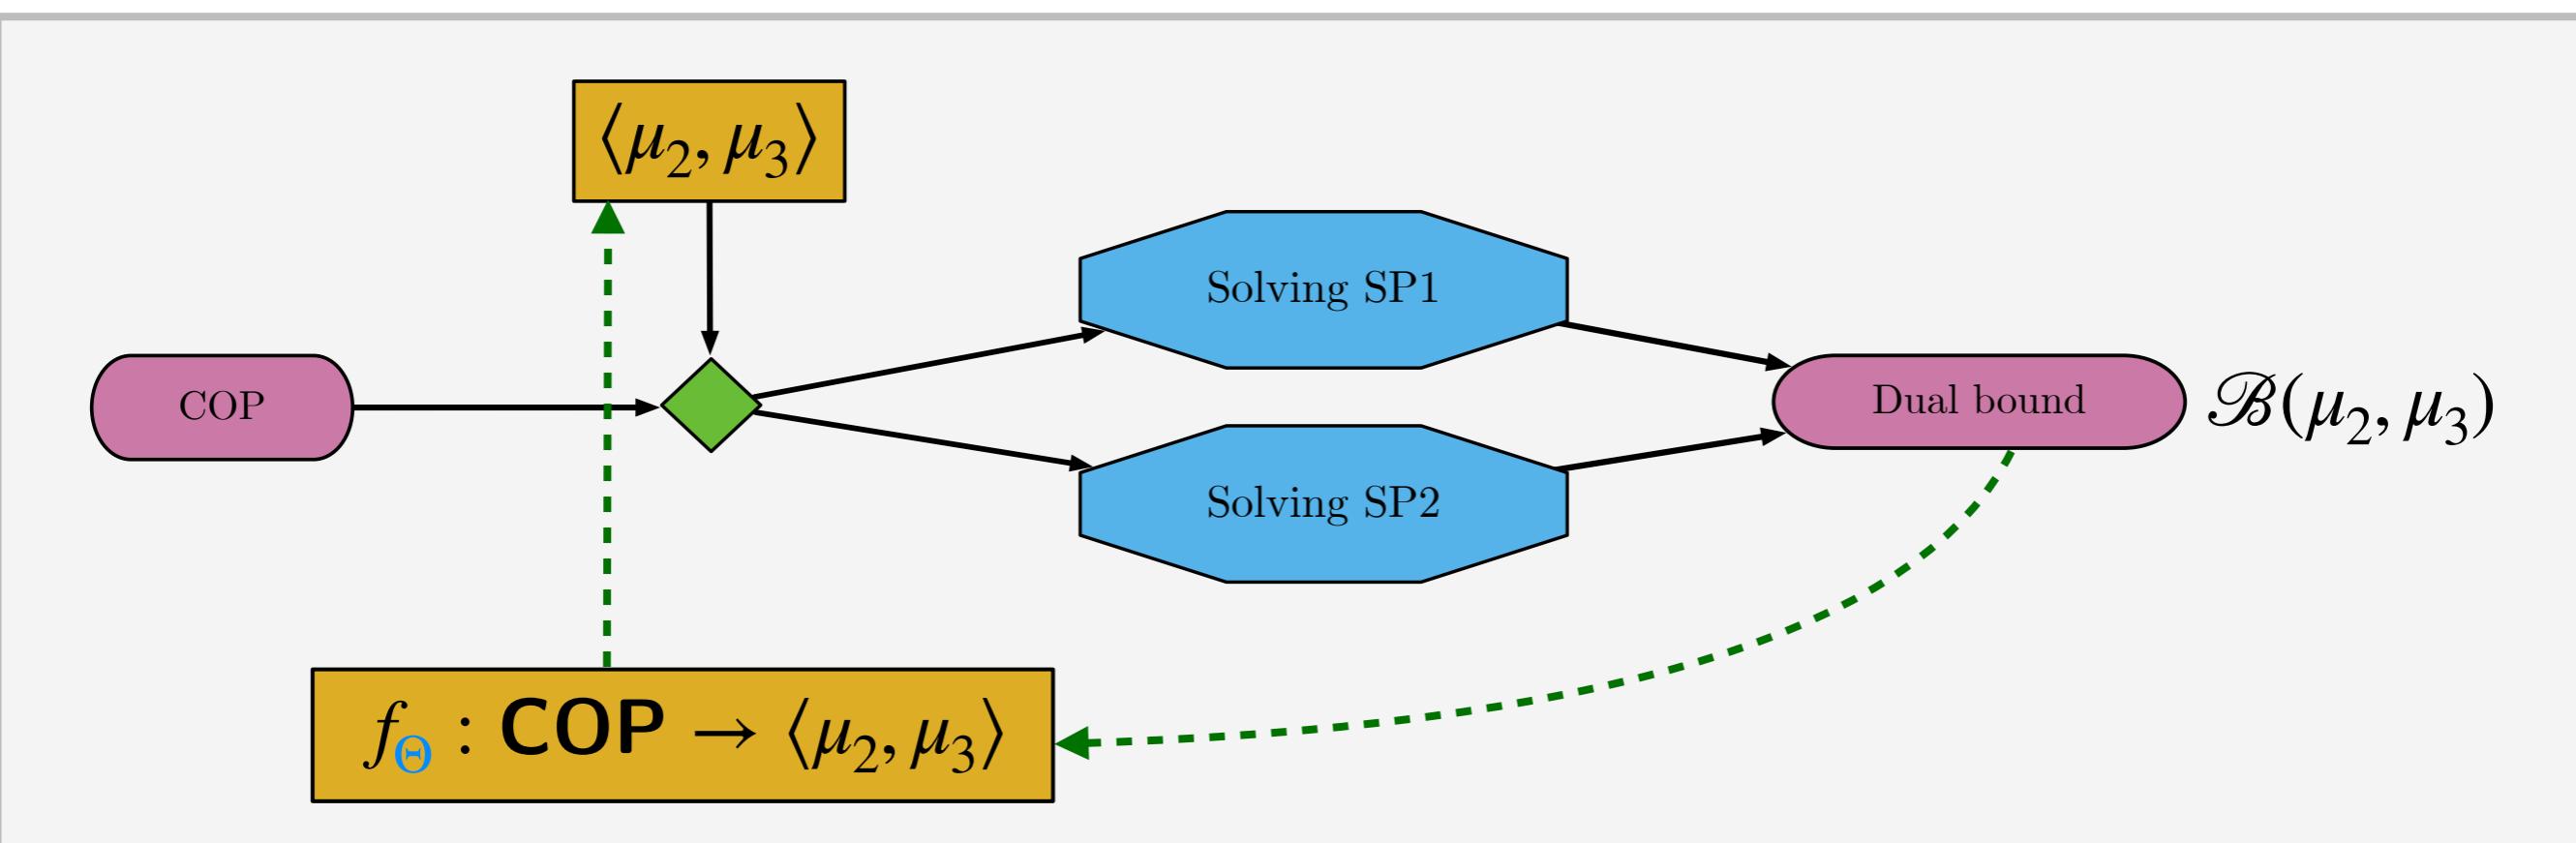
Step 2: the model is trained end-to-end by differentiating the bound



Intuition: the optimisation process is carried out offline during a training phase

# Learning multipliers for Lagrangian decomposition

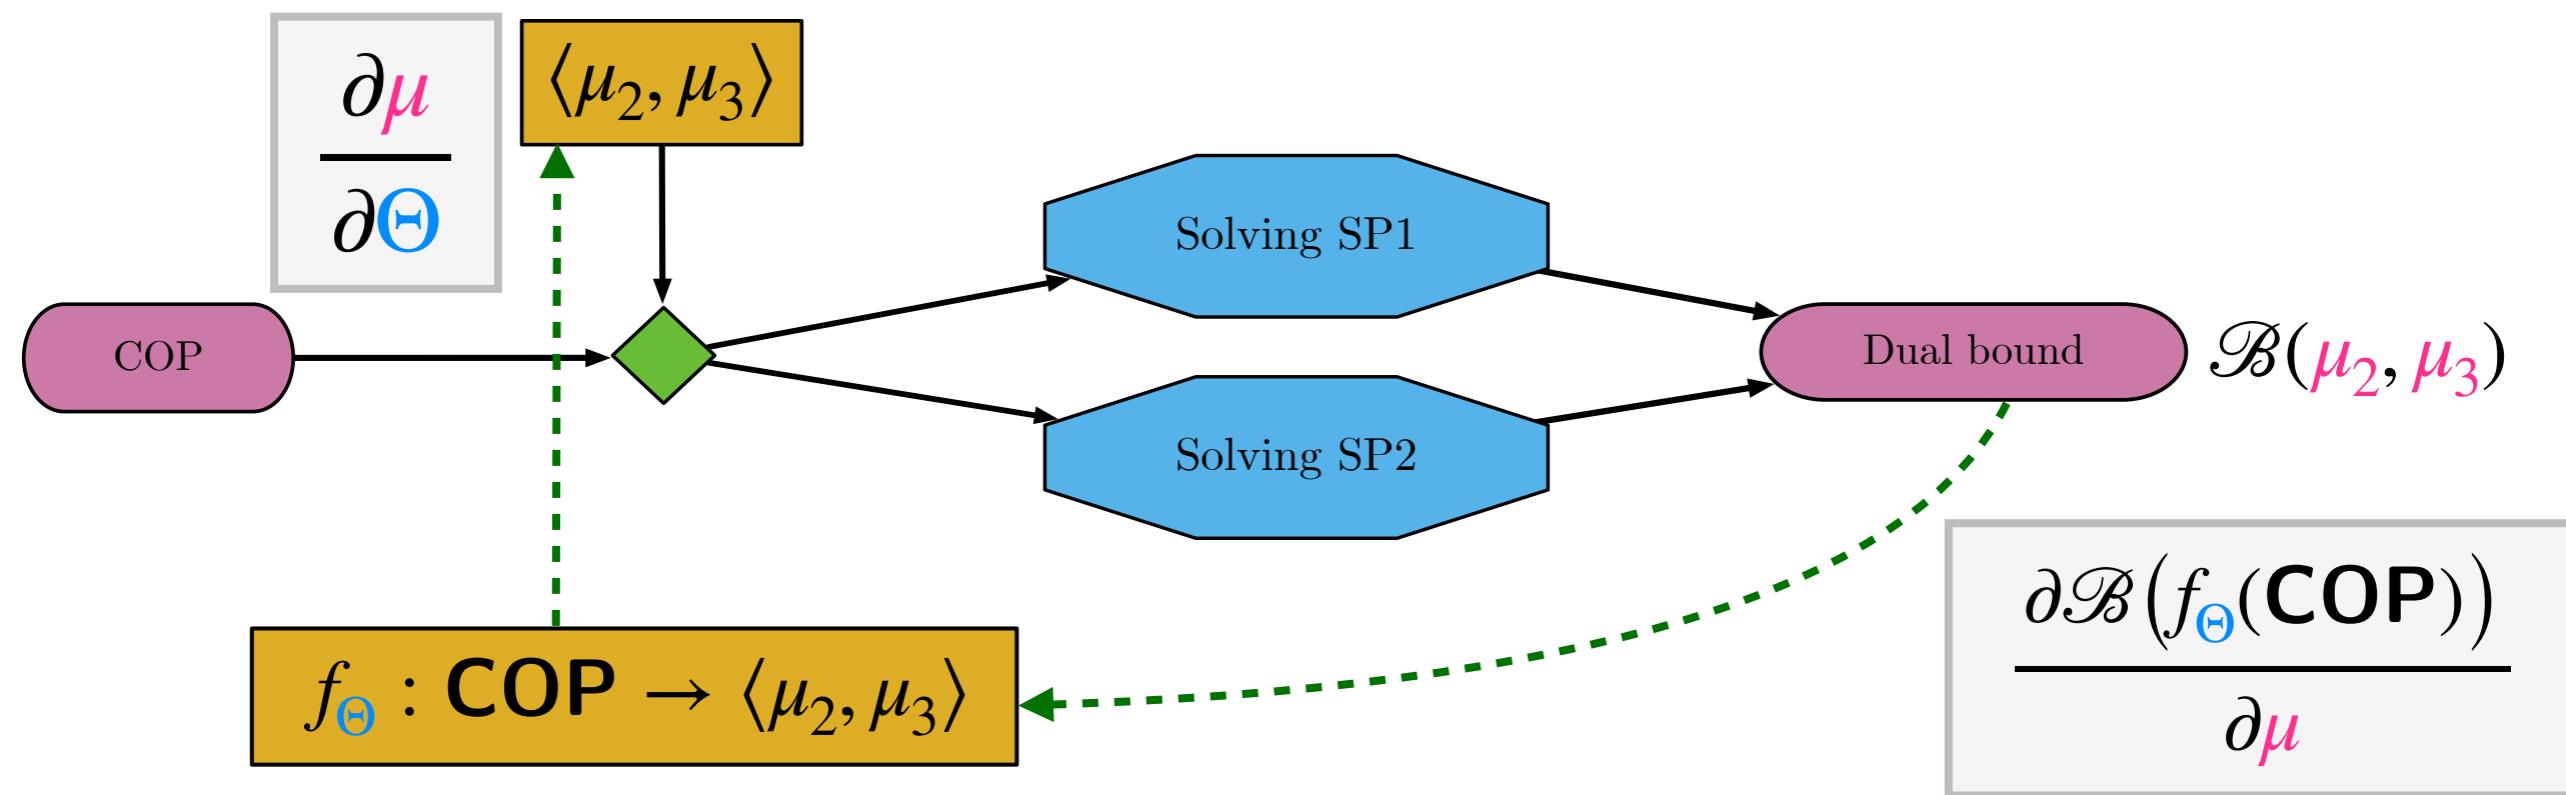
$$\mathcal{B}(\mu_2, \mu_3) = \begin{array}{l} \max \left( f(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) + \mu_2 \cdot \textcolor{violet}{X}_2 + \mu_3 \cdot \textcolor{green}{X}_3 \right) \\ \text{s.t. } C_1(X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3) \\ X_1, \textcolor{violet}{X}_2, \textcolor{green}{X}_3 \in \mathbb{N}^+ \end{array} + \begin{array}{l} \max \left( -\mu_2 \cdot \textcolor{magenta}{Y}_2 - \mu_3 \cdot \textcolor{blue}{Y}_3 \right) \\ \text{s.t. } C_2(\textcolor{magenta}{Y}_2, \textcolor{blue}{Y}_3) \\ \textcolor{magenta}{Y}_2, \textcolor{blue}{Y}_3 \in \mathbb{N}^+ \end{array}$$



**Gradient:**  $\nabla_{\Theta} \mathcal{B}(\mu)$

? How to compute the gradient of this bound ?

# Learning multipliers for Lagrangian decomposition



$$\nabla_{\Theta} \mathcal{B}(\mu) = \frac{\partial \mathcal{B}(f_\Theta(\text{COP}))}{\partial \mu} \times \frac{\partial \mu}{\partial \Theta}$$

$$= (X - Y) \times \frac{\partial \mu}{\partial \Theta}$$

**Step 1:** we use the **chain-rule** to uncover dependencies

**Step 2:** right-term is a simple **backpropagation** in the predictive model

**Step 3:** left-term reuses the **initial sub-gradient expression**

**Training:** differentiating on training instances (no need to have labels)

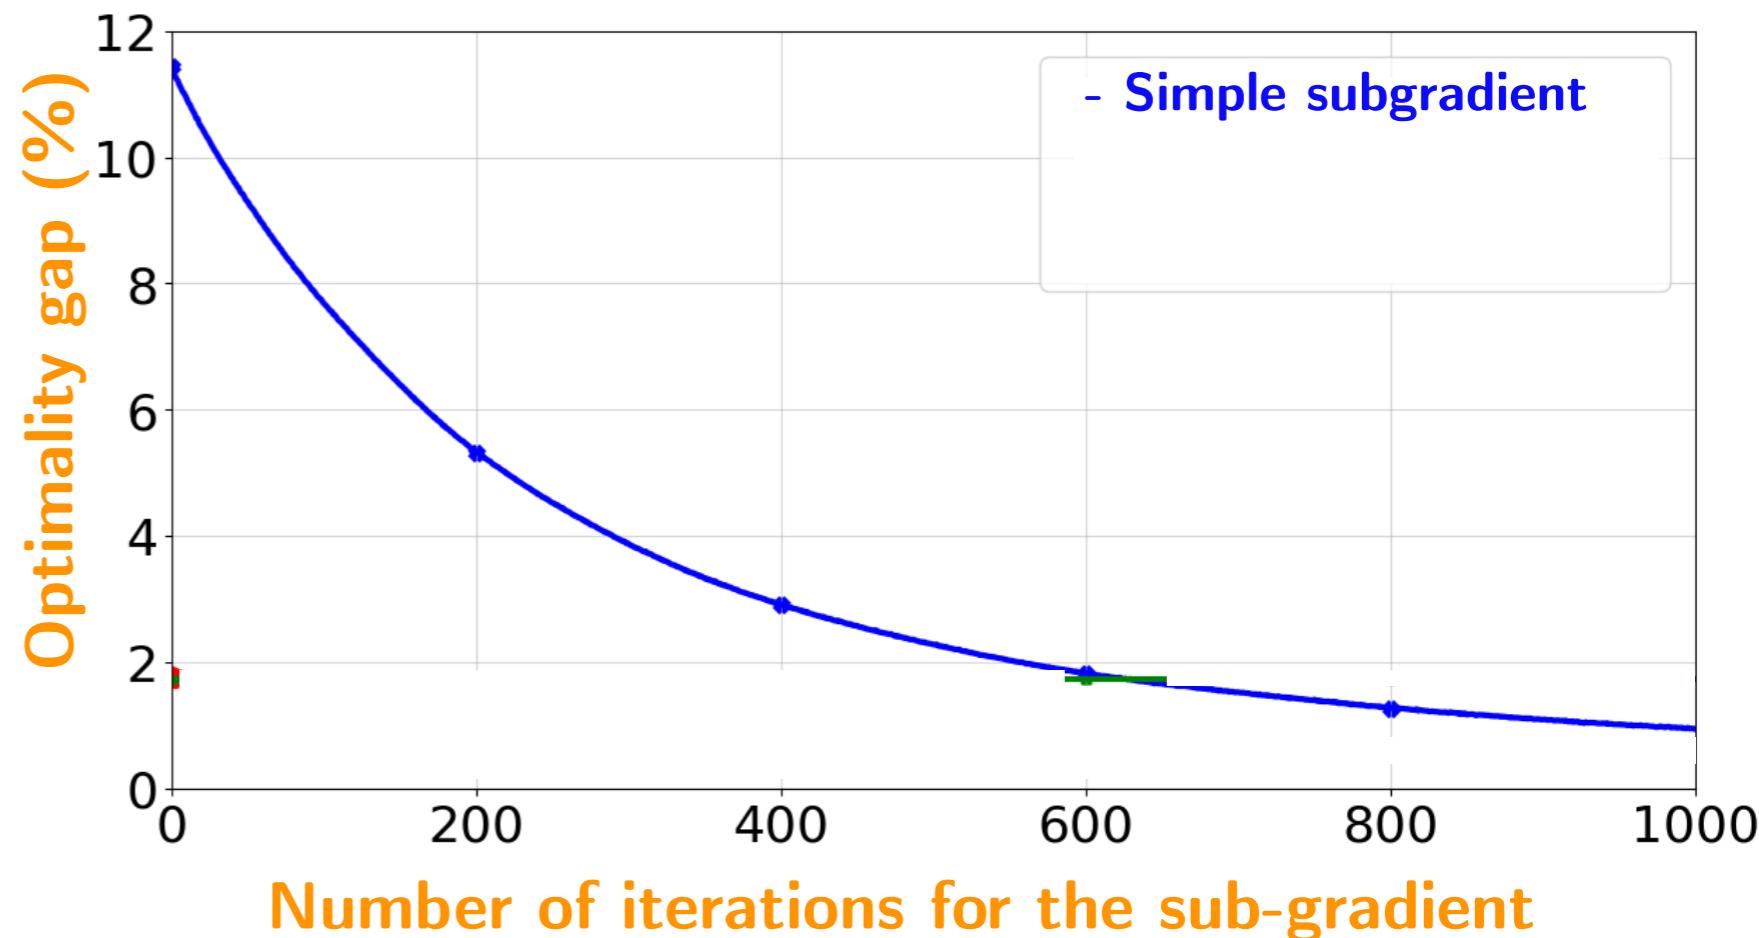
# Qualitative results

## Multidimensional knapsack



Size: 100 items

Dimension: 5 constraints



Metric: quality of the **dual bound** obtained at the root node

**Observation 1:** simple subgradient requires a lot of iterations to find good bounds

**Observation 2:** learning alone manages to get directly good bounds

**Observation 3:** learning to initialize sub-gradient quickly gives better bounds

Our results showed both the interest of learning alone or with sub-gradient

# My current thoughts in learning to bound


I think that learning to bound in CP is a promising direction

- (1) The lack of generic relaxation is a shortcoming of CP (compared to MIP solvers)
- (2) Lagrangian relaxation (or decomposition) ensures the validity of the bound
- (3) Learning can fully replace the subgradient (less filtering but cheaper)
- (4) Learning can also only initialize subgradient (better filtering but more expensive)
- (5) Results are encouraging in terms of execution time (more than for branching)

But there are still questions and challenges to address...

- (1) Which predictive model should we use? (is GNN a right choice?)
- (2) Can we have a generic representation of any COP (Boisvert et al., CPAIOR 2024)
- (3) How to handle problems with only few representative instances?
- (4) There are still the subproblems that need to be solved efficiently

# Conclusion with my final notes



## My take-home messages

- (1) Learning something meaningful does not always result in improved performances
- (2) Combining learning with existing tools mitigate the shortcomings of learning
- (3) Promising research direction, but a lot of challenges to handle for a practical use

# Many thanks!!



# Reference list (work we carried out)

## Learning to branch

- ❖ Cappart, Q., Moisan, T., Rousseau, L. M., Prémont-Schwarz, I., & Cire, A. A.  
Combining reinforcement learning and constraint programming for combinatorial optimization. [\[AAAI 2021\]](#)
- ❖ Chalumeau, F., Coulon, I., Cappart, Q., & Rousseau, L. M.  
Seapearl: A constraint programming solver guided by reinforcement learning. [\[CPAIOR 2021\]](#)
- ❖ Marty, T., François, T., Tessier, P., Gautier, L., Rousseau, L. M., & Cappart, Q.  
Learning a Generic Value-Selection Heuristic Inside a Constraint Programming Solver. [\[CP 2023 - distinguished paper\]](#)

## Learning to bound

- ❖ Cappart, Q., Bergman, D., Rousseau, L. M., Prémont-Schwarz, I., & Parjadis, A.  
Improving variable orderings of approximate decision diagrams using reinforcement learning. [\[IJOC 2022\]](#)
- ❖ Parjadis, A., Cappart, Q., Dilkina, B., Ferber, A., & Rousseau, L. M.  
Learning Lagrangian Multipliers for the Travelling Salesman Problem. [\[CP 2024 - Best ML paper award\]](#)
- ❖ Dabert, D., Bessa, S., Bourgeat, M., Rousseau, L.M., Cappart, Q.  
Learning Valid Dual Bounds in Constraint Programming [\[ArXivPreprint 2024\]](#)

## Learning to model and graph neural networks

- ❖ Cappart, Q., Chételat, D., Khalil, E. B., Lodi, A., Morris, C., & Veličković, P.  
Combinatorial optimization and reasoning with graph neural networks. [\[IJCAI 2021, JMLR 2023\]](#)
- ❖ Boisvert, L., Verhaeghe, H., & Cappart, Q.  
Towards a Generic Representation of Combinatorial Problems for Learning-Based Approaches. [\[CPAIOR 2024\]](#)
- ❖ Barral, H., Gaha, M., Dems, A., Côté, A., Nguewouo, F., & Cappart, Q.  
Acquiring Constraints for a Non-linear Transmission Maintenance Scheduling Problem. [\[CPAIOR 2024\]](#)

# Reference list (other related works)

## Learning to branch (also in MIP)

- ❖ Khalil, E., Le Bodic, P., Song, L., Nemhauser, G., & Dilkina, B.  
Learning to branch in mixed integer programming. [AAAI 2016]
- ❖ Gasse, M., Chételat, D., Ferroni, N., Charlin, L., & Lodi, A.  
Exact combinatorial optimization with graph convolutional neural networks. [NeurIPS 2019]
- ❖ Song, W., Cao, Z., Zhang, J., Xu, C., & Lim, A.  
Learning variable ordering heuristics for solving constraint satisfaction problems. [EAAA 2022]

## Foundations of our work on learning to prune

- ❖ Held, M., & Karp, R. The traveling-salesman problem and minimum spanning trees [Operations Research 1970]
- ❖ Guignard, M., & Kim, S. Lagrangean decomposition: A model yielding stronger Lagrangean bounds. [Math. Prog. 1987]
- ❖ Focacci, F., Lodi, A., & Milano, M. Cost-based domain filtering [CP 1999]
- ❖ Bergman, D., Cire, A. A., & van Hoeve, W. J. Improved constraint propagation via lagrangian decomposition. [CP 2015]
- ❖ Hà, M. H., Quimper, C. G., & Rousseau, L. M. General bounding mechanism for constraint programs. [CP 2015]

## Learning dual bounds (only outside CP?)

- ❖ Deng, Y., Kong, S., Liu, C., & An, B.  
Deep attentive belief propagation: Integrating reasoning and learning for solving constraint optimization problems. [NeurIPS 2022]
- ❖ Abbas, A., & Swoboda, P. Doge-train: Discrete optimization on gpu with end-to-end training. [AAAI 2024]

Please reach me if you know other works using learning to get dual bounds :-)




Research group: [corail-research.github.io](https://corail-research.github.io)

Personal page: [qcappart.github.io](https://qcappart.github.io)

Slides: QR code (or via my personal page)

Email: [quentin.cappart@polymtl.ca](mailto:quentin.cappart@polymtl.ca)



POLYTECHNIQUE  
MONTRÉAL

CORAIL



Combinatorial Optimization and  
Reasoning in  
Artificial Intelligence  
Laboratory



ACP  
Quentin Cappart